

Cloud Computing Technologies

DAT515 - Fall 2024

Docker

Prof. Hein Meling

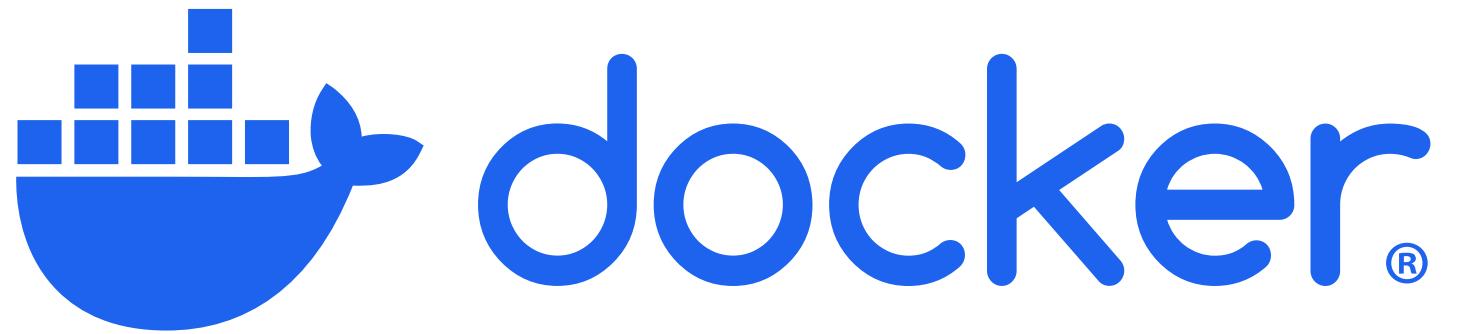




docker

16:9

What is Docker?



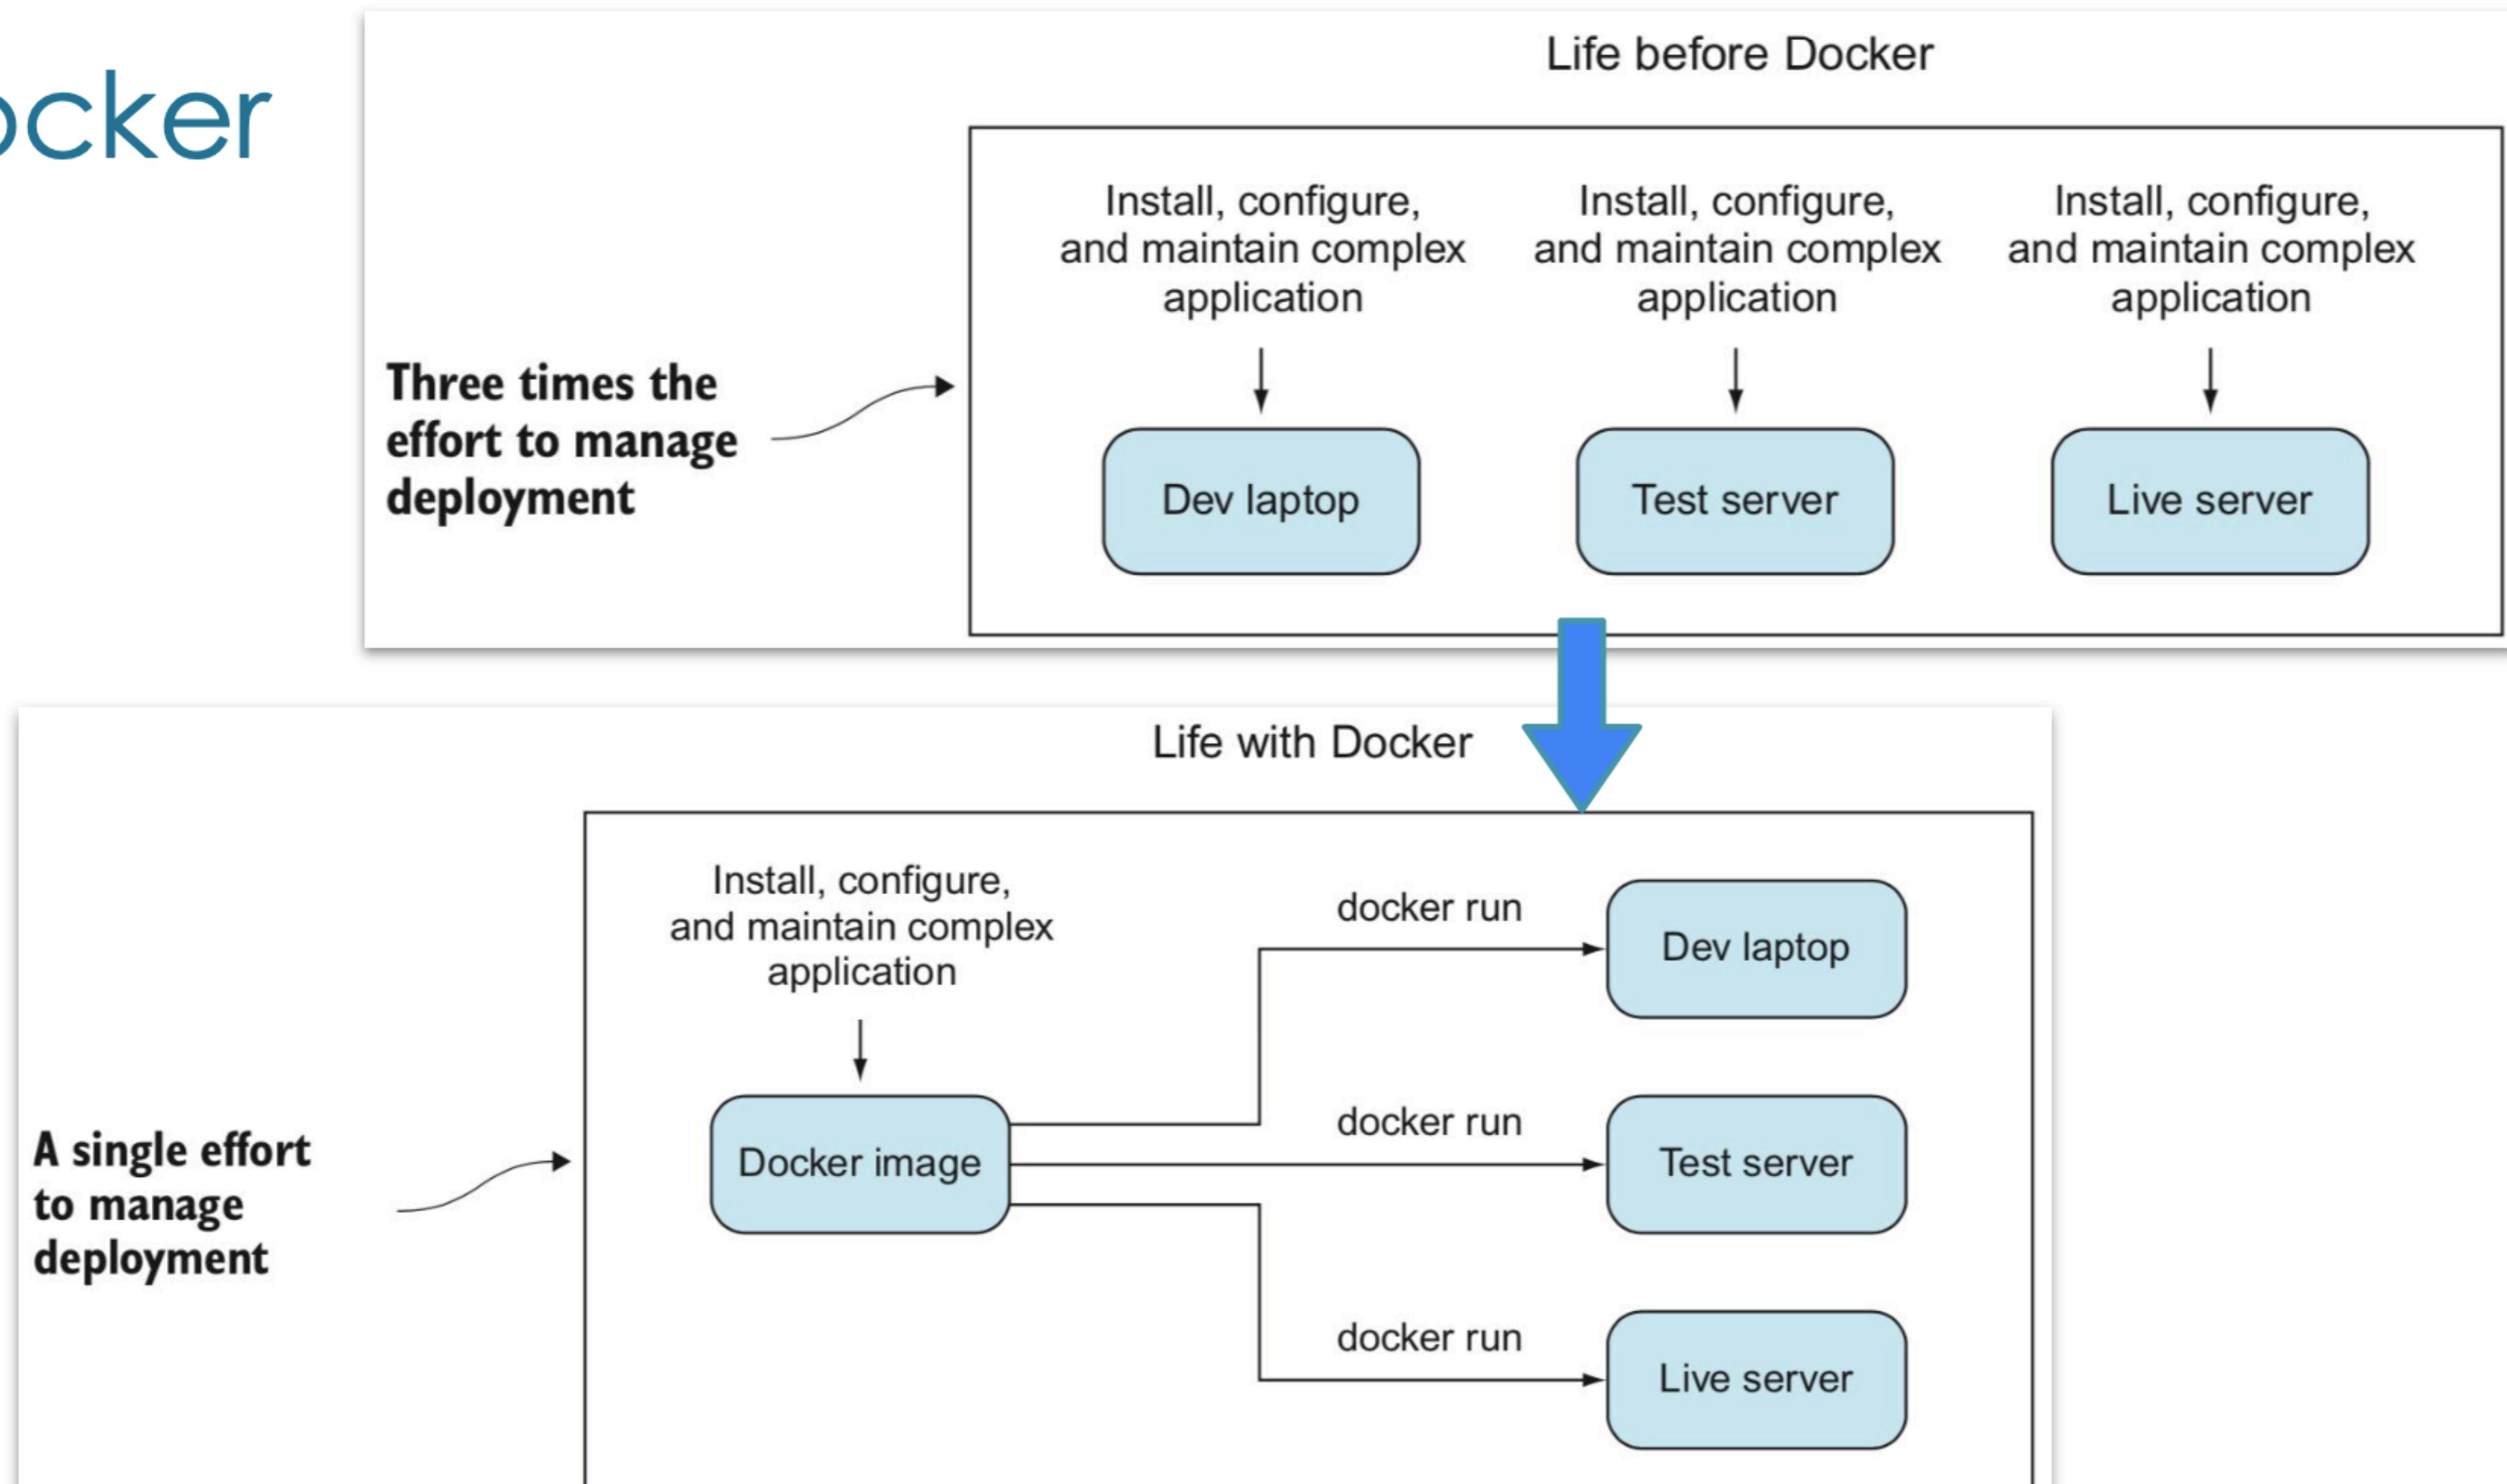
- Docker was introduced by Solomon Hykes in 2013
- Initially seen as yet another virtualization tool
- But Docker is a containerization technology
- Revolutionized development and deployment workflows
- Automating and streamlining the deployment process

The word “Docker” comes from a British colloquialism meaning dock worker – somebody who loads and unloads ships.

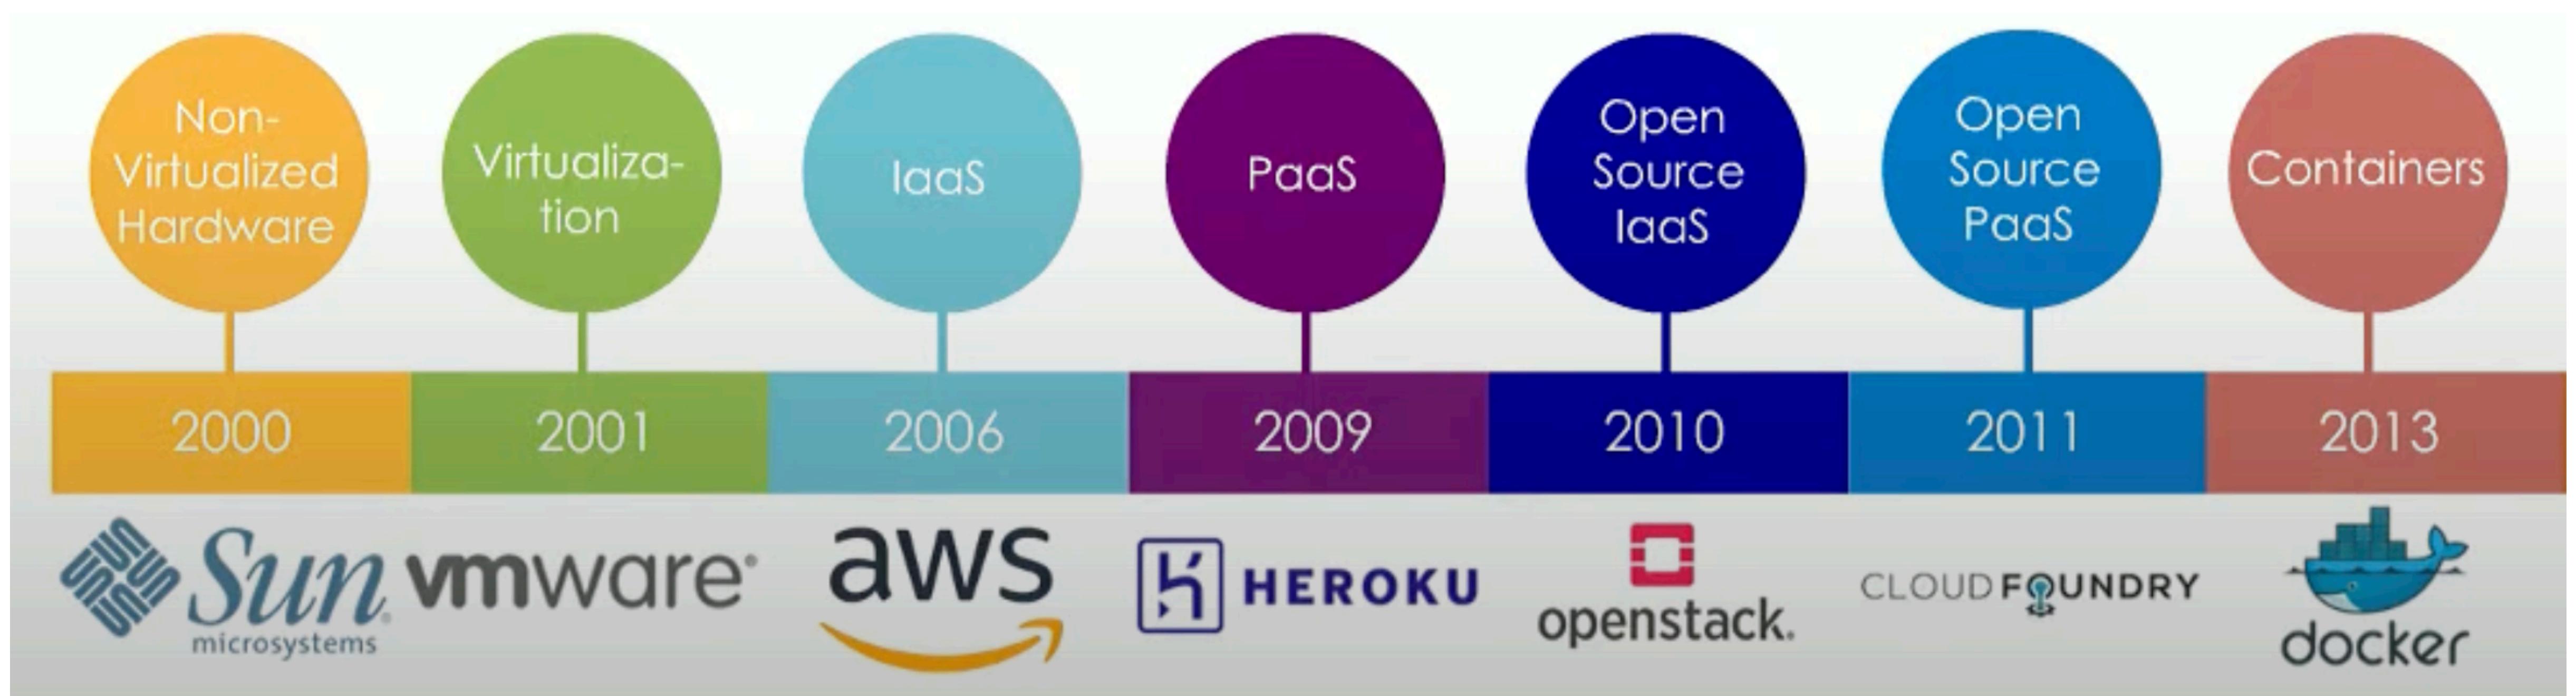
Why is Docker Important?

- Docker provides an abstraction to simplify development and deployment
- Cross platform and open source
 - Runs on Windows, macOS, and Linux...
 - Runs on several processor architectures (amd64, arm64, ...)
- Community and company adoption
 - Amazon, Microsoft, and Google

Docker



Another View of History



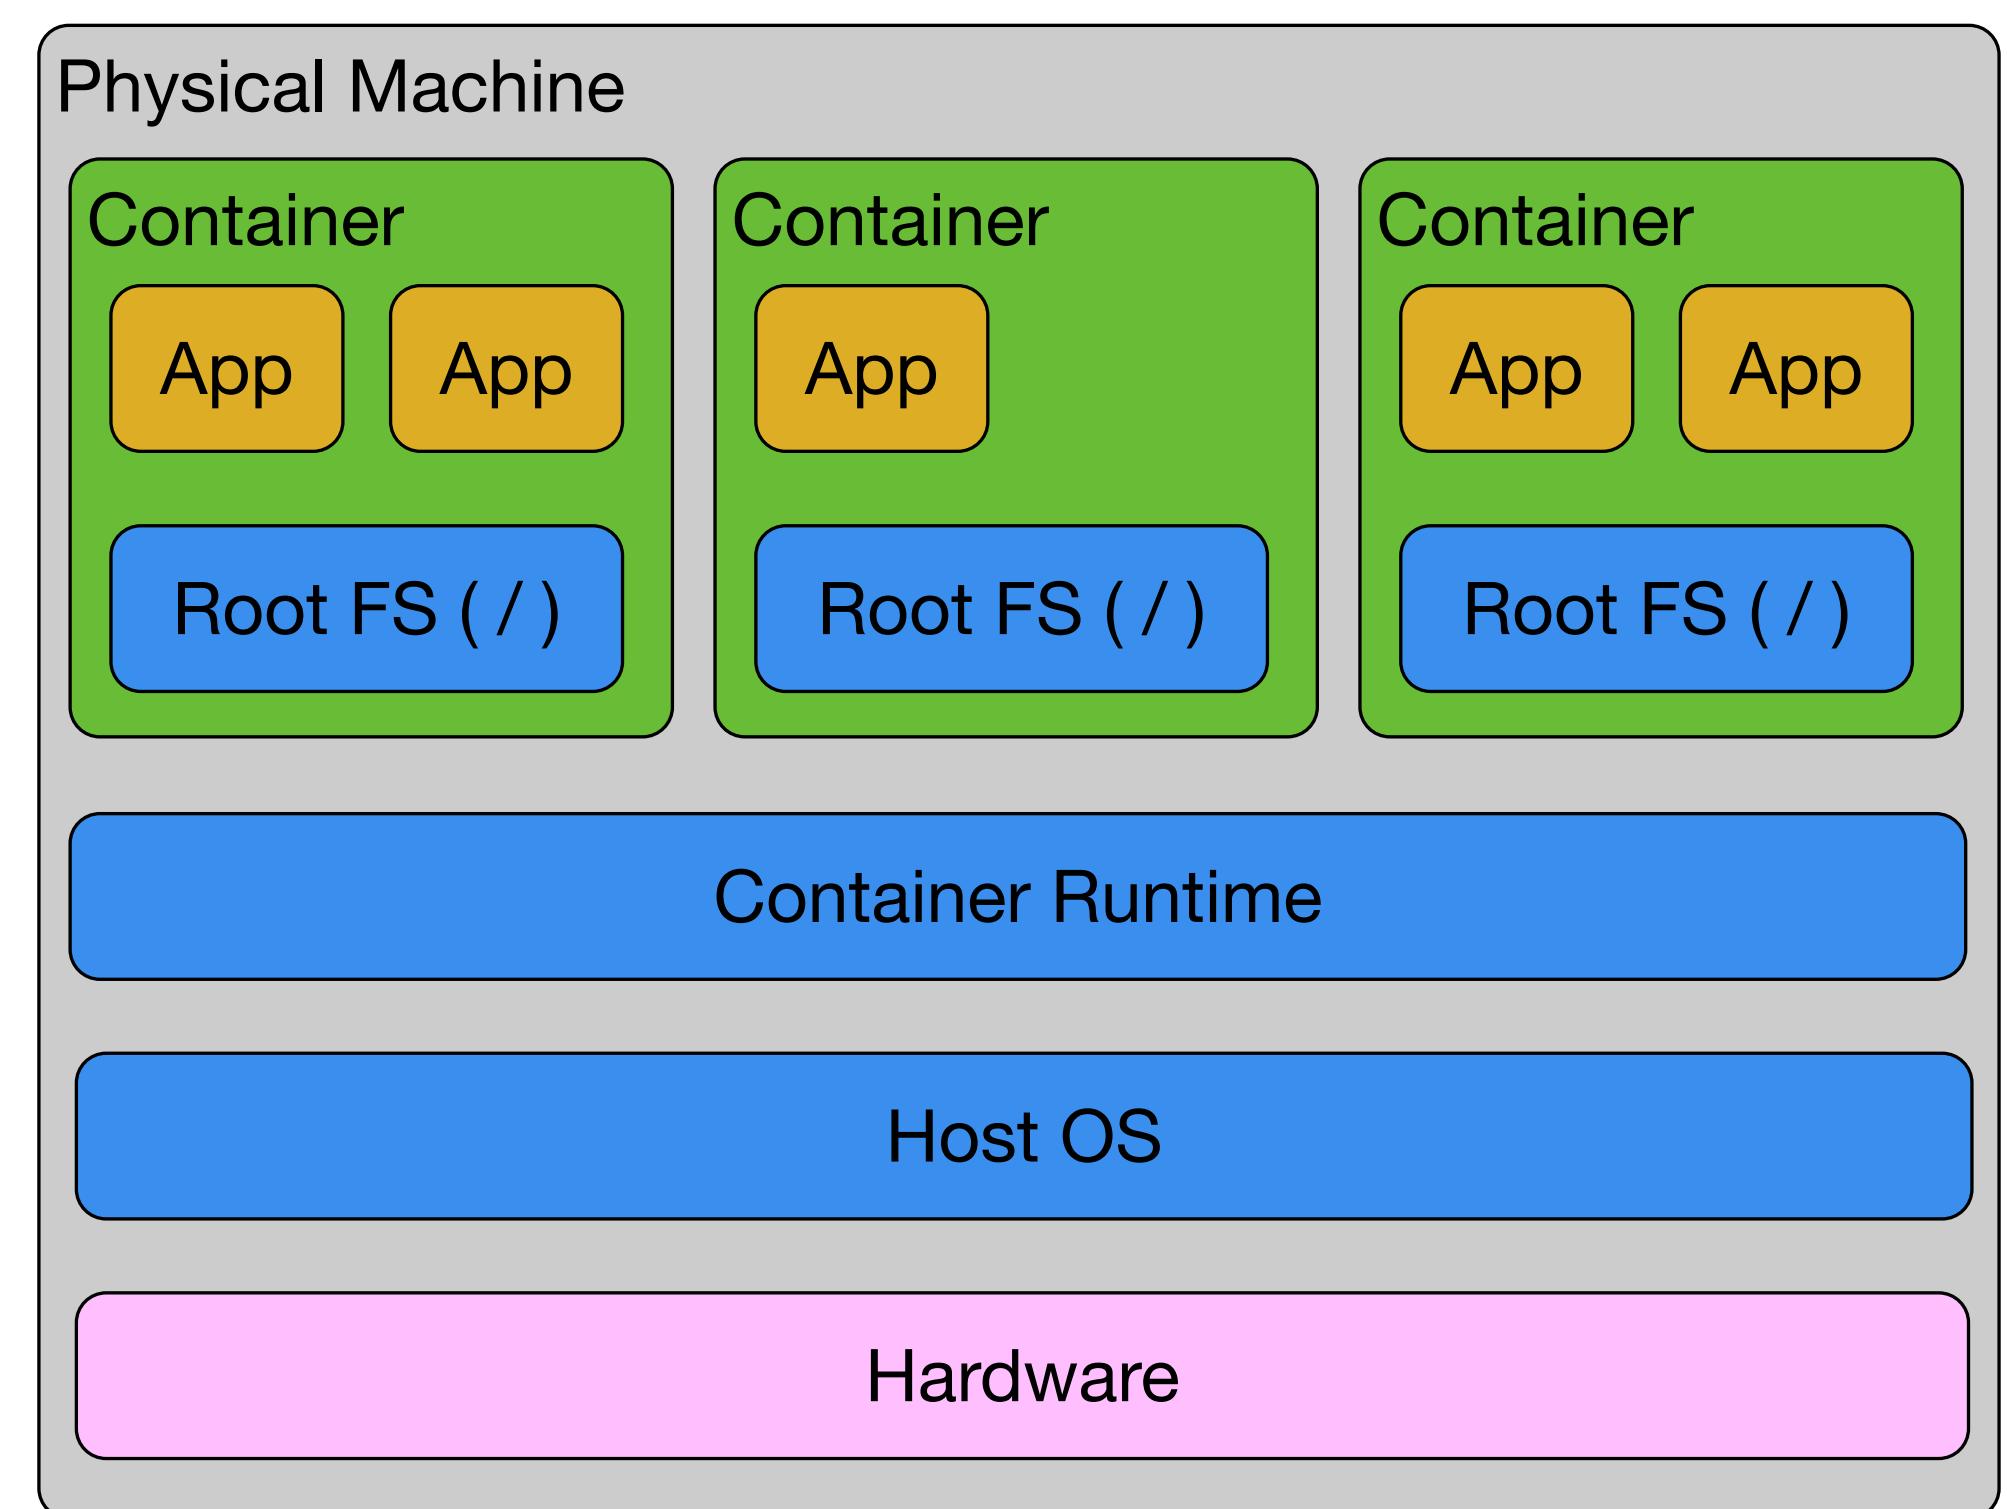
Docker has become
synonymous with Containers

What is Docker, Technically?

Lightweight, standalone, and executable software package

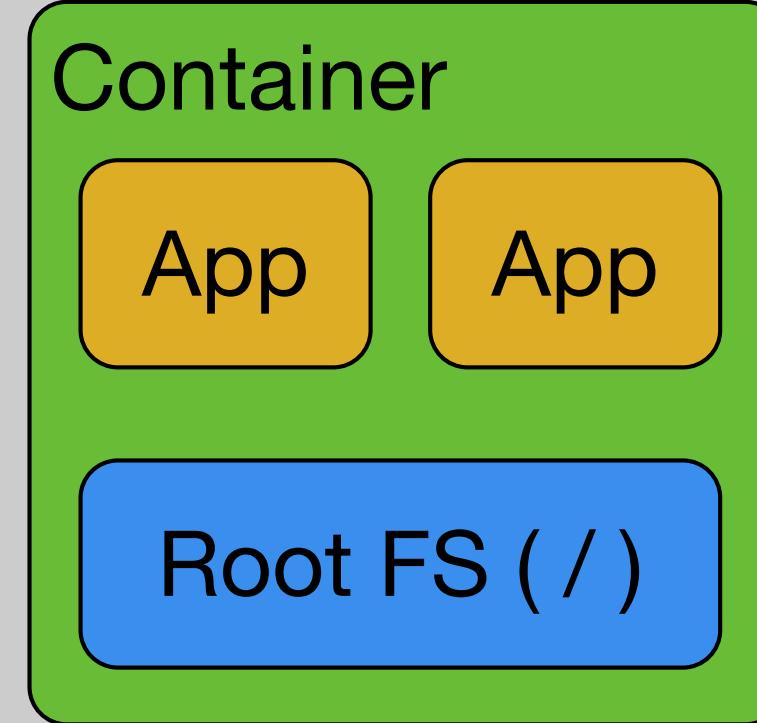
- Application code, Configuration files, Runtime libraries, Dependencies
- Pre-configured runtime environment
- Executable with a single command

Write once, run anywhere!

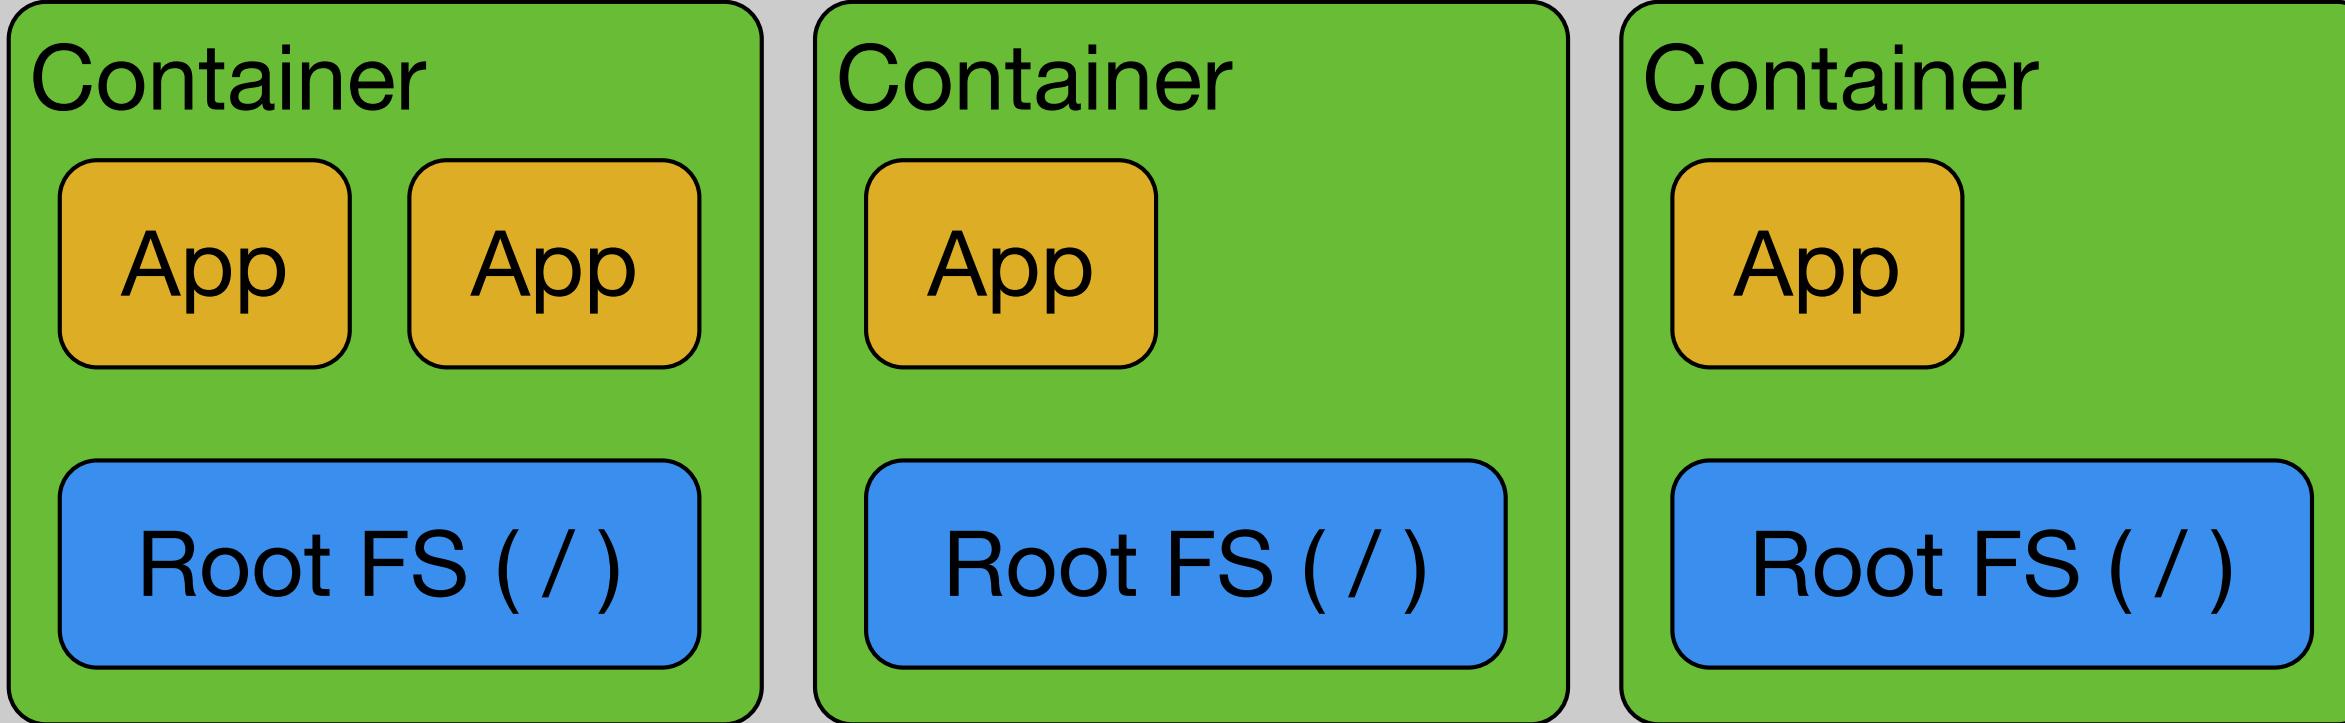


Containers inside Virtual Machines

Virtual Machine



Virtual Machine



OS

Virtual Hardware

OS

Virtual Hardware

Hypervisor

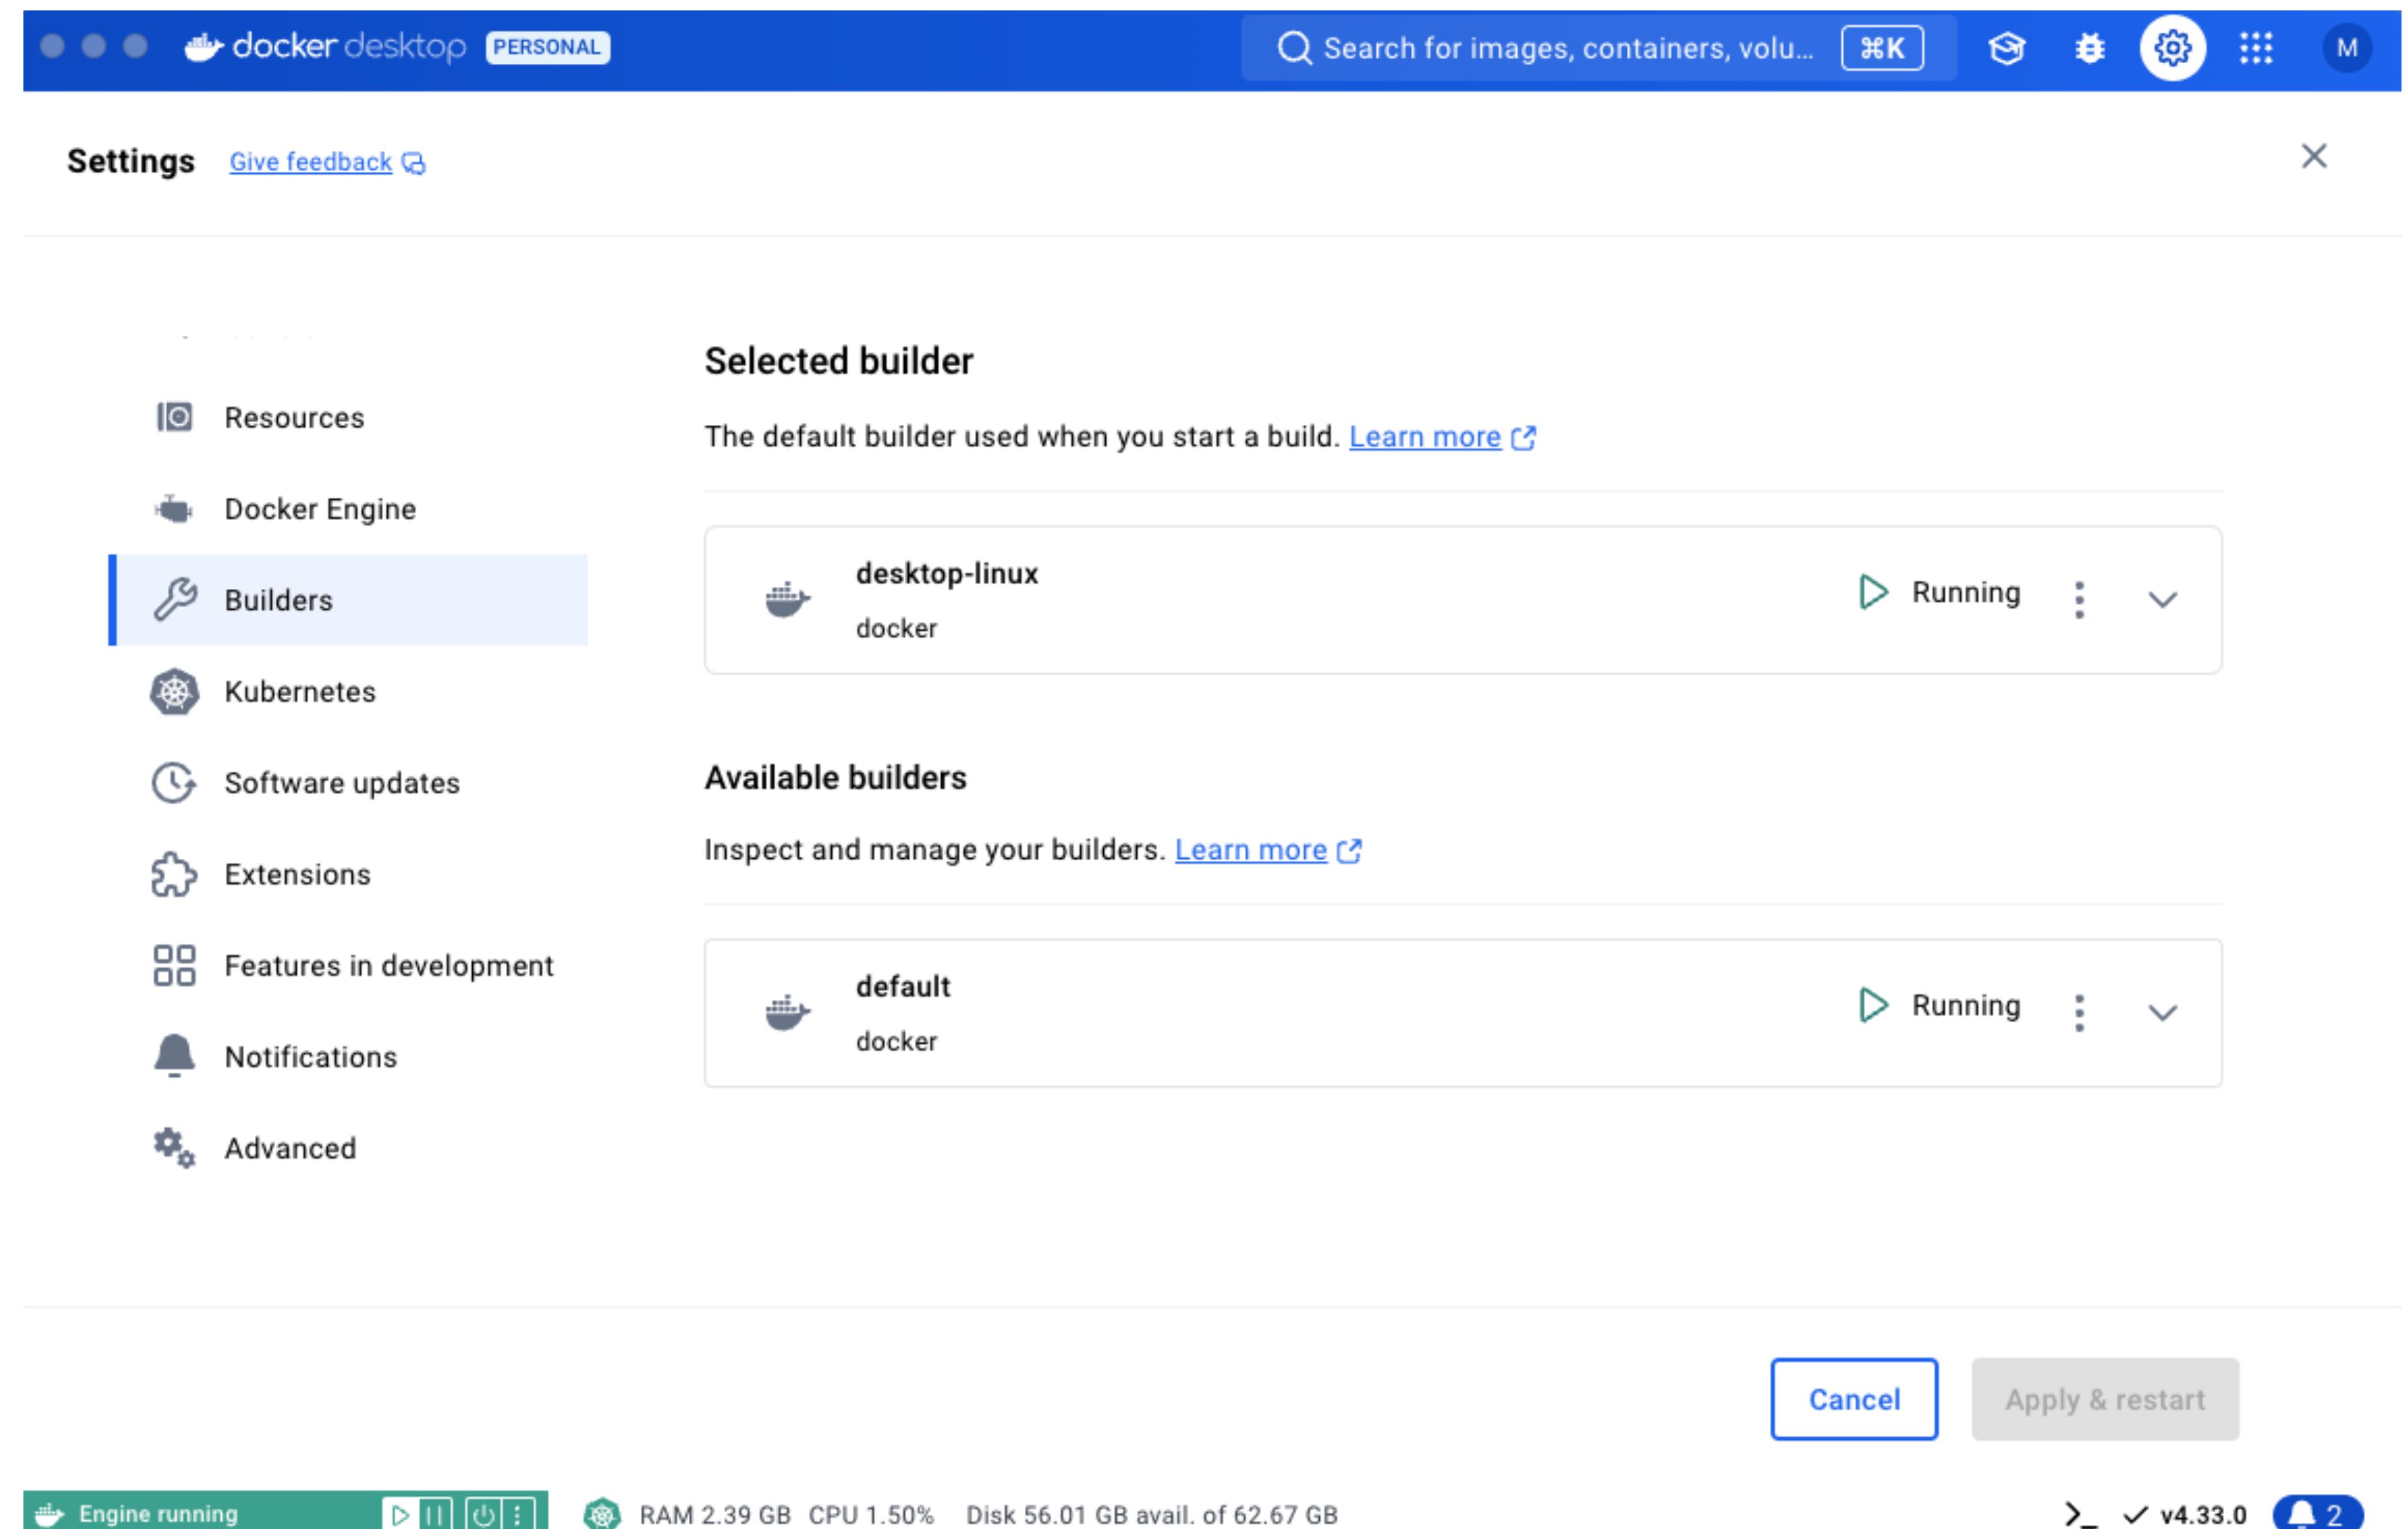
Hardware

What's the Difference?

- **Containers:** lightweight packaging of an application and its dependencies
 - Can run consistently across different environments
 - Share the host system's kernel, making them more efficient
- **Docker** is a platform and set of tools
 - Simplify the creation, deployment, and management of containers
 - Easy-to-use interface, standardized container format
 - Comprehensive ecosystem for building, distributing, and running containers

What's the Difference?

- Docker is based on the **runC** container runtime
 - Portable and compliant with the Open Container Initiative (OCI) standards
- LXC is tightly coupled to Linux environments (not portable)
- Docker is more than LXC
 - Command line tools and Desktop Client



The screenshot shows the Docker Desktop Settings interface. The left sidebar has a "Builders" section selected, indicated by a blue highlight. Other options include Resources, Docker Engine, Kubernetes, Software updates, Extensions, Features in development, Notifications, and Advanced.

Selected builder
The default builder used when you start a build. [Learn more ↗](#)

 desktop-linux	 Running	
docker		

Available builders
Inspect and manage your builders. [Learn more ↗](#)

 default	 Running	
docker		

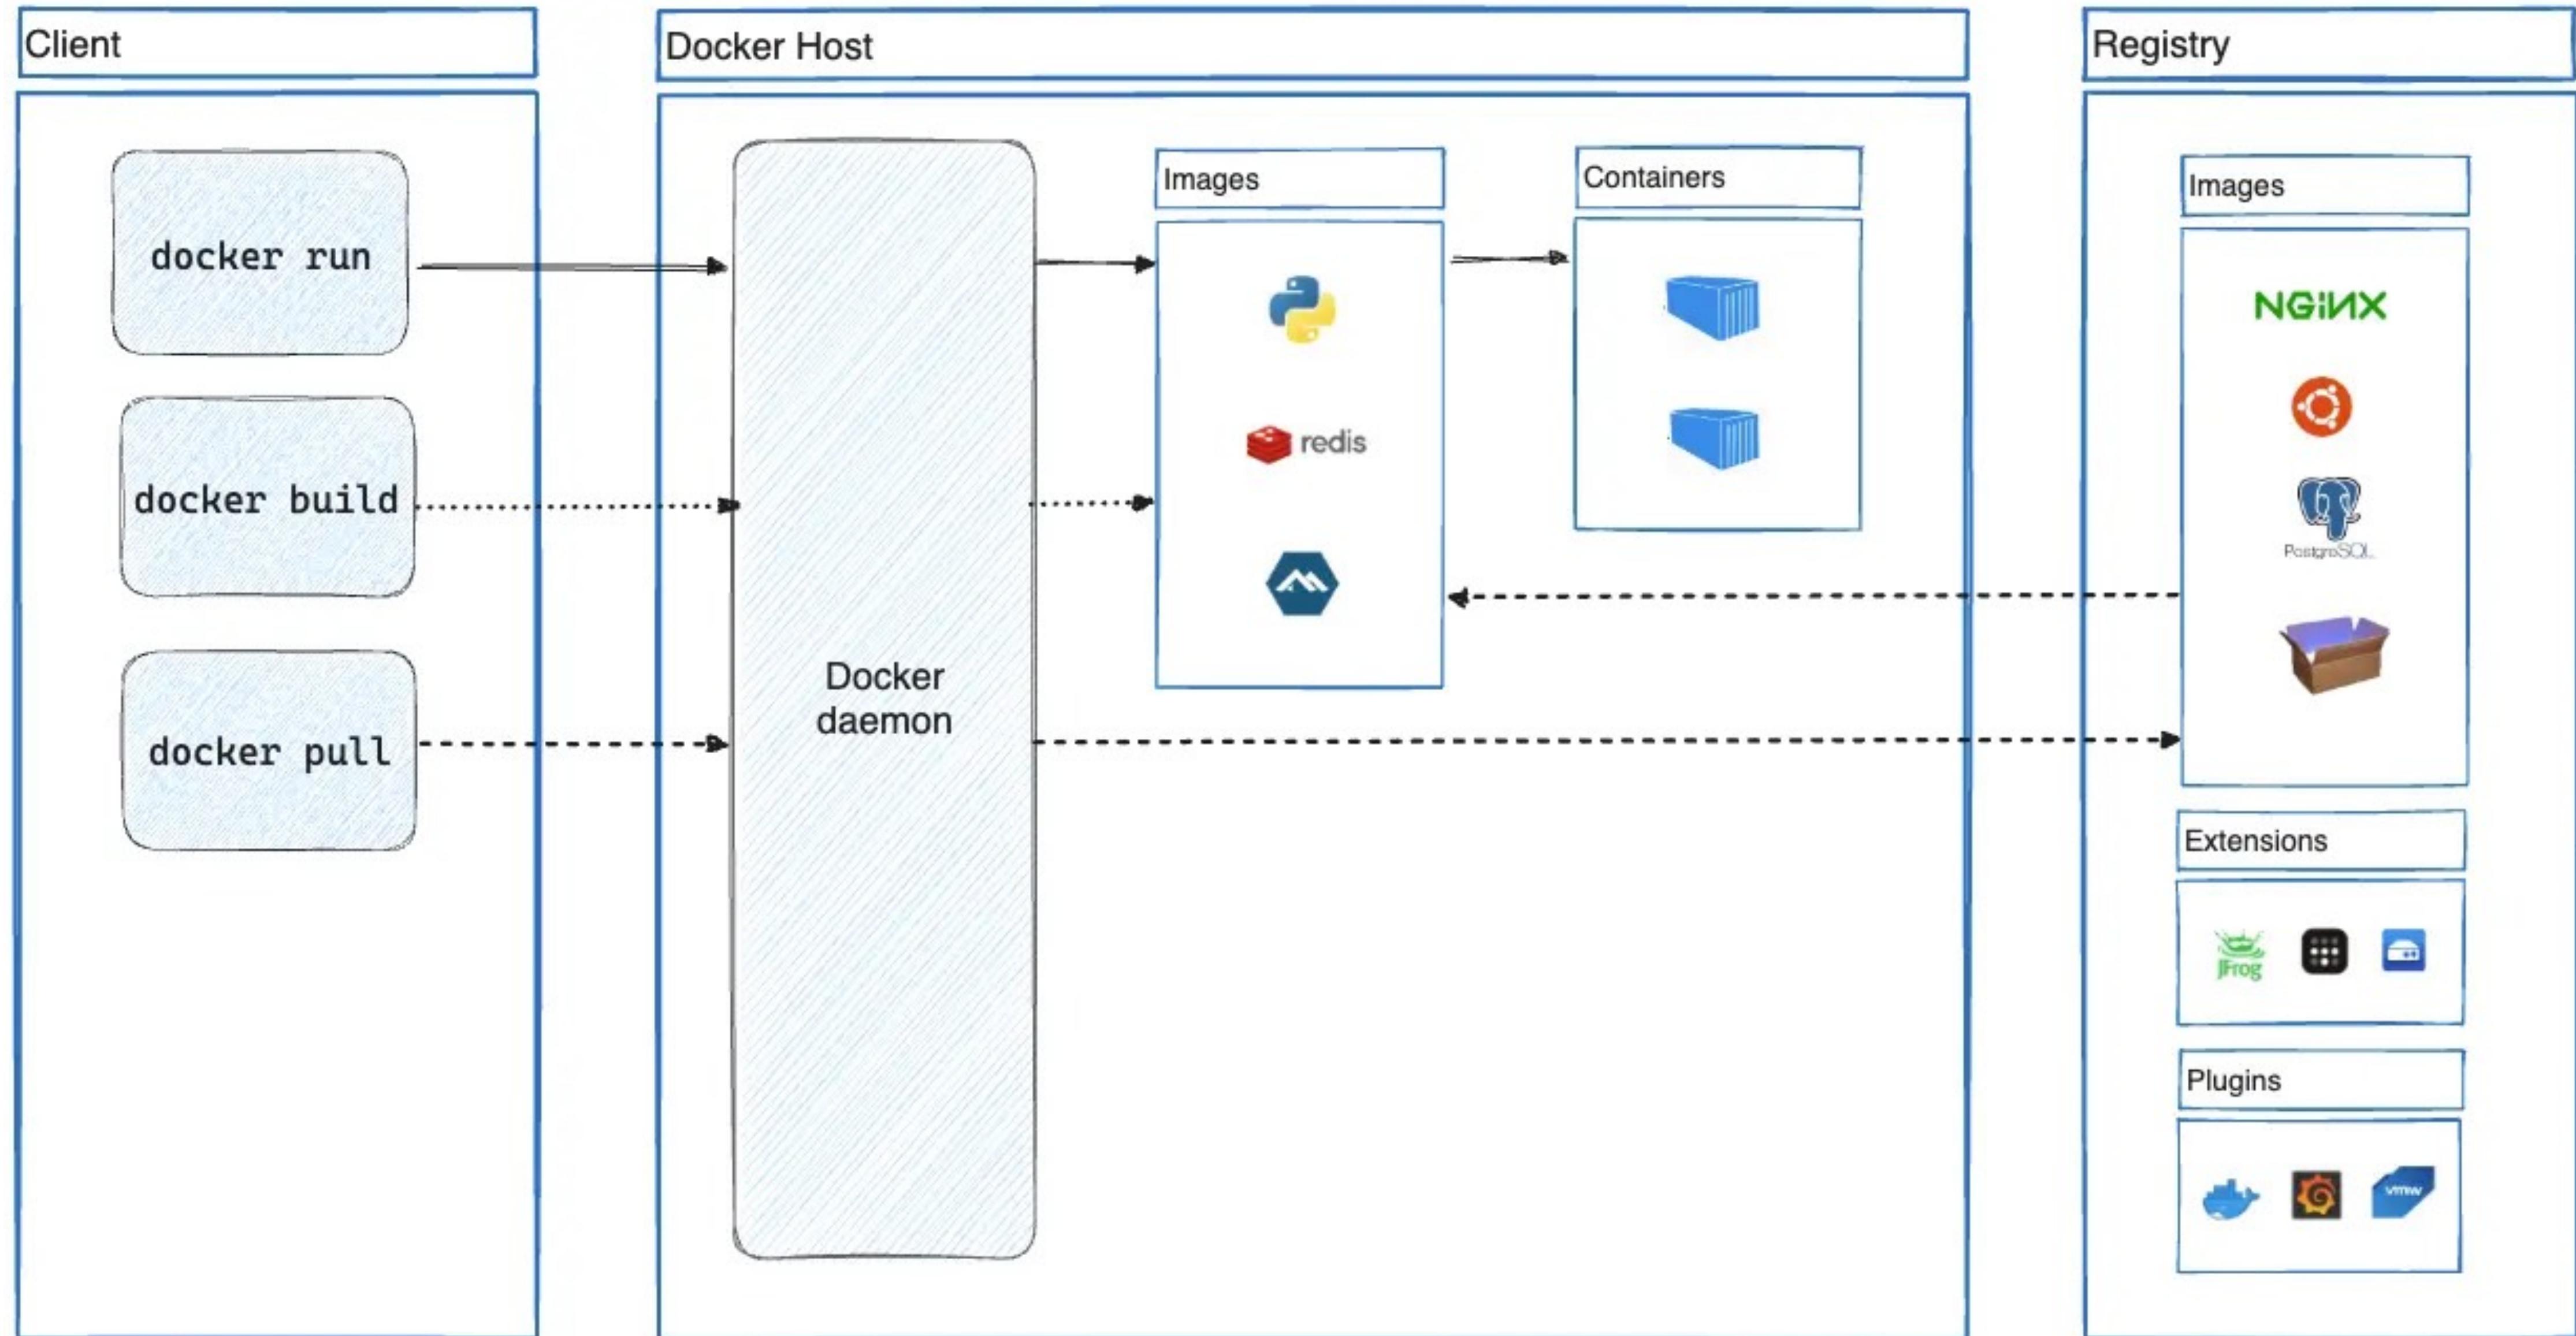
Buttons

[Cancel](#) [Apply & restart](#)

System status

Engine running | RAM 2.39 GB CPU 1.50% Disk 56.01 GB avail. of 62.67 GB | v4.33.0 | 2 notifications

Docker Architecture



Joke time

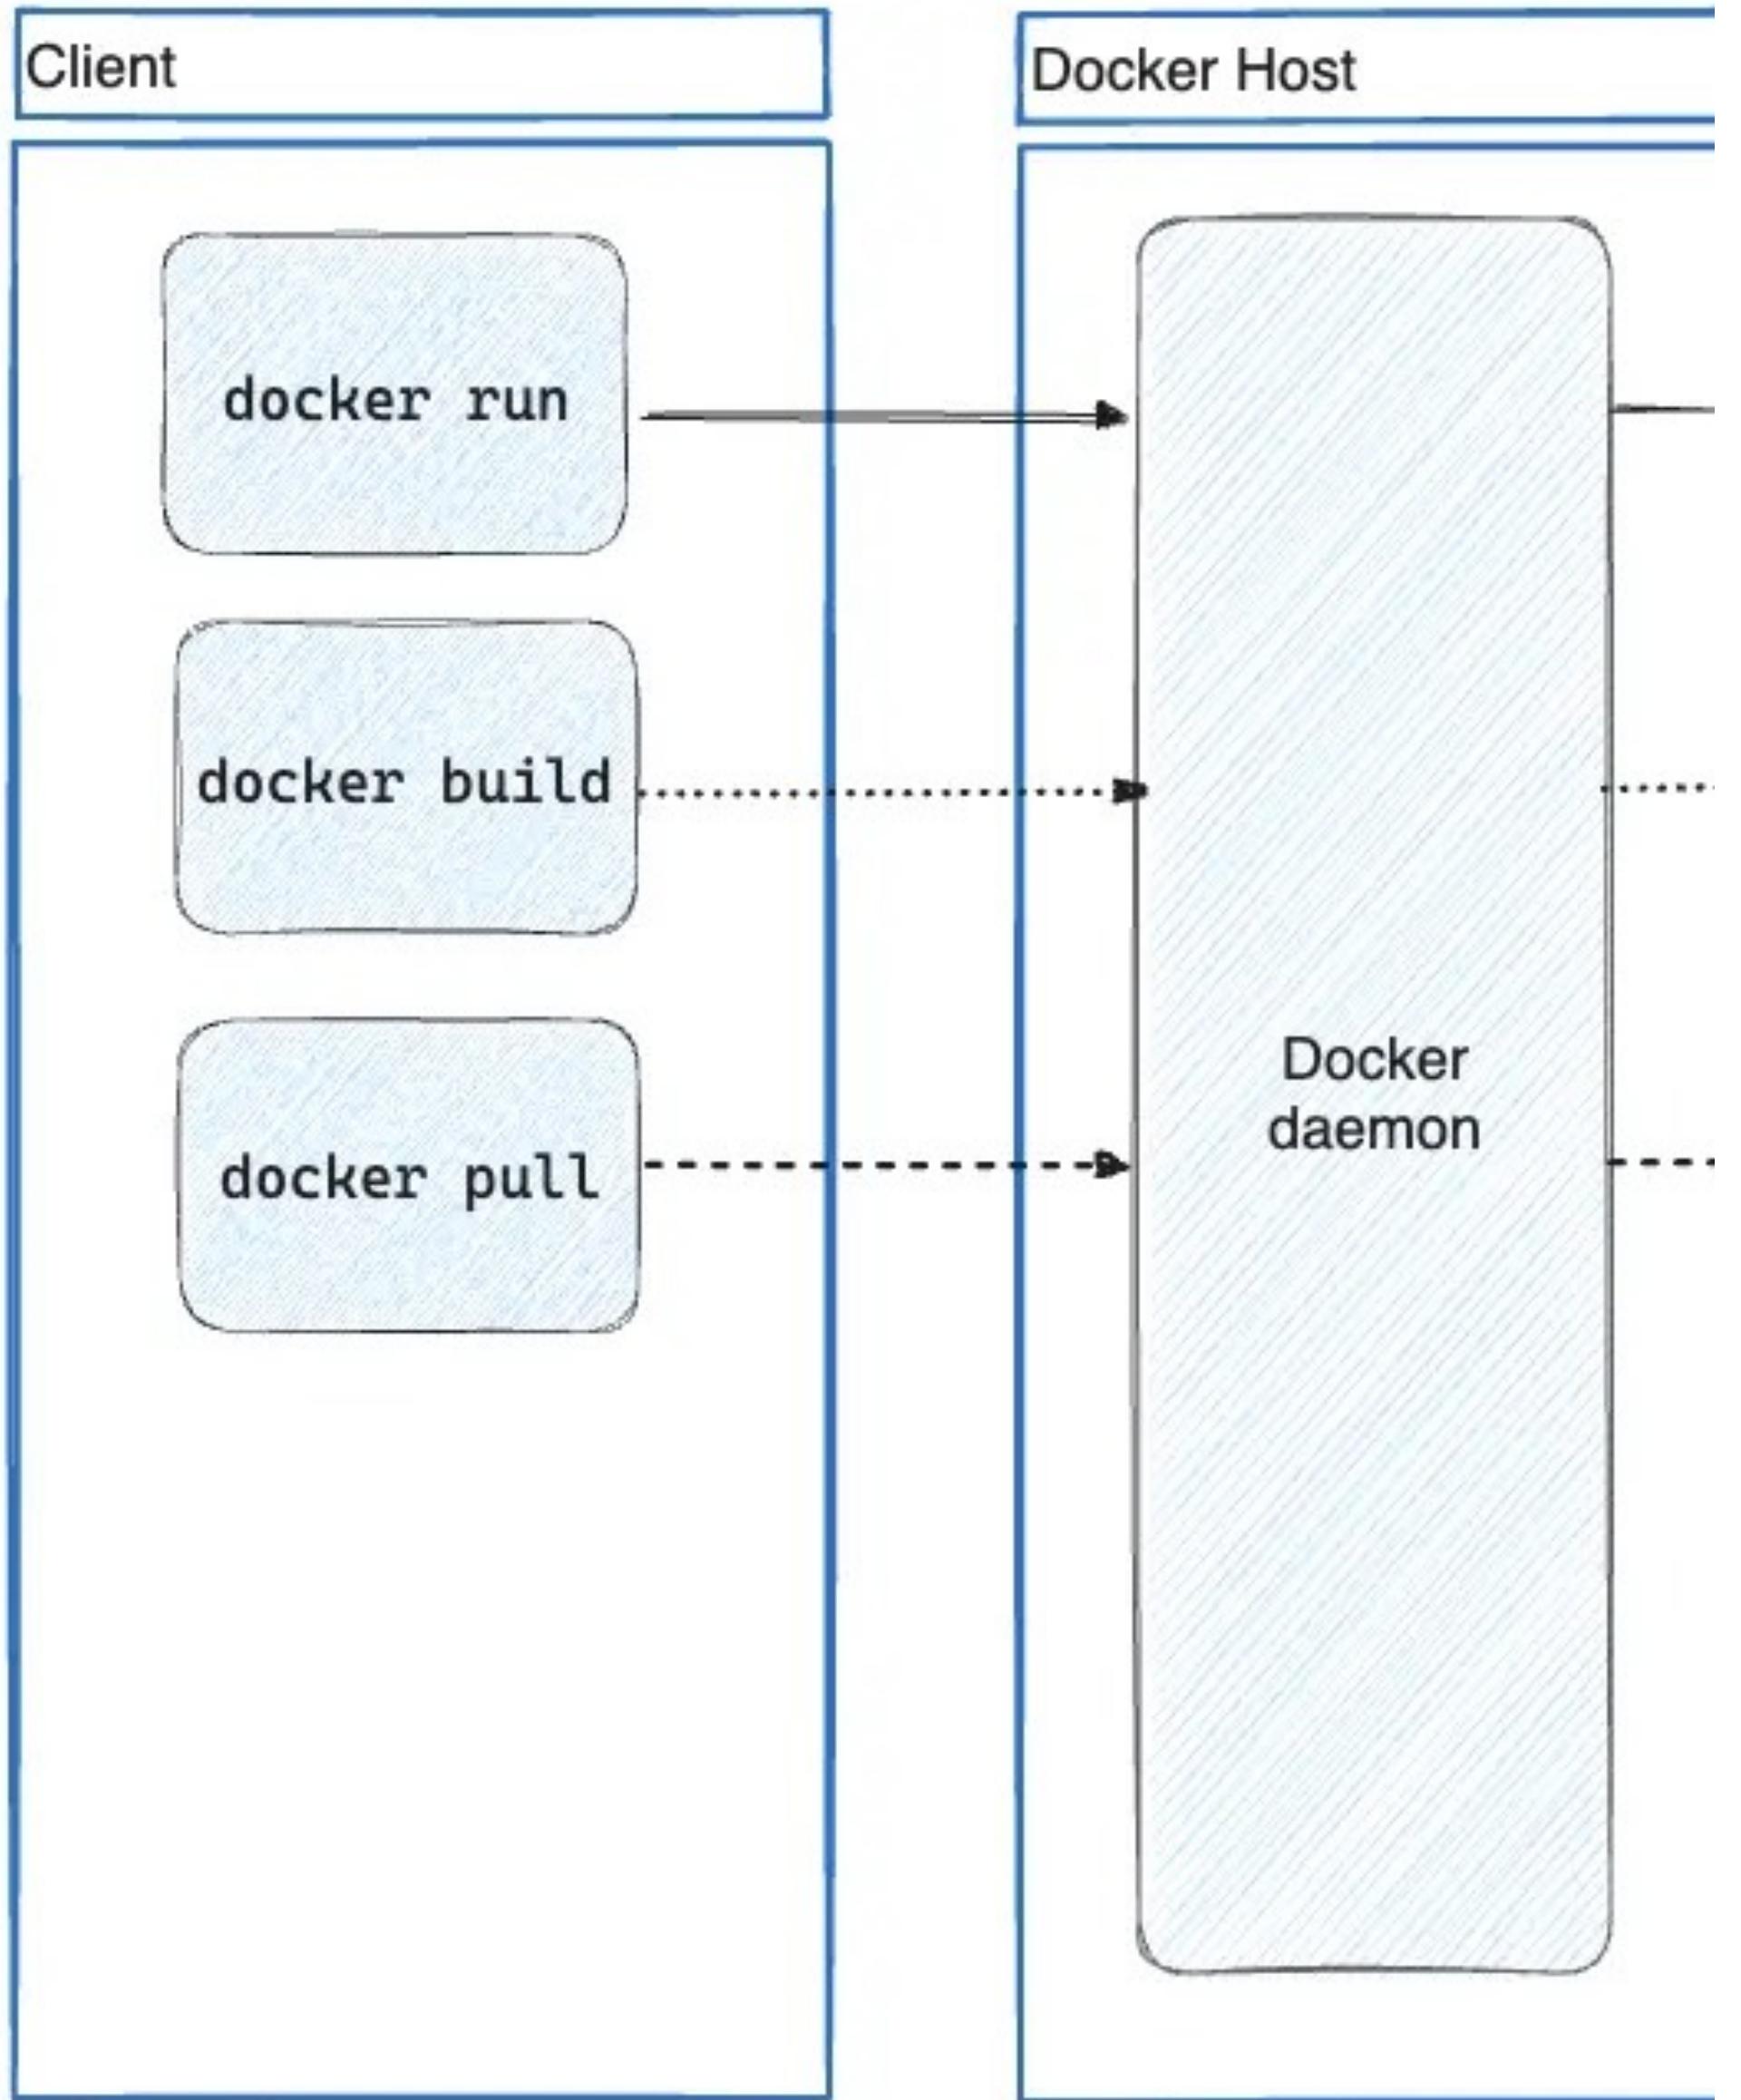
Why did the Docker container
break up with the virtual machine?

It needed some space and
felt they were too heavy

Docker Components

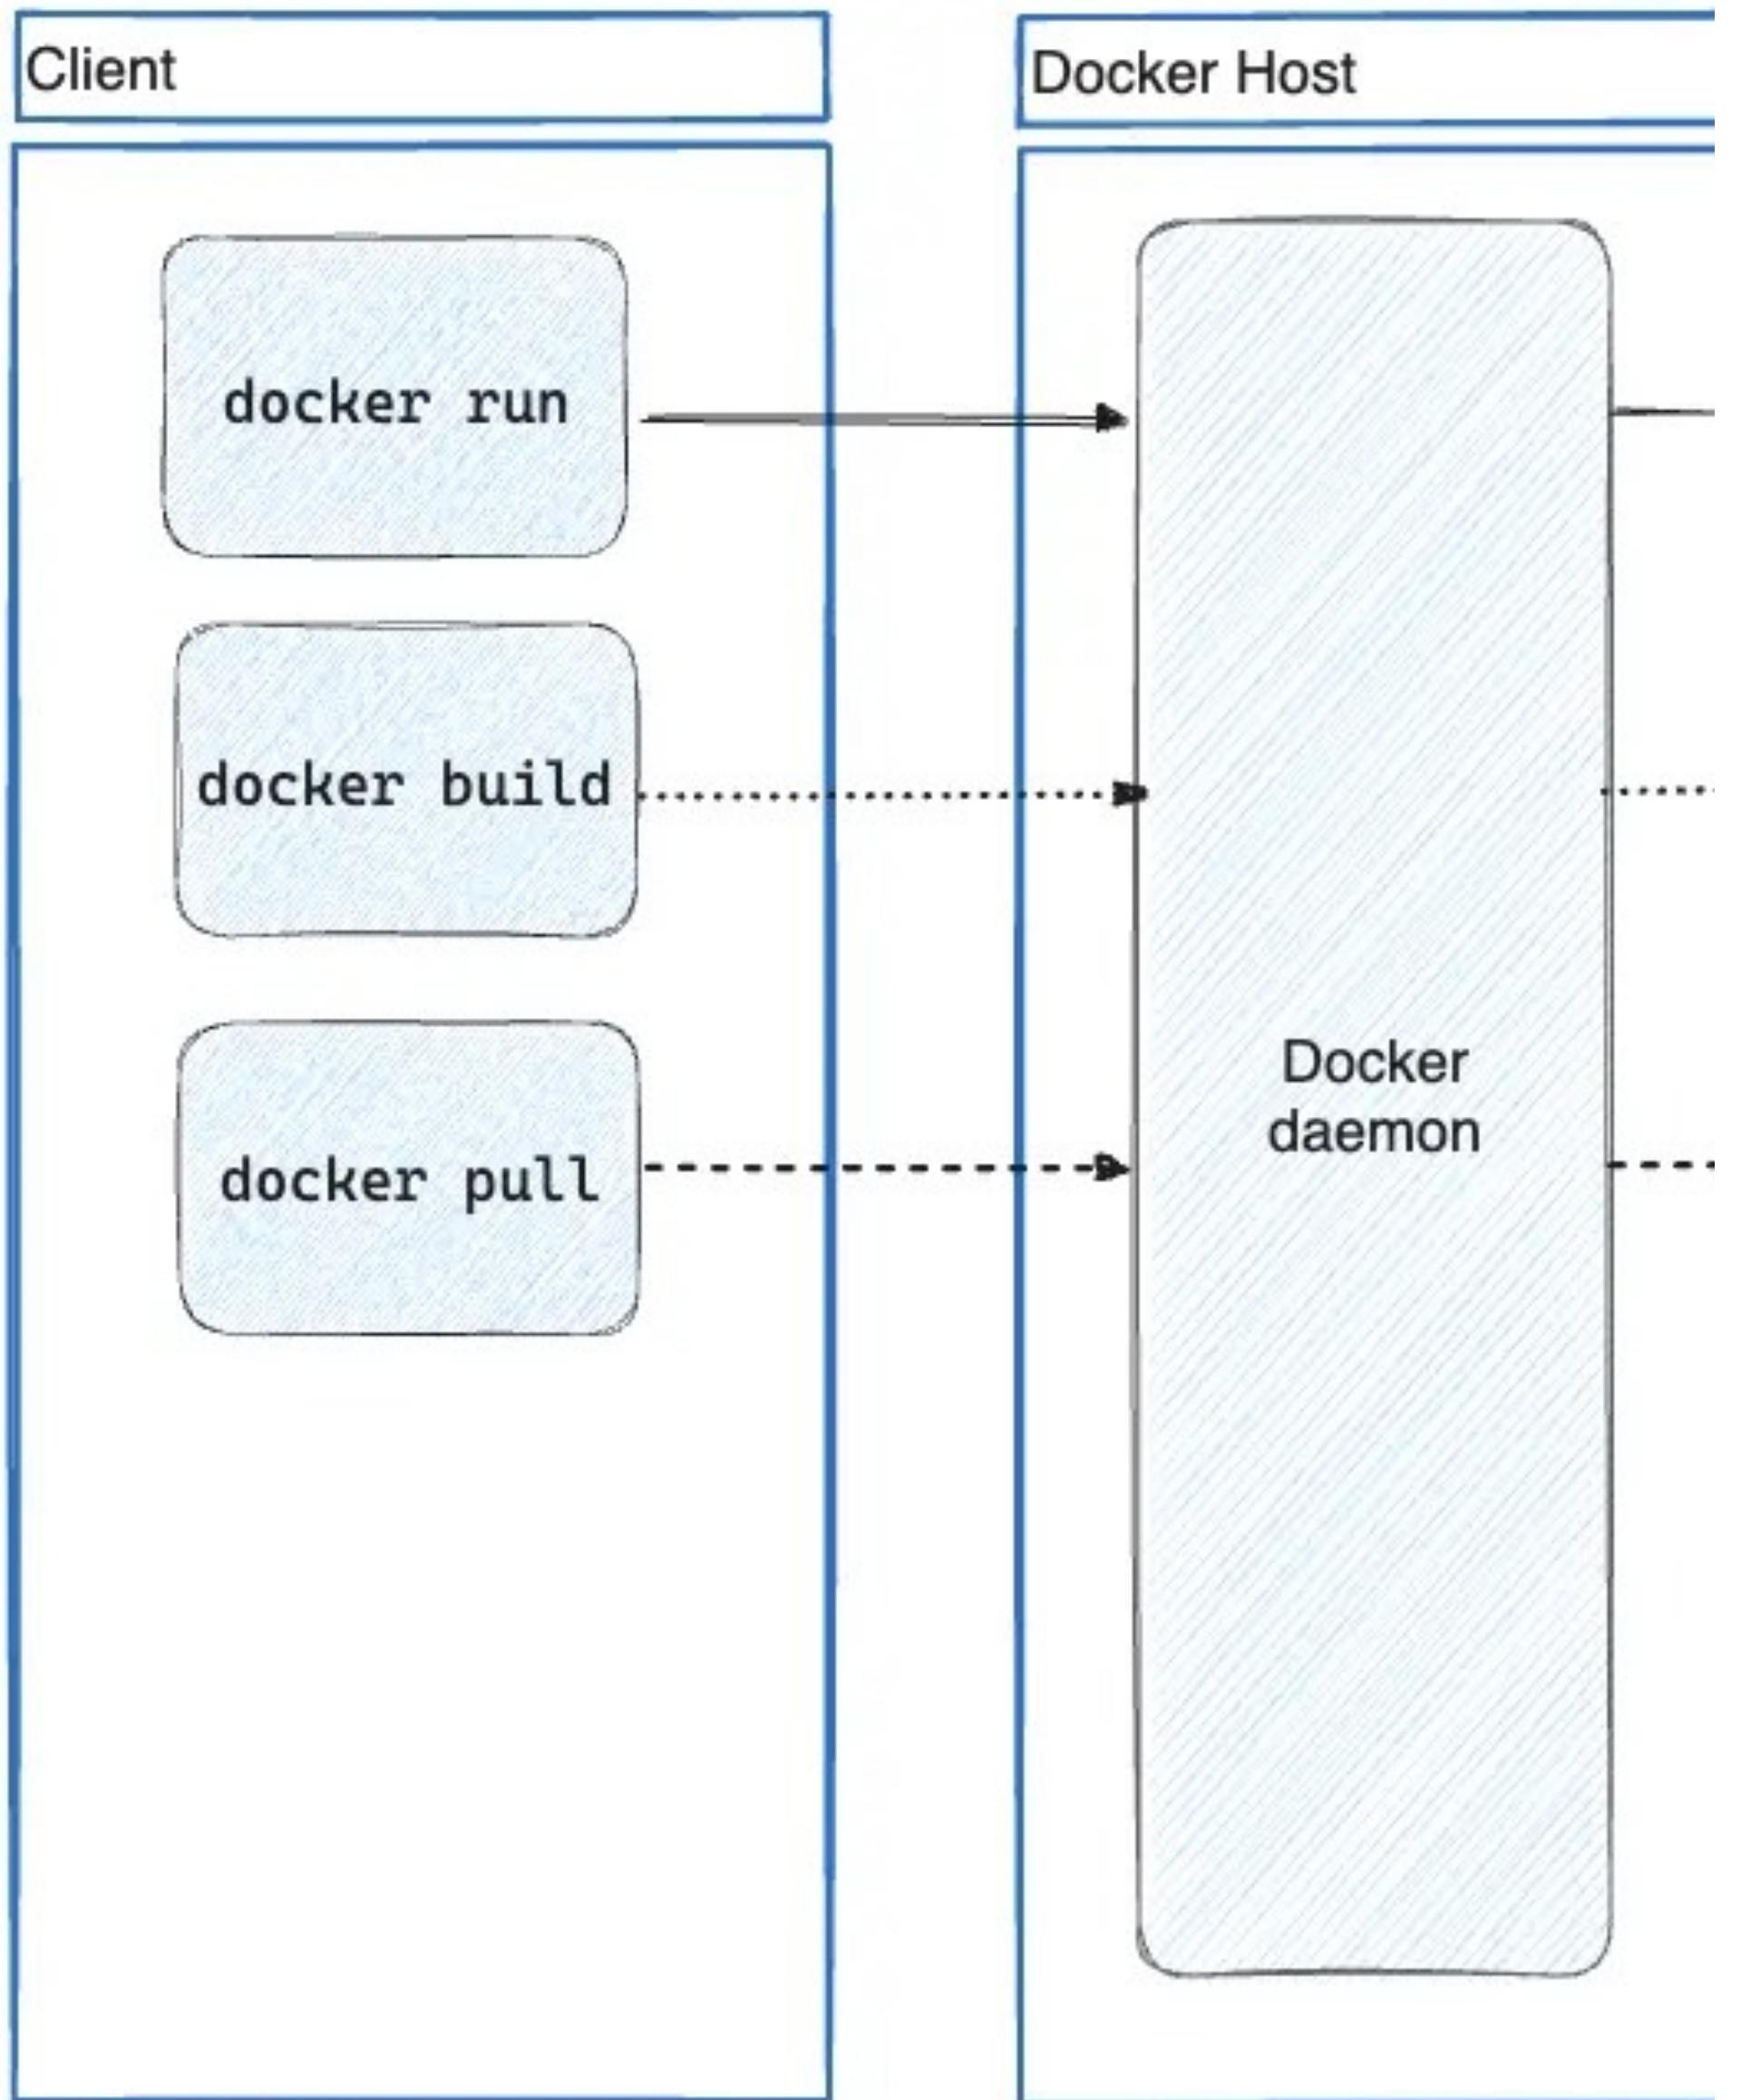
The Docker Daemon (Server)

- Runs and manages containers on a host system
 - Including the container runtime
- Client-Server Architecture
 - Clients sends commands to the daemon which handles container operations
- Supports building, running, and managing containers



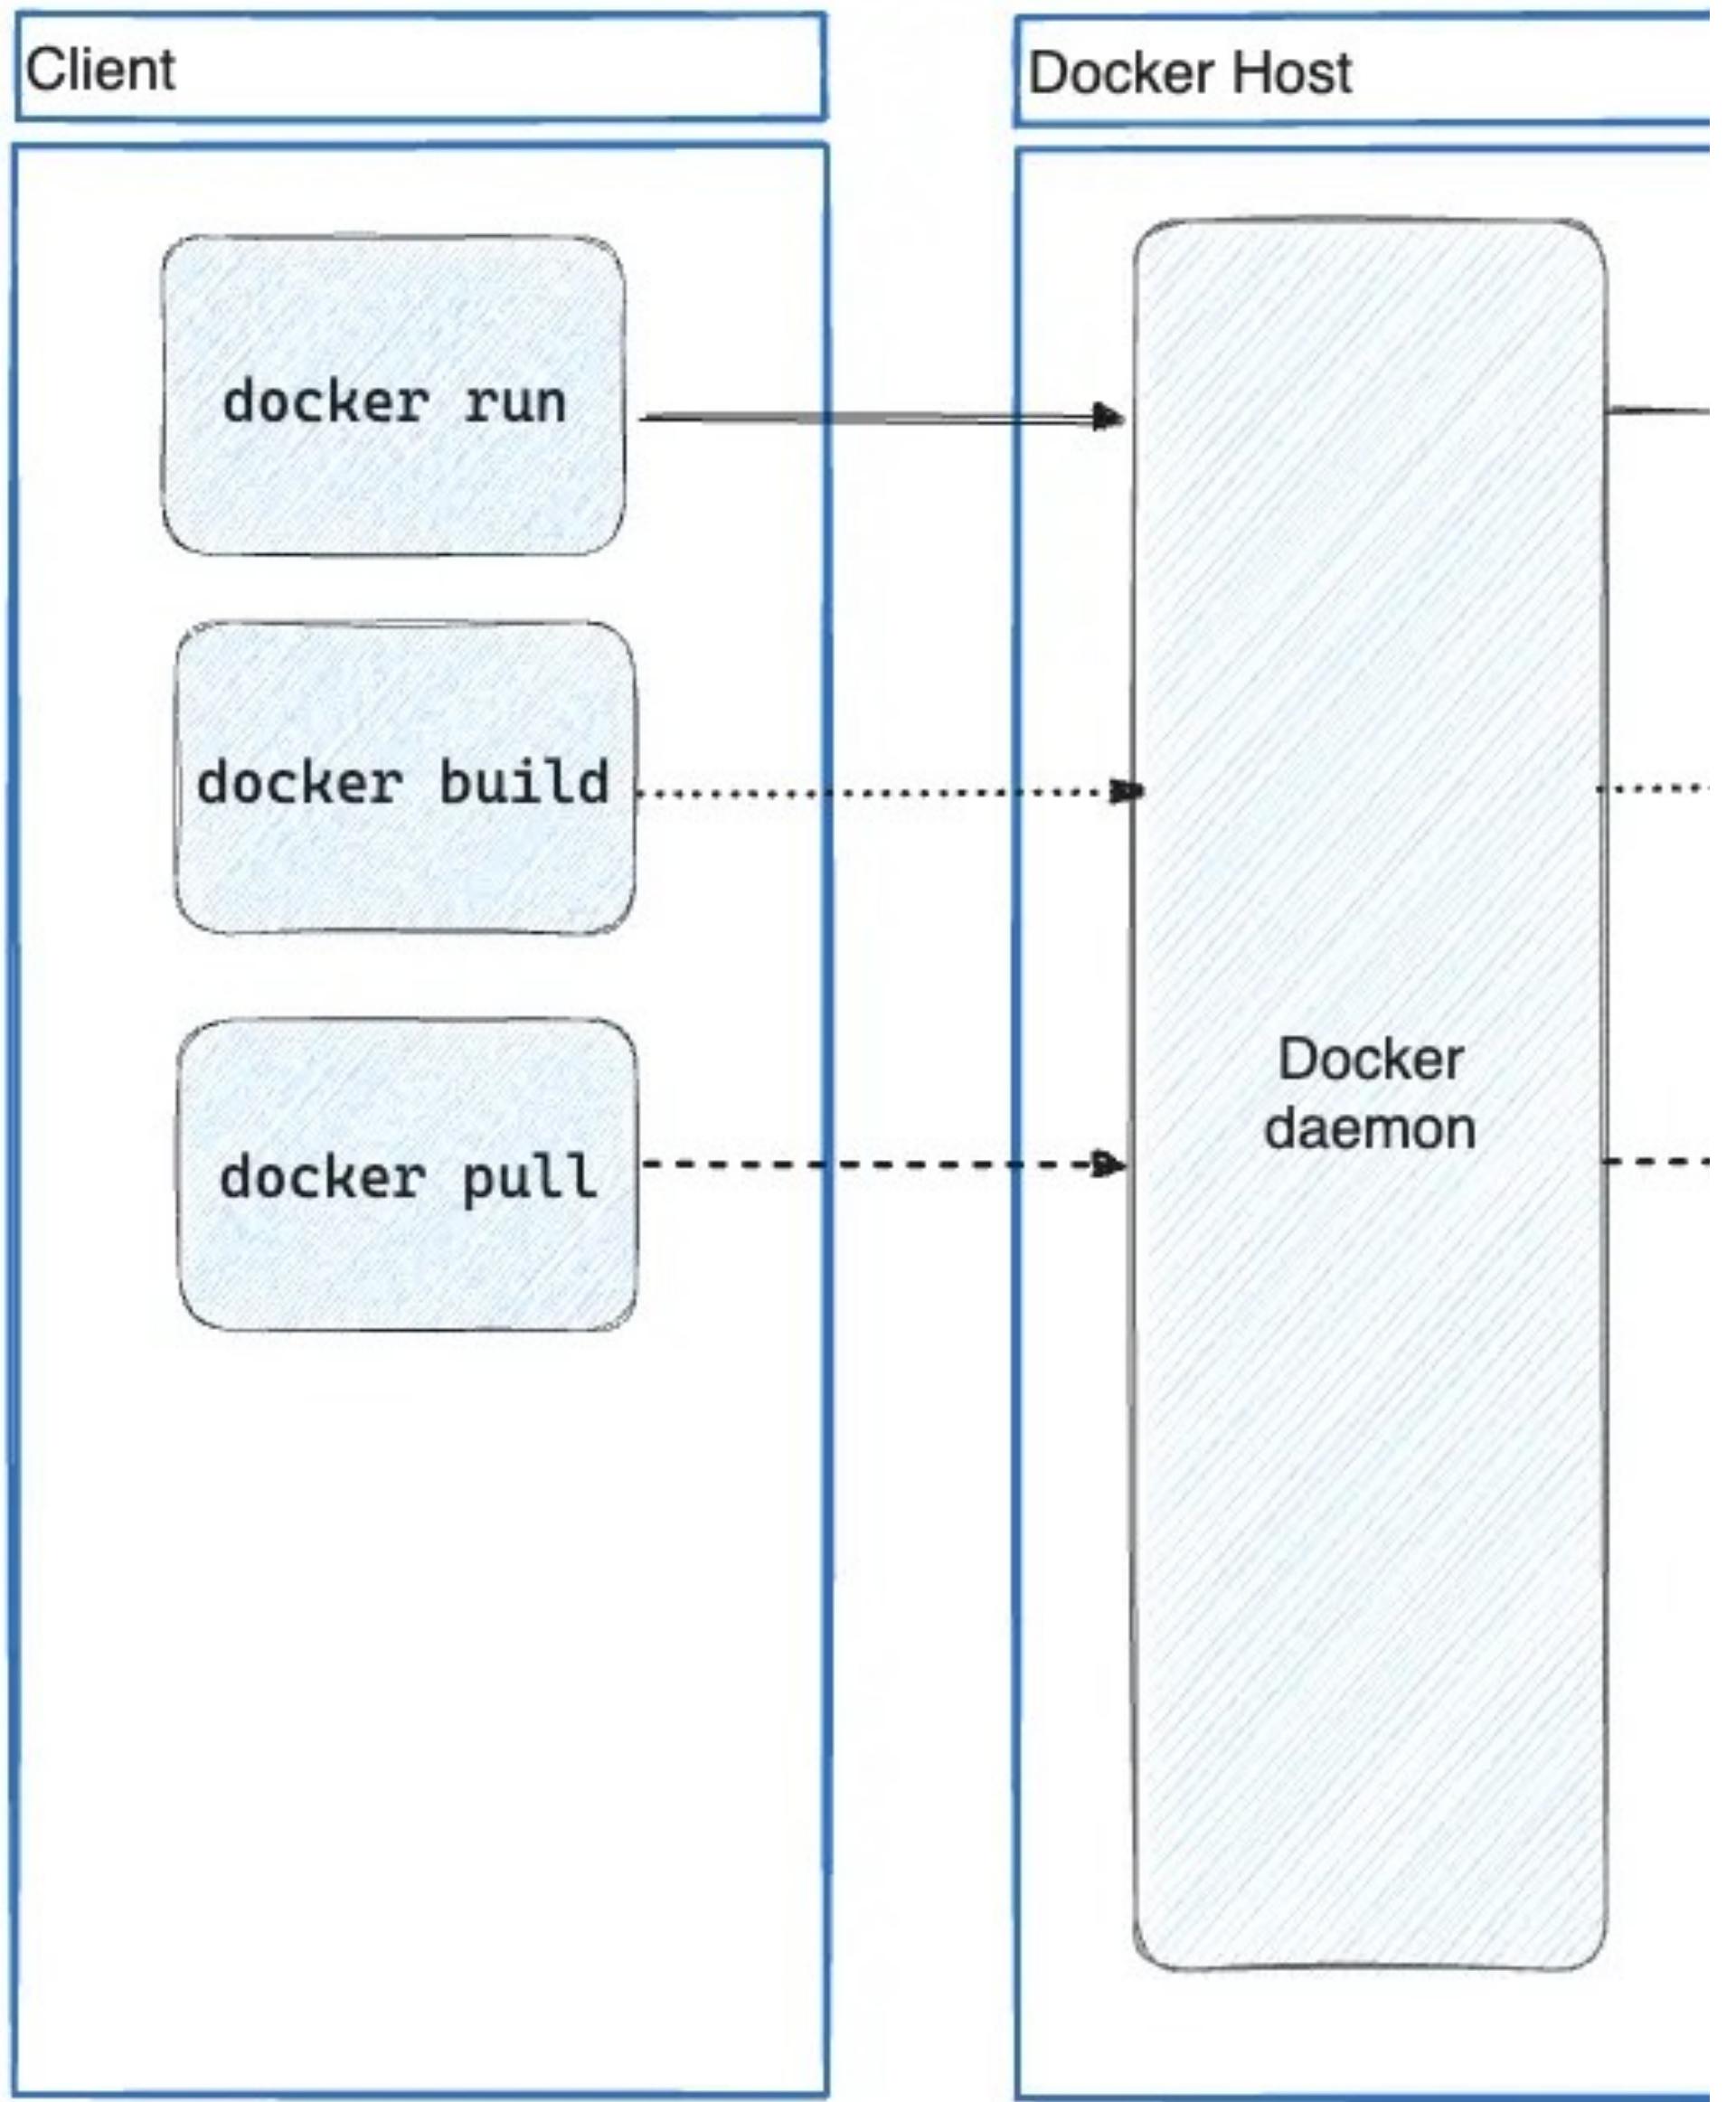
The Docker Command Line Client

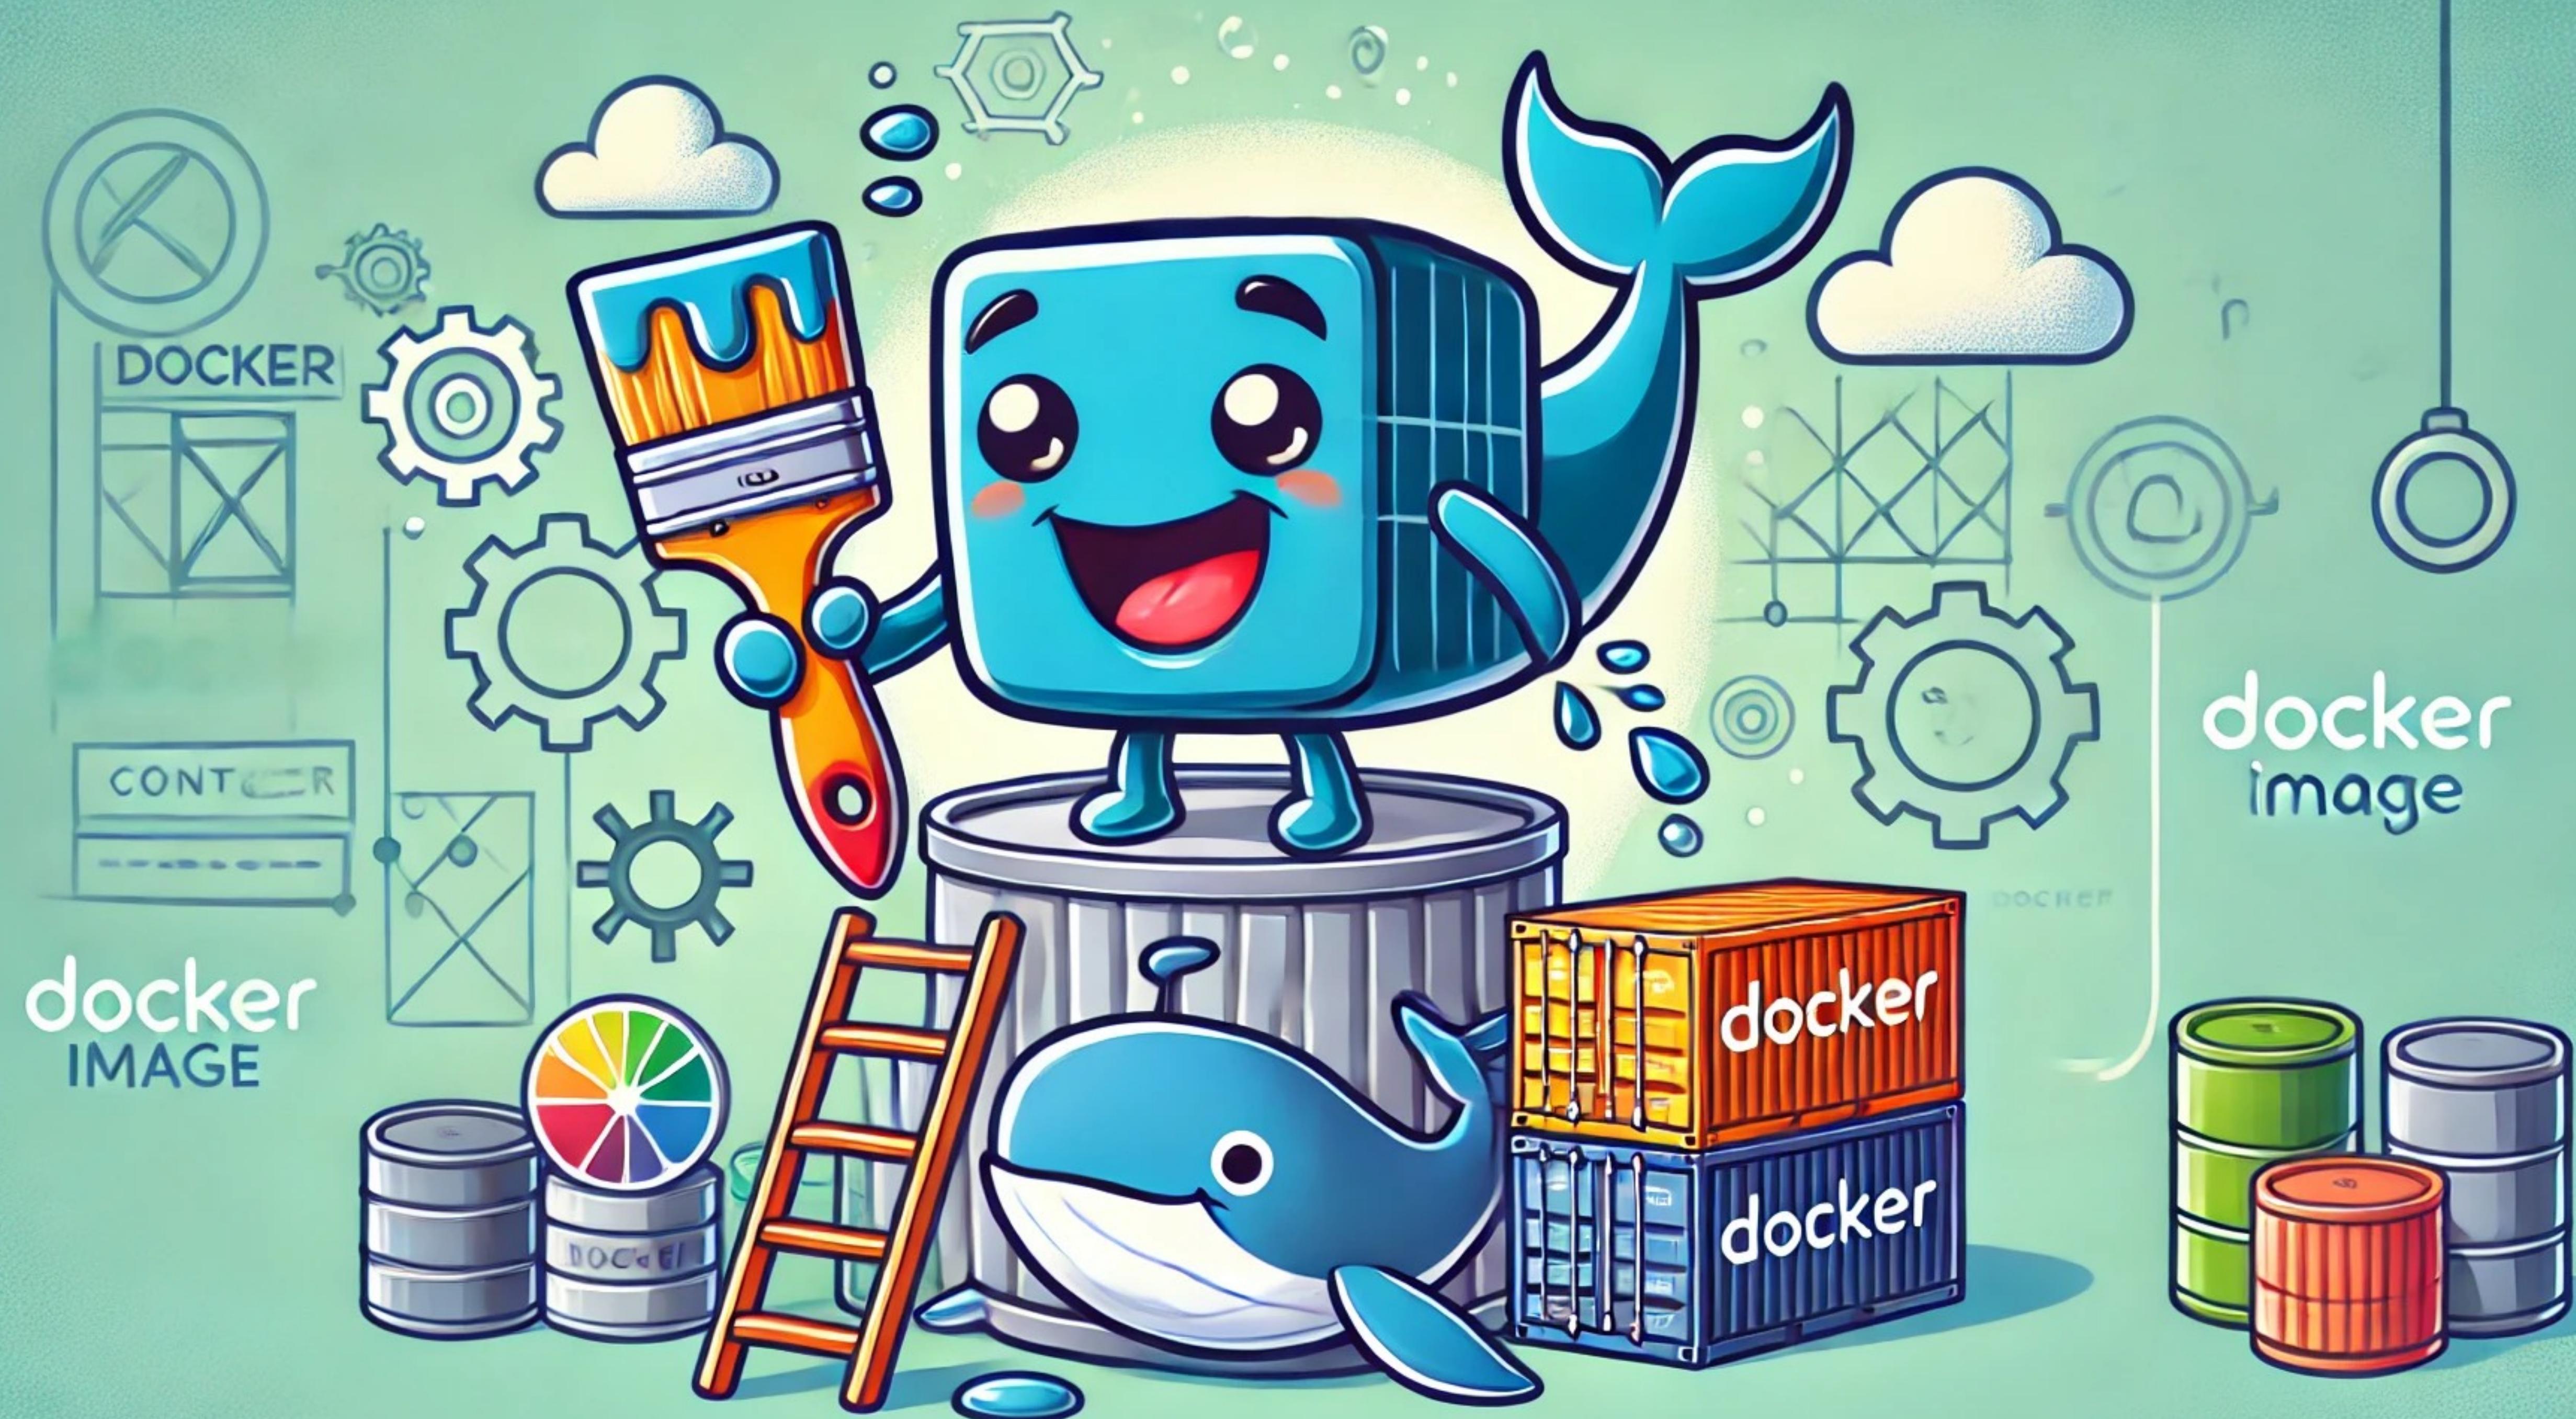
- Docker command uses the Docker APIs
 - Interact with the Daemon
- Docker is written in Go



REST-based API

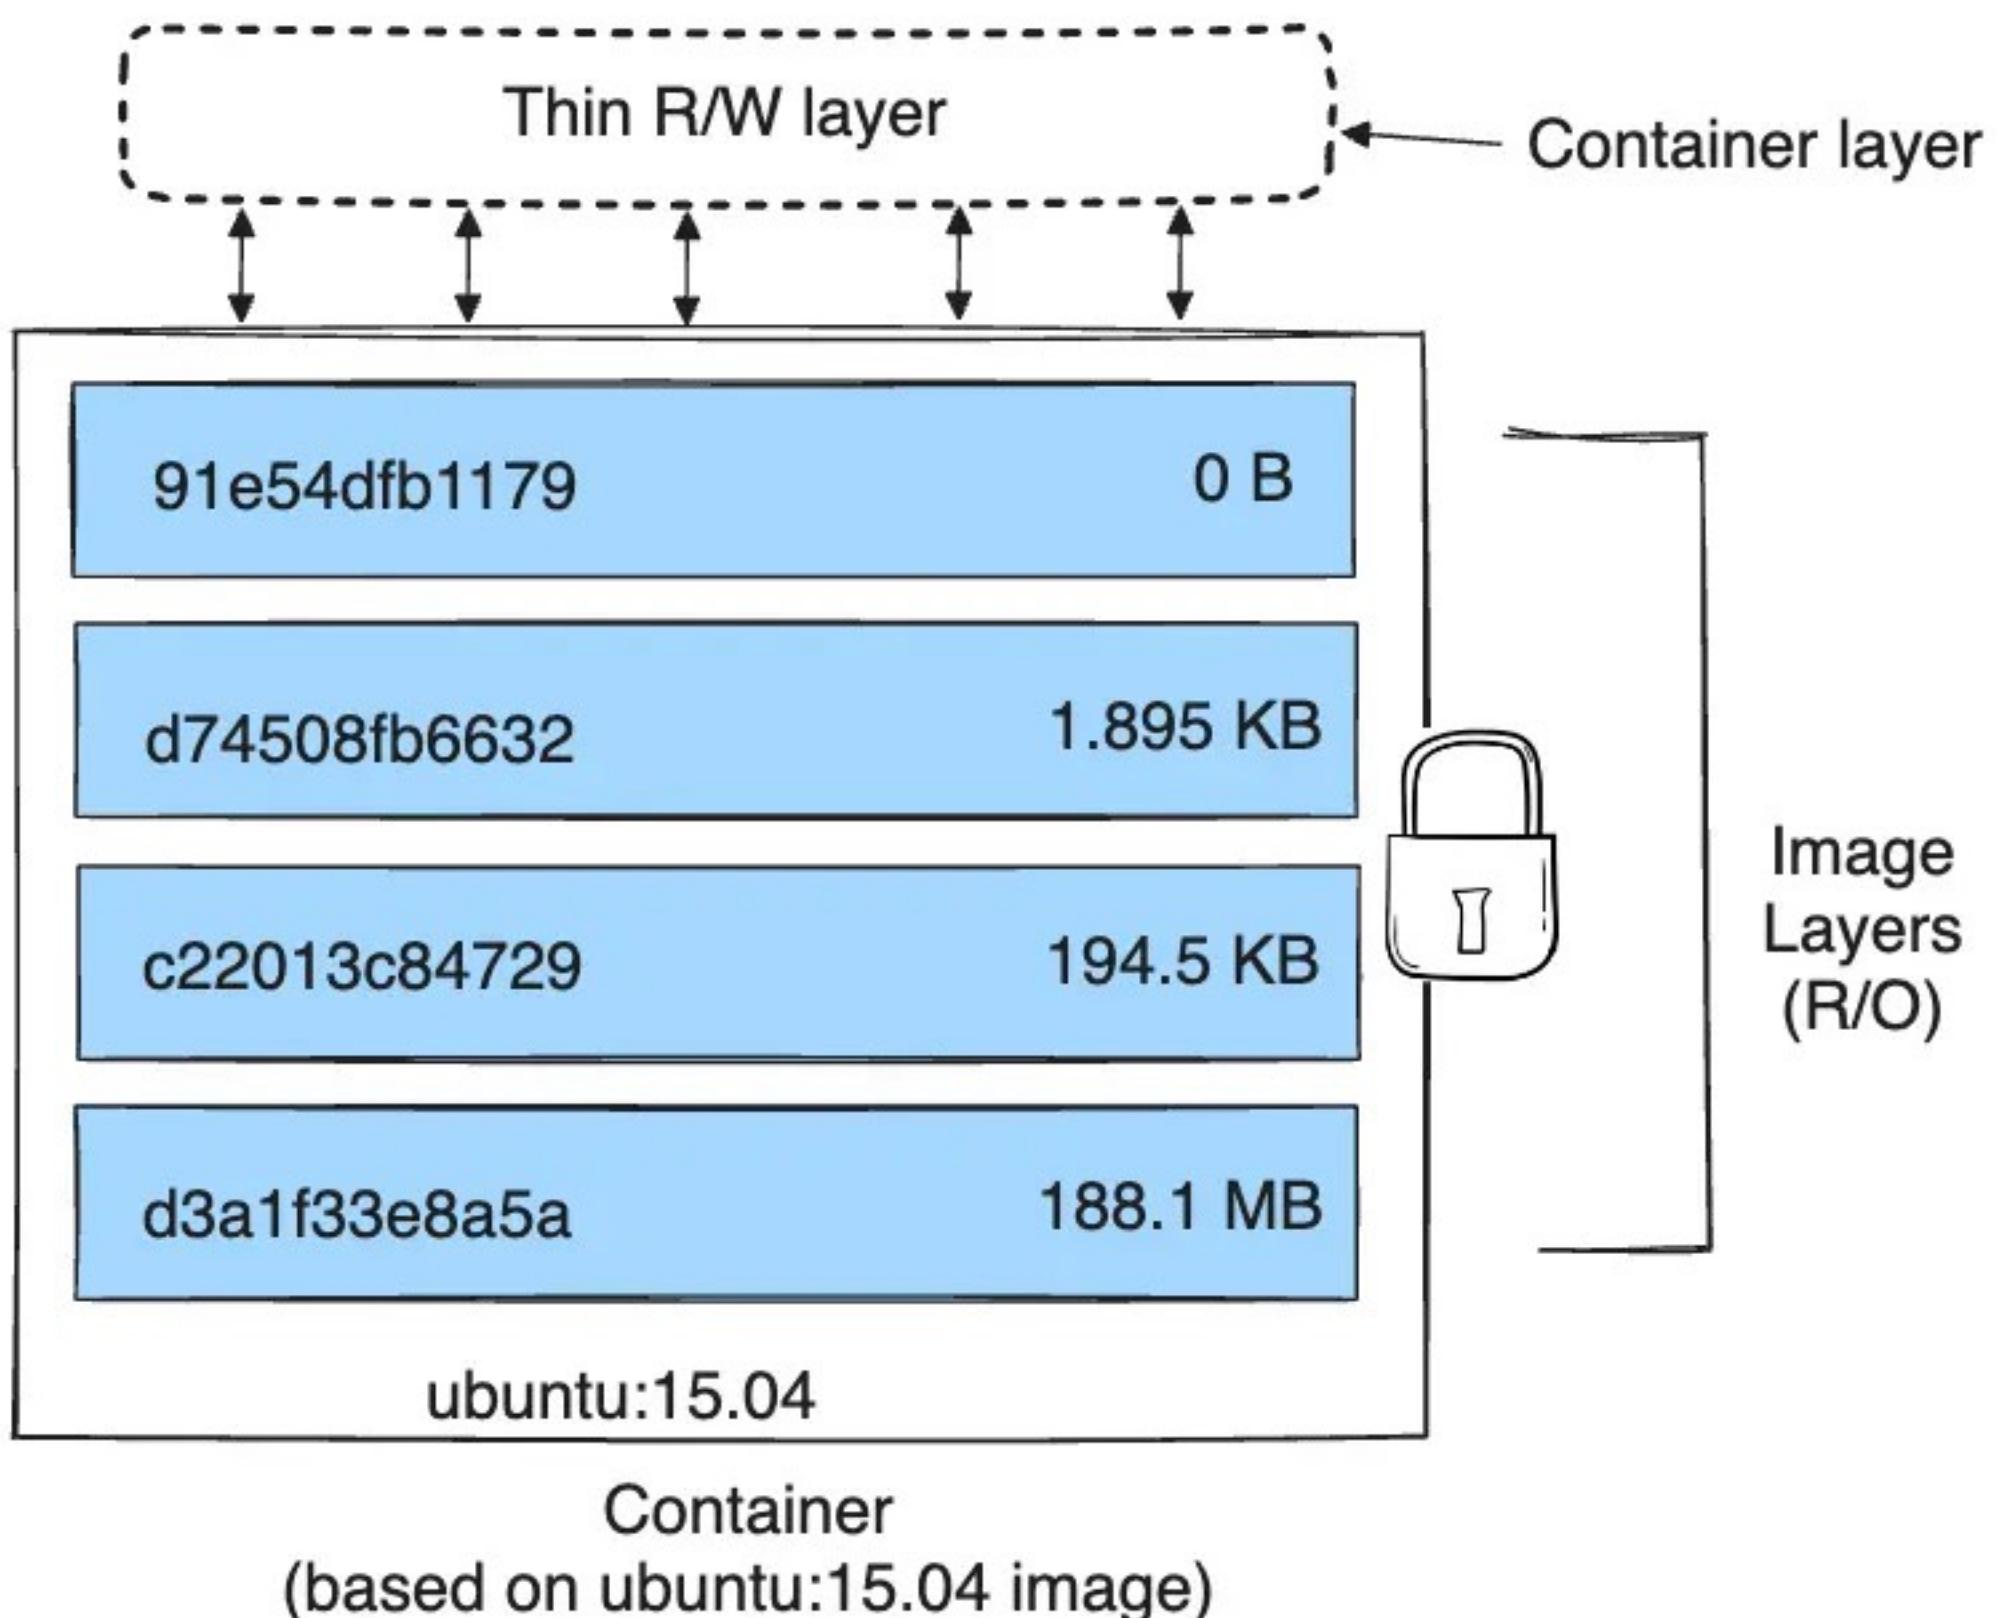
- Client libraries in most languages
 - Go SDK:
 - `go get github.com/docker/docker/client`
 - Python
 - `pip install docker`
 - Unofficial: C#, Java, Ruby, Node.js, ...





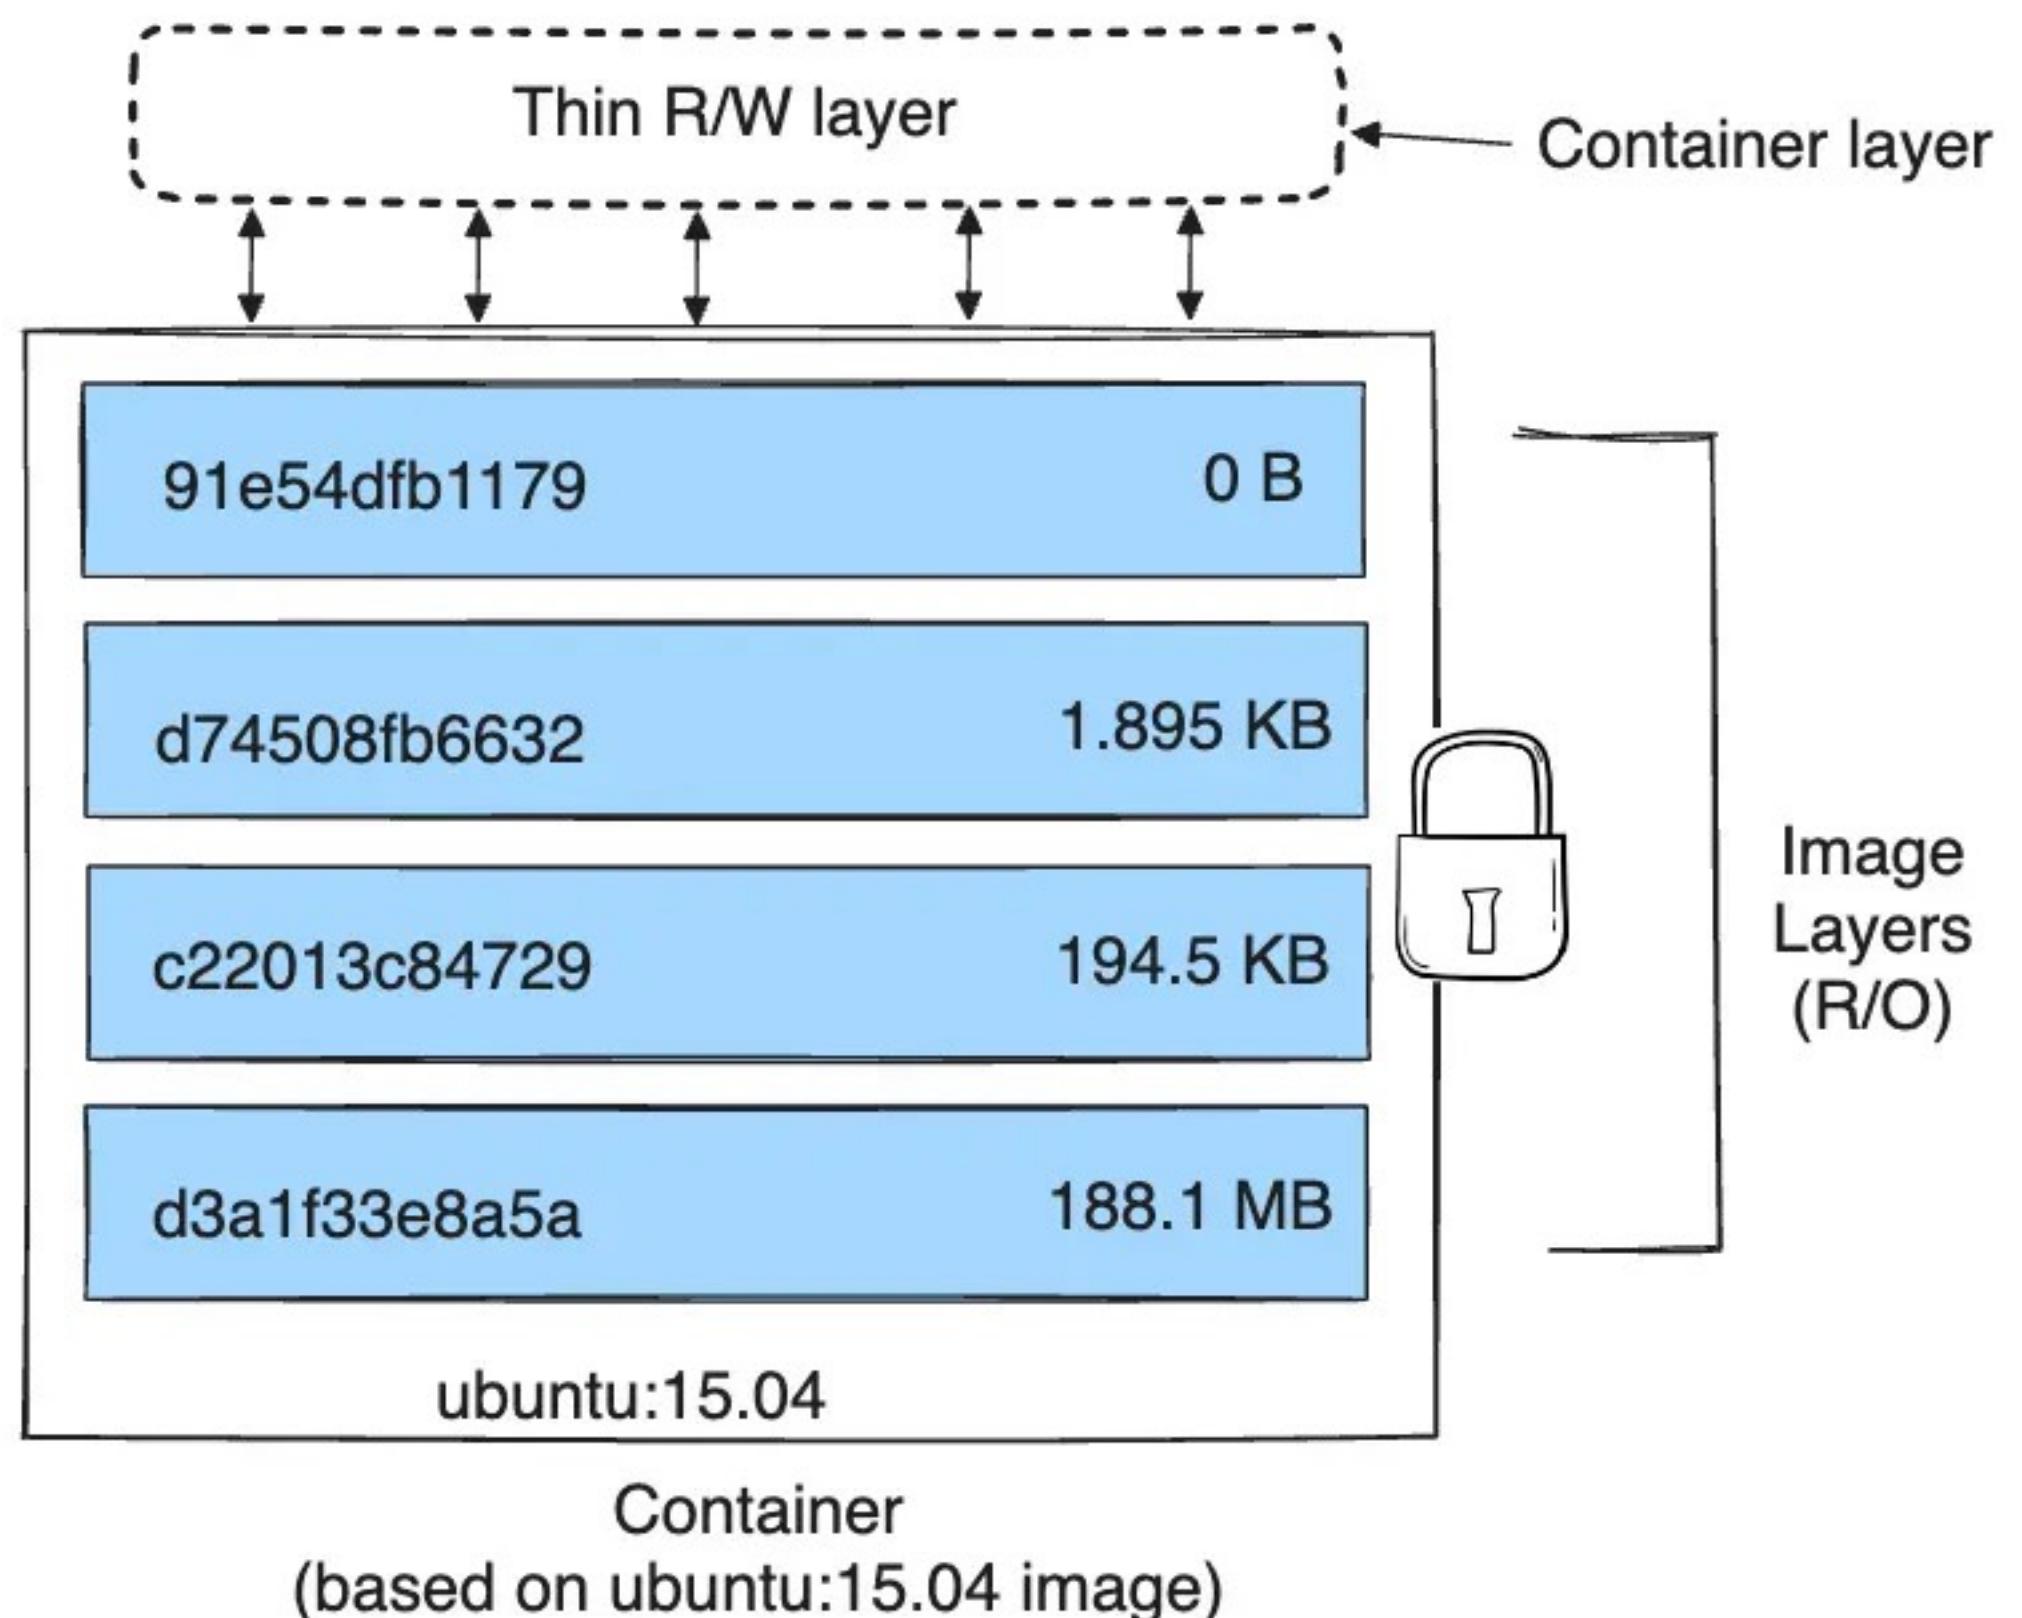
Immutable snapshot of an application and its dependencies

- Used to create containers
- Composed of layers
 - Each layer represents a change or addition to a base layer
- Enhancing efficiency through caching



Immutable snapshot of an application and its dependencies

- Reusable: Images can be shared and reused across environments and projects
 - Images are built from scratch or pulled from repositories
- Versioned: Tagged with versions
 - Ease management and rollback



Container
(based on ubuntu:15.04 image)

How to Identify a Docker Image?

Immutable snapshot of an application and its dependencies

- Format: **[repository]:[tag]**
- Repository Name
 - Examples: nginx, golang, python, myuser/myapp
- Tag
 - Default: latest
 - Examples: nginx:1.19, golang:1.23.0-alpine3.20

How to Identify a Docker Image?

Immutable snapshot of an application and its dependencies

```
docker pull golang
```

```
docker pull golang:1.20-alpine
```

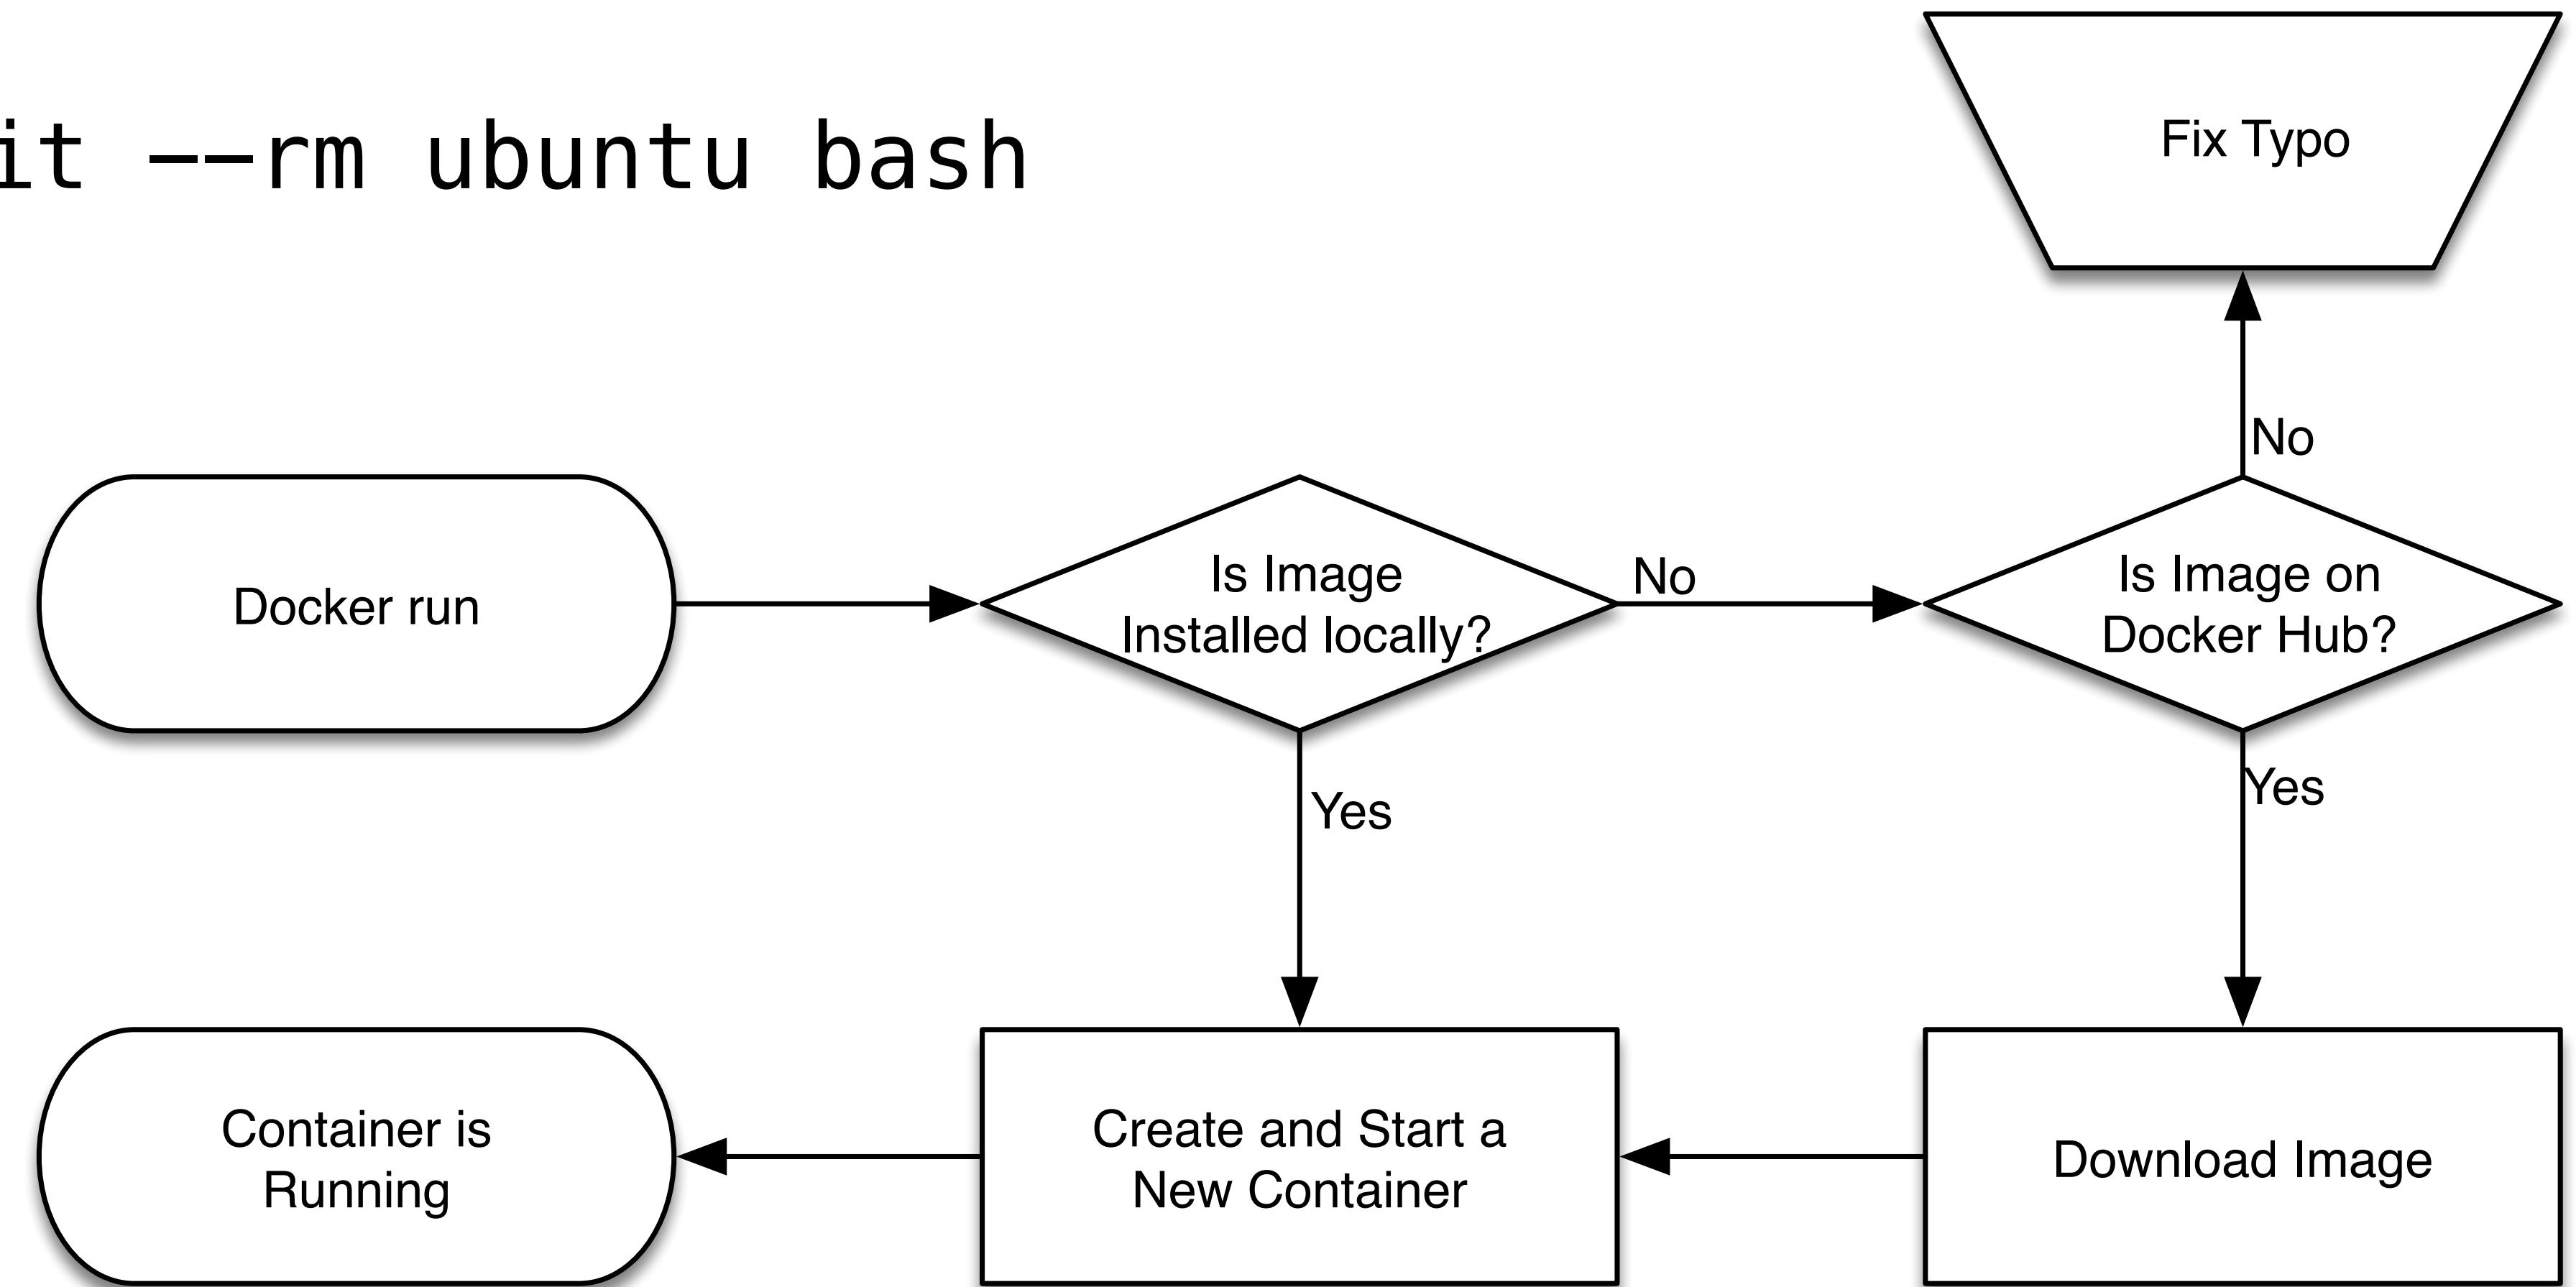
```
docker pull python
```

```
docker images
```

How to Identify a Docker Image?

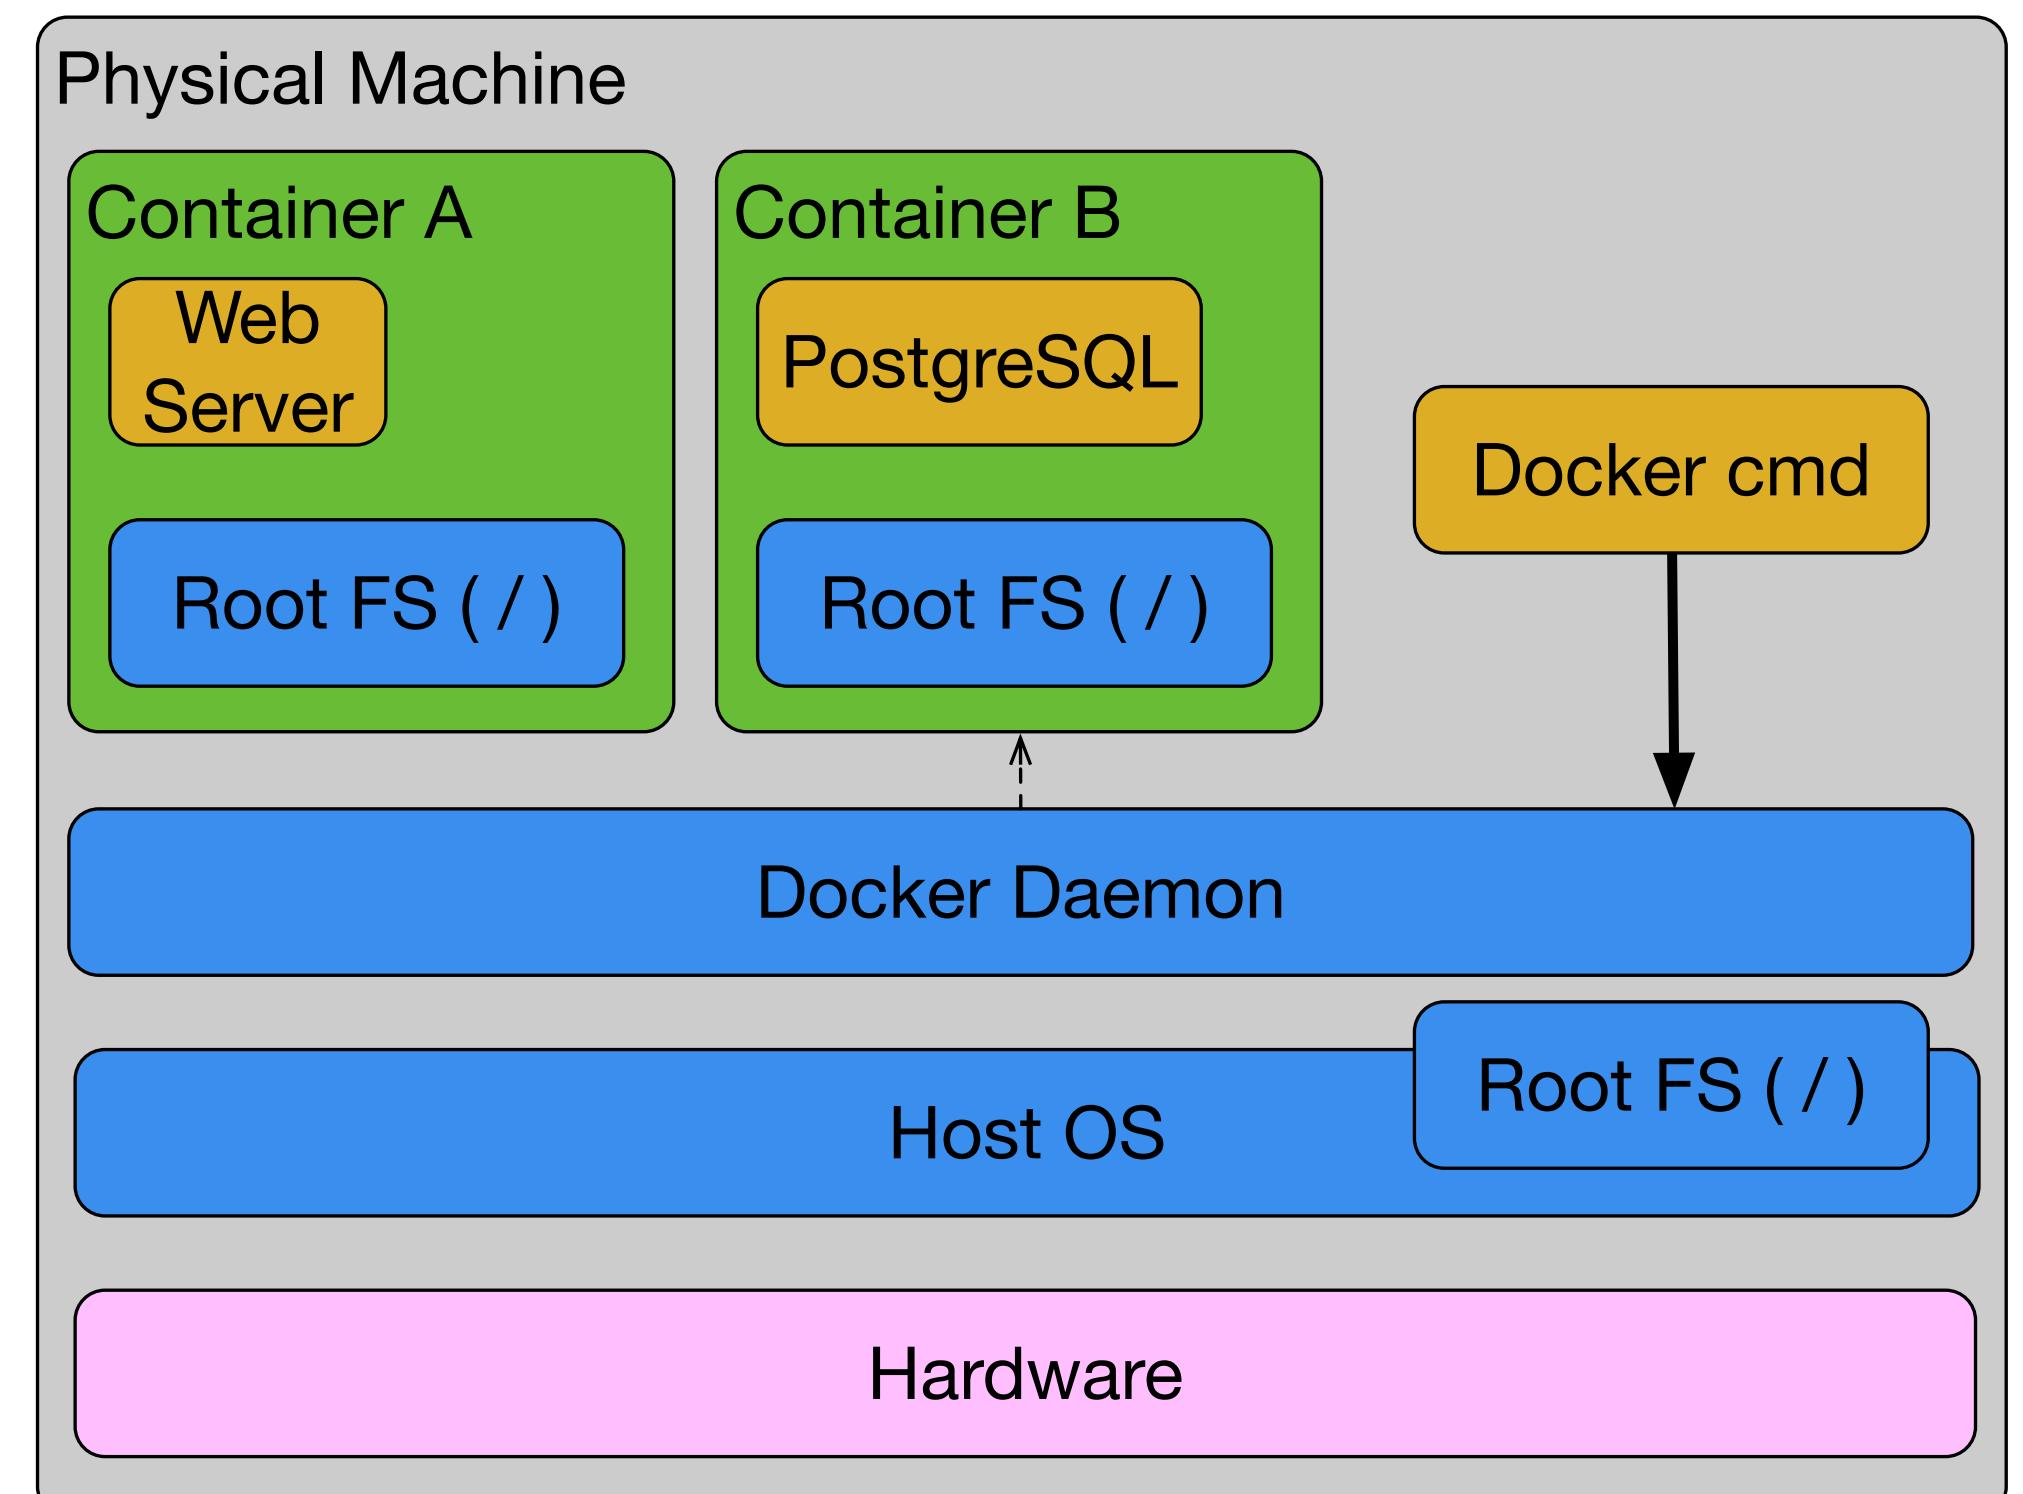
Immutable snapshot of an application and its dependencies

`docker run -it --rm ubuntu bash`



Running instance of an image

- Isolated environment for an application
- Lightweight and portable
 - Share host OS kernel
 - Minimize resource overhead
- Containers are live and ephemeral
 - Users can interact with them
 - Administrators can adjust their settings using Docker commands



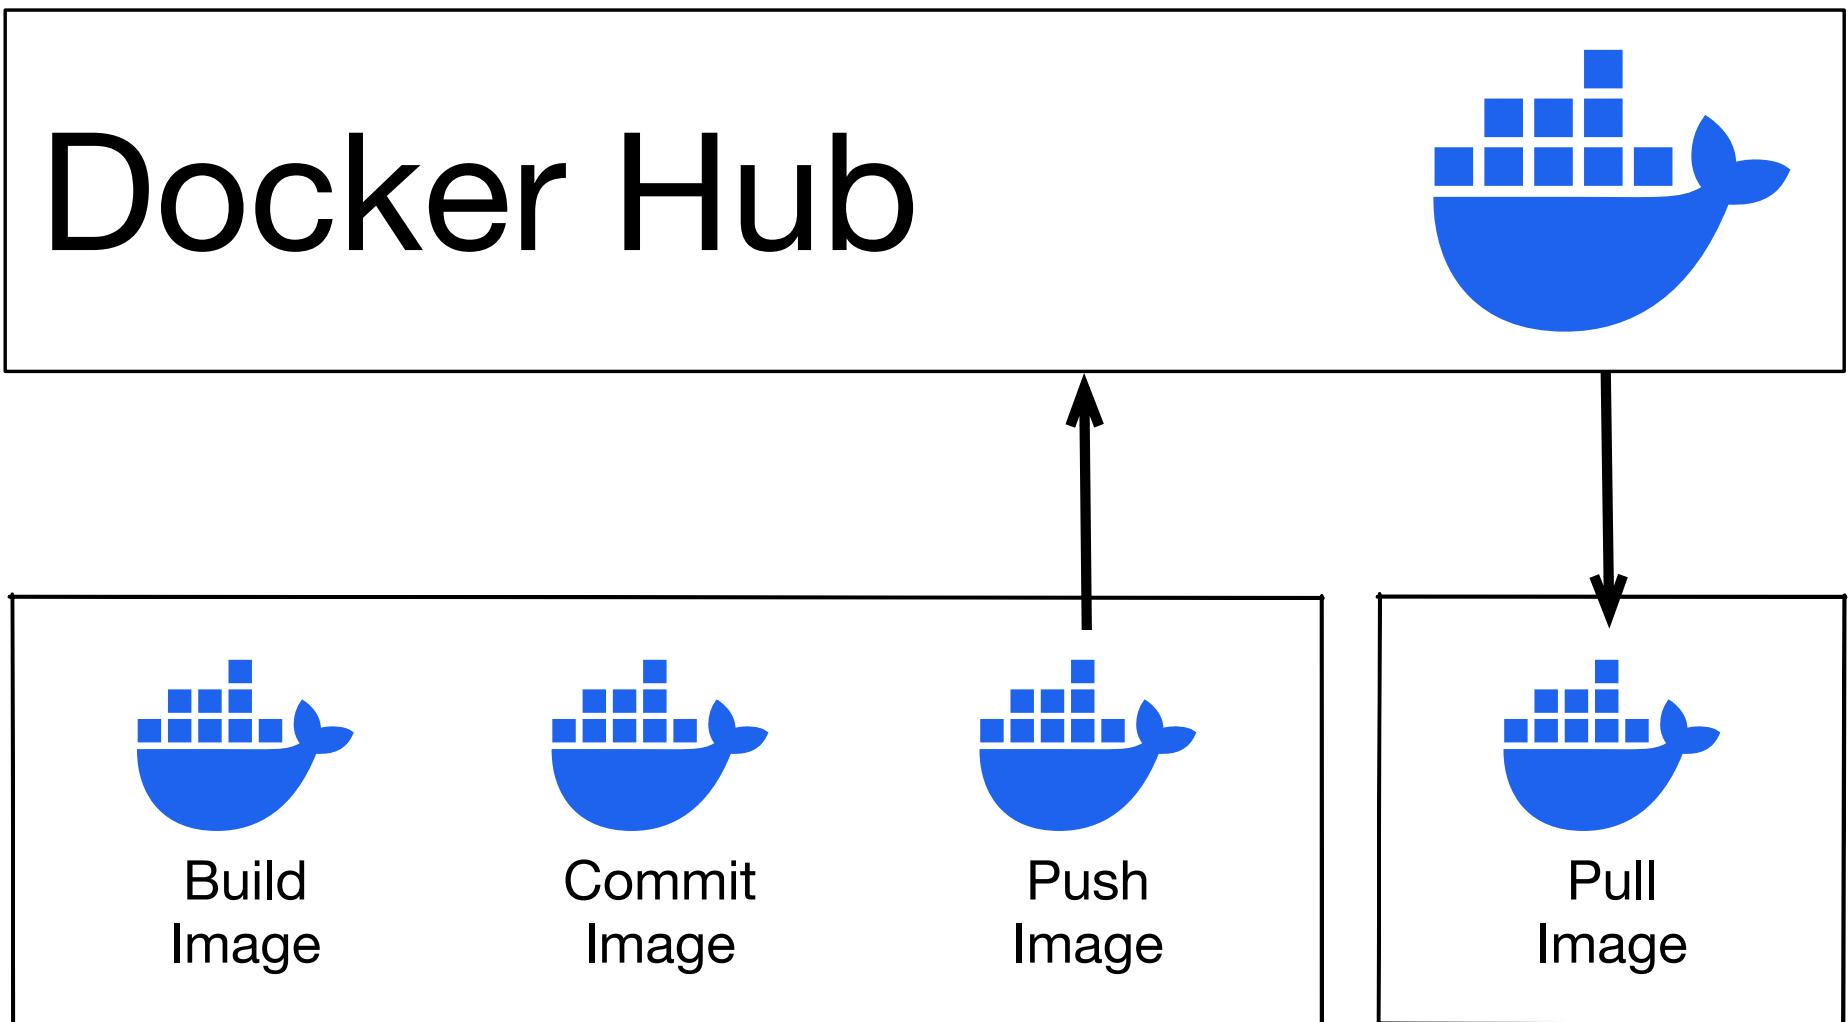
Registry, Repositories, and Images

- Registry is a centralized location that stores and manages container images
- Repository is a collection of related container images
 - Each repository contains one or more container images
 - Tagged for identification



Registry Service Provided by Docker

- Repositories for sharing container images with Docker community
- Official images
 - Clear doc, best practices, scanned for vulnerabilities...
- Publish images (external vendors)
- Private/Company images





The screenshot shows the Docker Hub search interface. The search bar at the top contains the query "Message Queues". Below the search bar, there is a "Suggested" dropdown menu. The main content area displays a list of 1 - 25 of 668 available results for "Message Queues". Each result card includes the repository name, a small icon, the number of pulls, and a pull chart. The repositories listed are rabbitmq, nats, eclipse-mosquitto, and emqx.

Filters (1) [Clear All](#)**Products**

- Images
- Extensions
- Plugins

Trusted Content

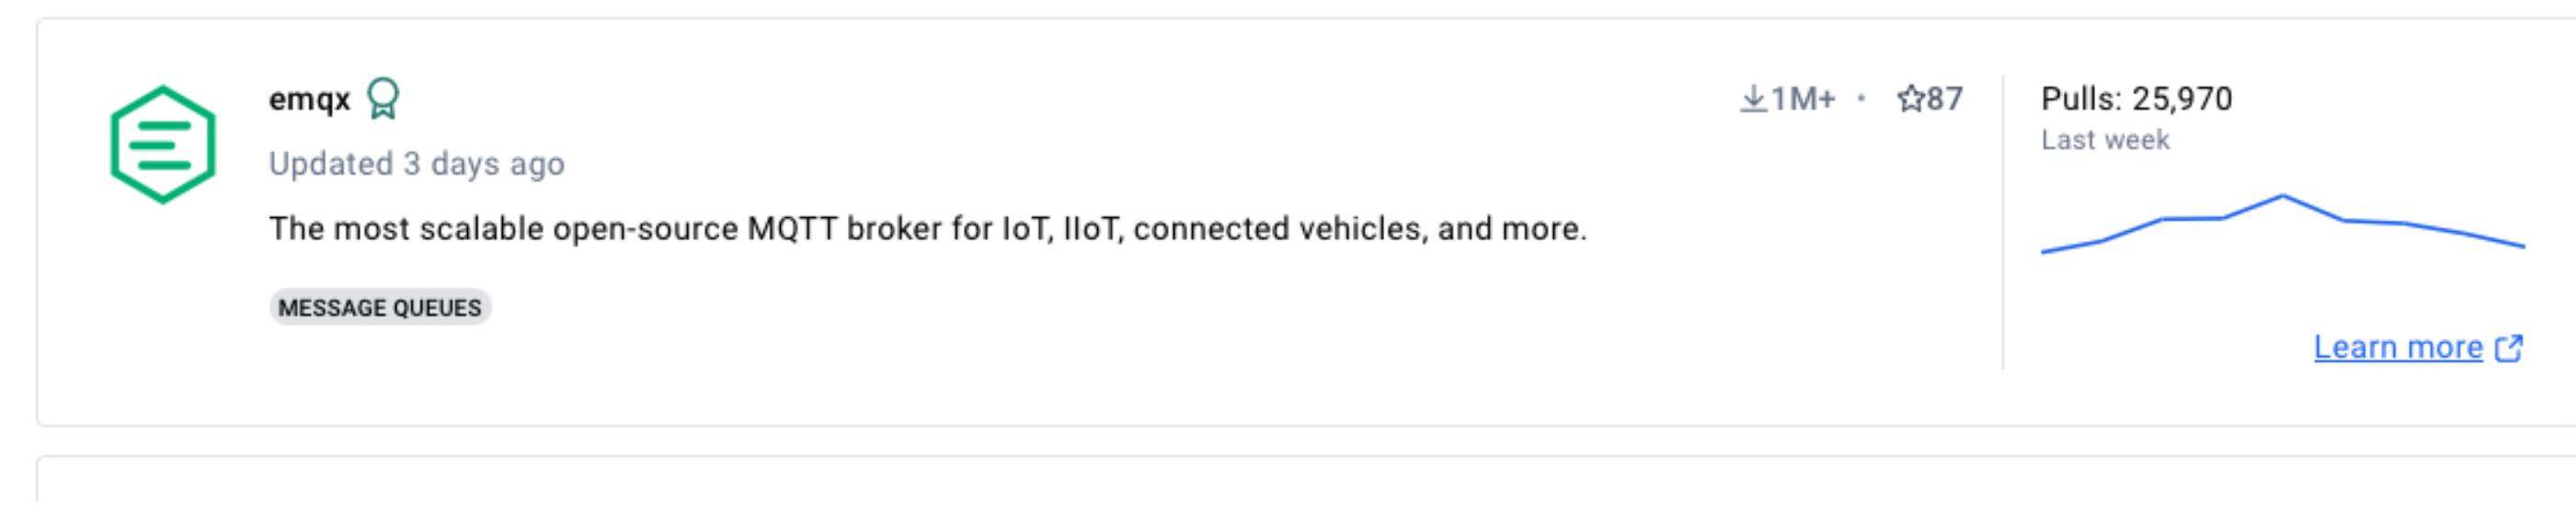
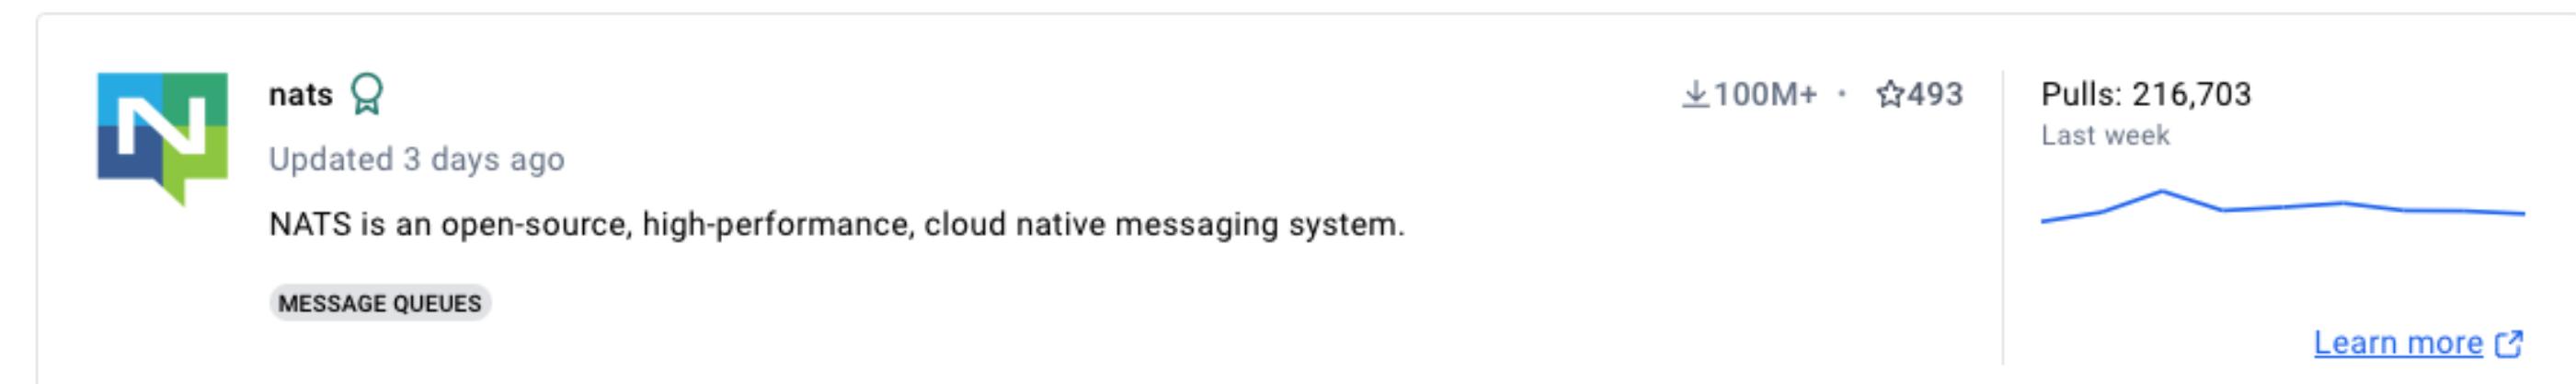
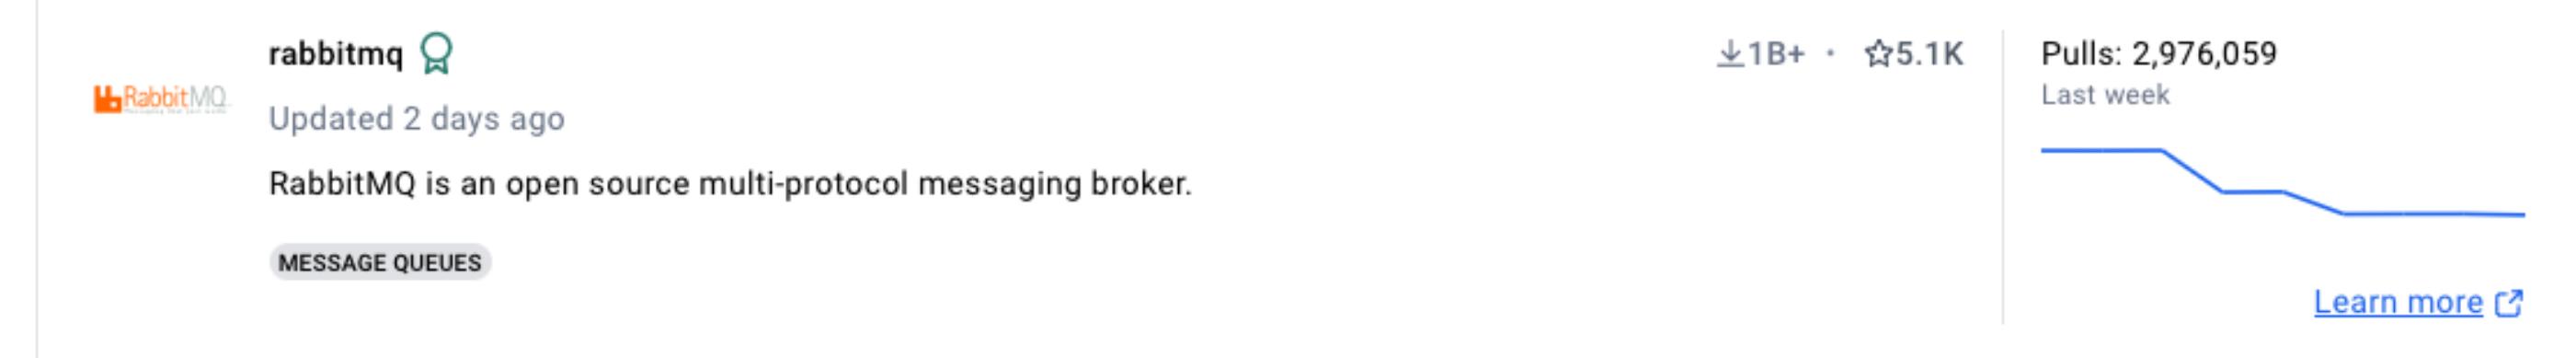
- Docker Official Image ⓘ
- Verified Publisher ⓘ
- Sponsored OSS ⓘ

Categories

- API Management
- Content Management System
- Data Science
- Databases & Storage
- Developer Tools
- Integration & Delivery
- Internet Of Things
- Languages & Frameworks
- Machine Learning & AI
- Message Queues
- Monitoring & Observability

1 - 25 of 668 available results.

Message Queues ×

Suggested ▾

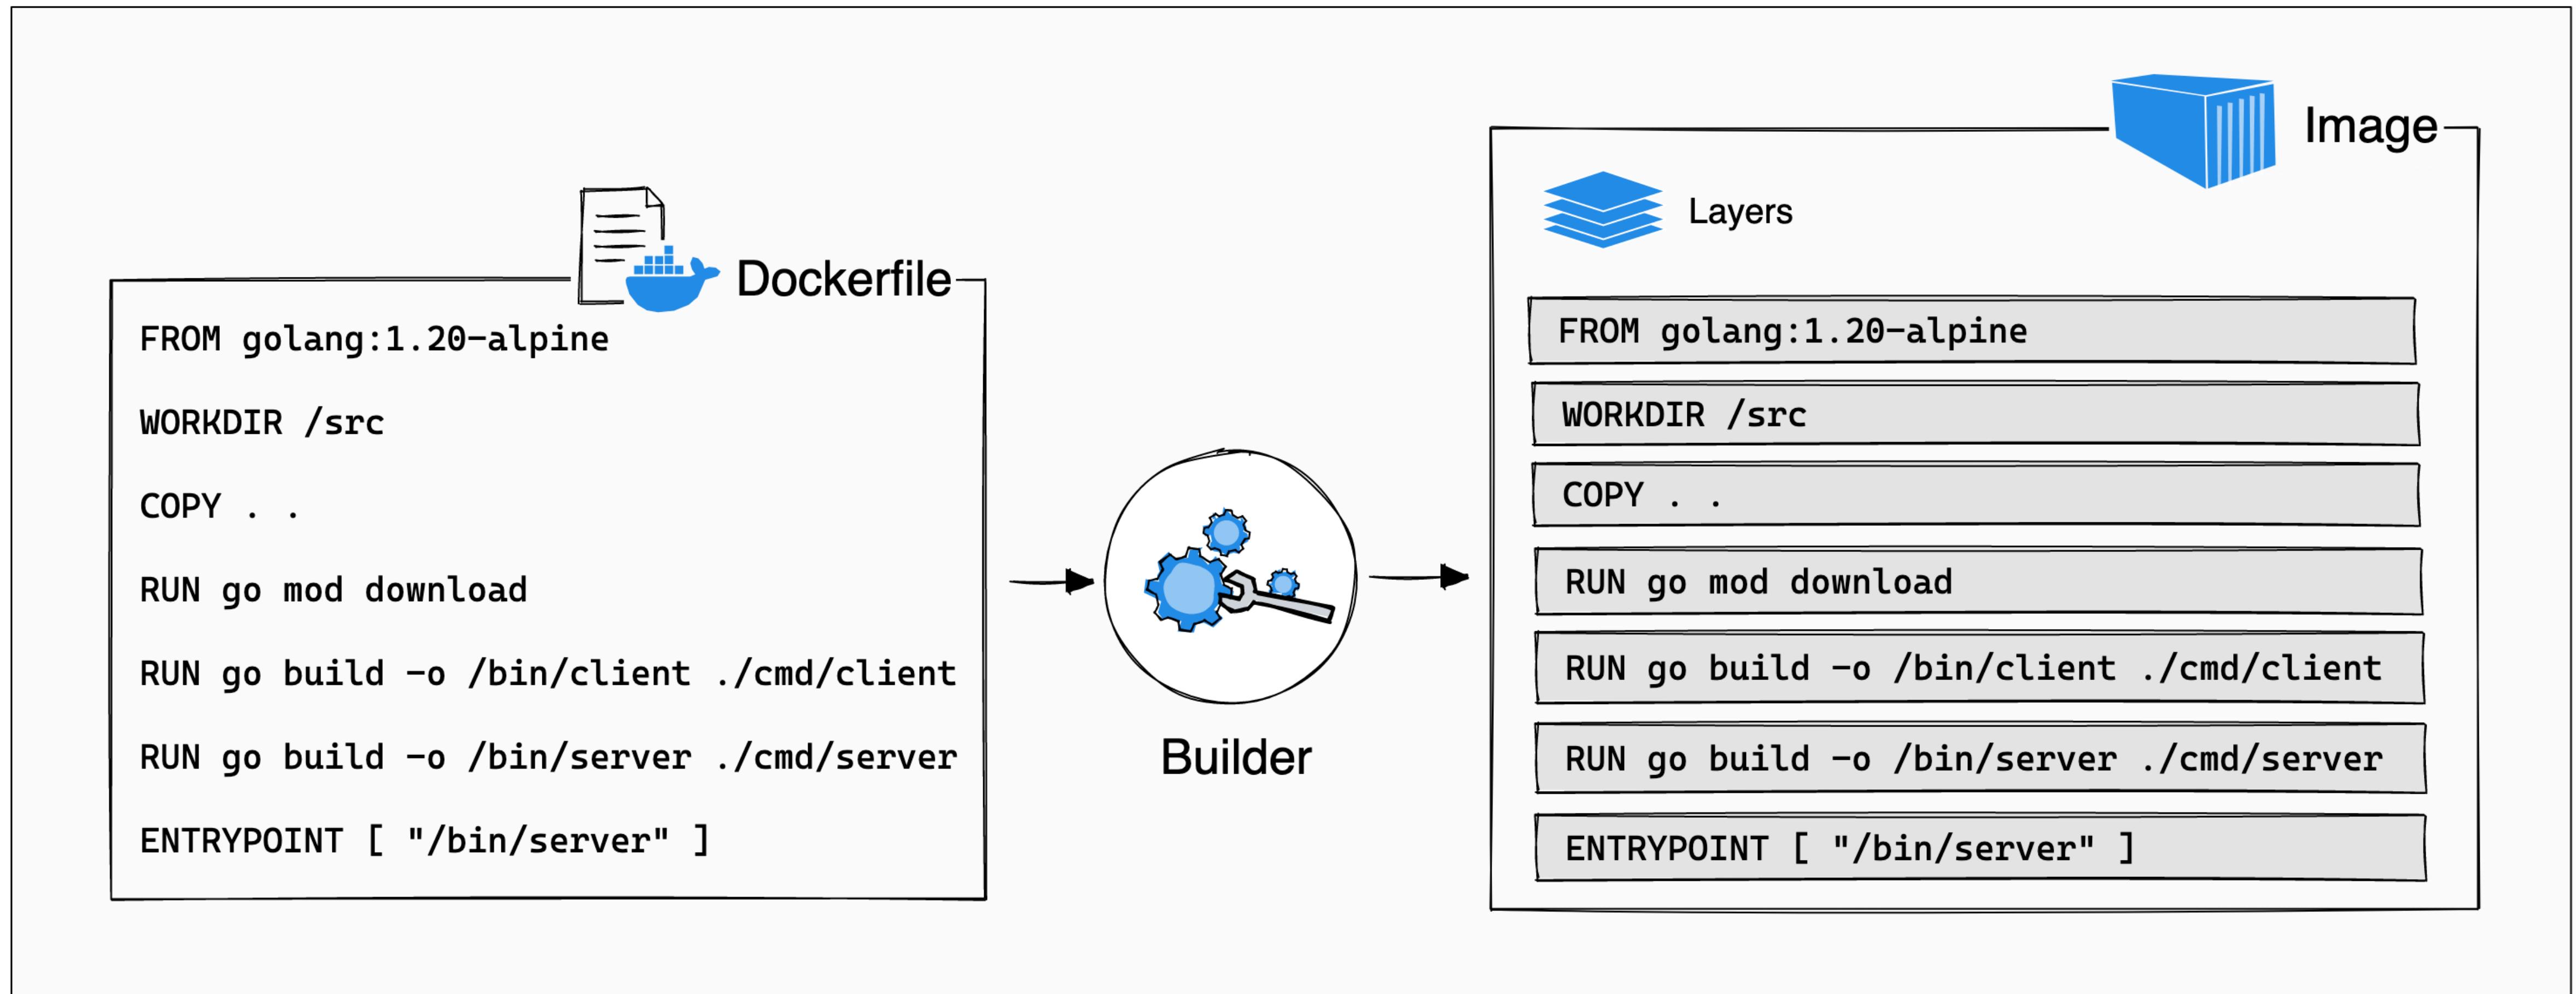
The Dockerfile

Text File for Image Construction

- Docker builds images based on Dockerfile
- Standard set of commands

Instruction	Description
ADD	Add local or remote files and directories.
ARG	Use build-time variables.
CMD	Specify default commands.
COPY	Copy files and directories.
ENTRYPOINT	Specify default executable.
ENV	Set environment variables.
EXPOSE	Describe which ports your application is listening on.
FROM	Create a new build stage from a base image.
HEALTHCHECK	Check a container's health on startup.
LABEL	Add metadata to an image.
MAINTAINER	Specify the author of an image.
ONBUILD	Specify instructions for when the image is used in a build.
RUN	Execute build commands.
SHELL	Set the default shell of an image.
STOPSIG	Specify the system call signal for exiting a container.

Text File for Image Construction



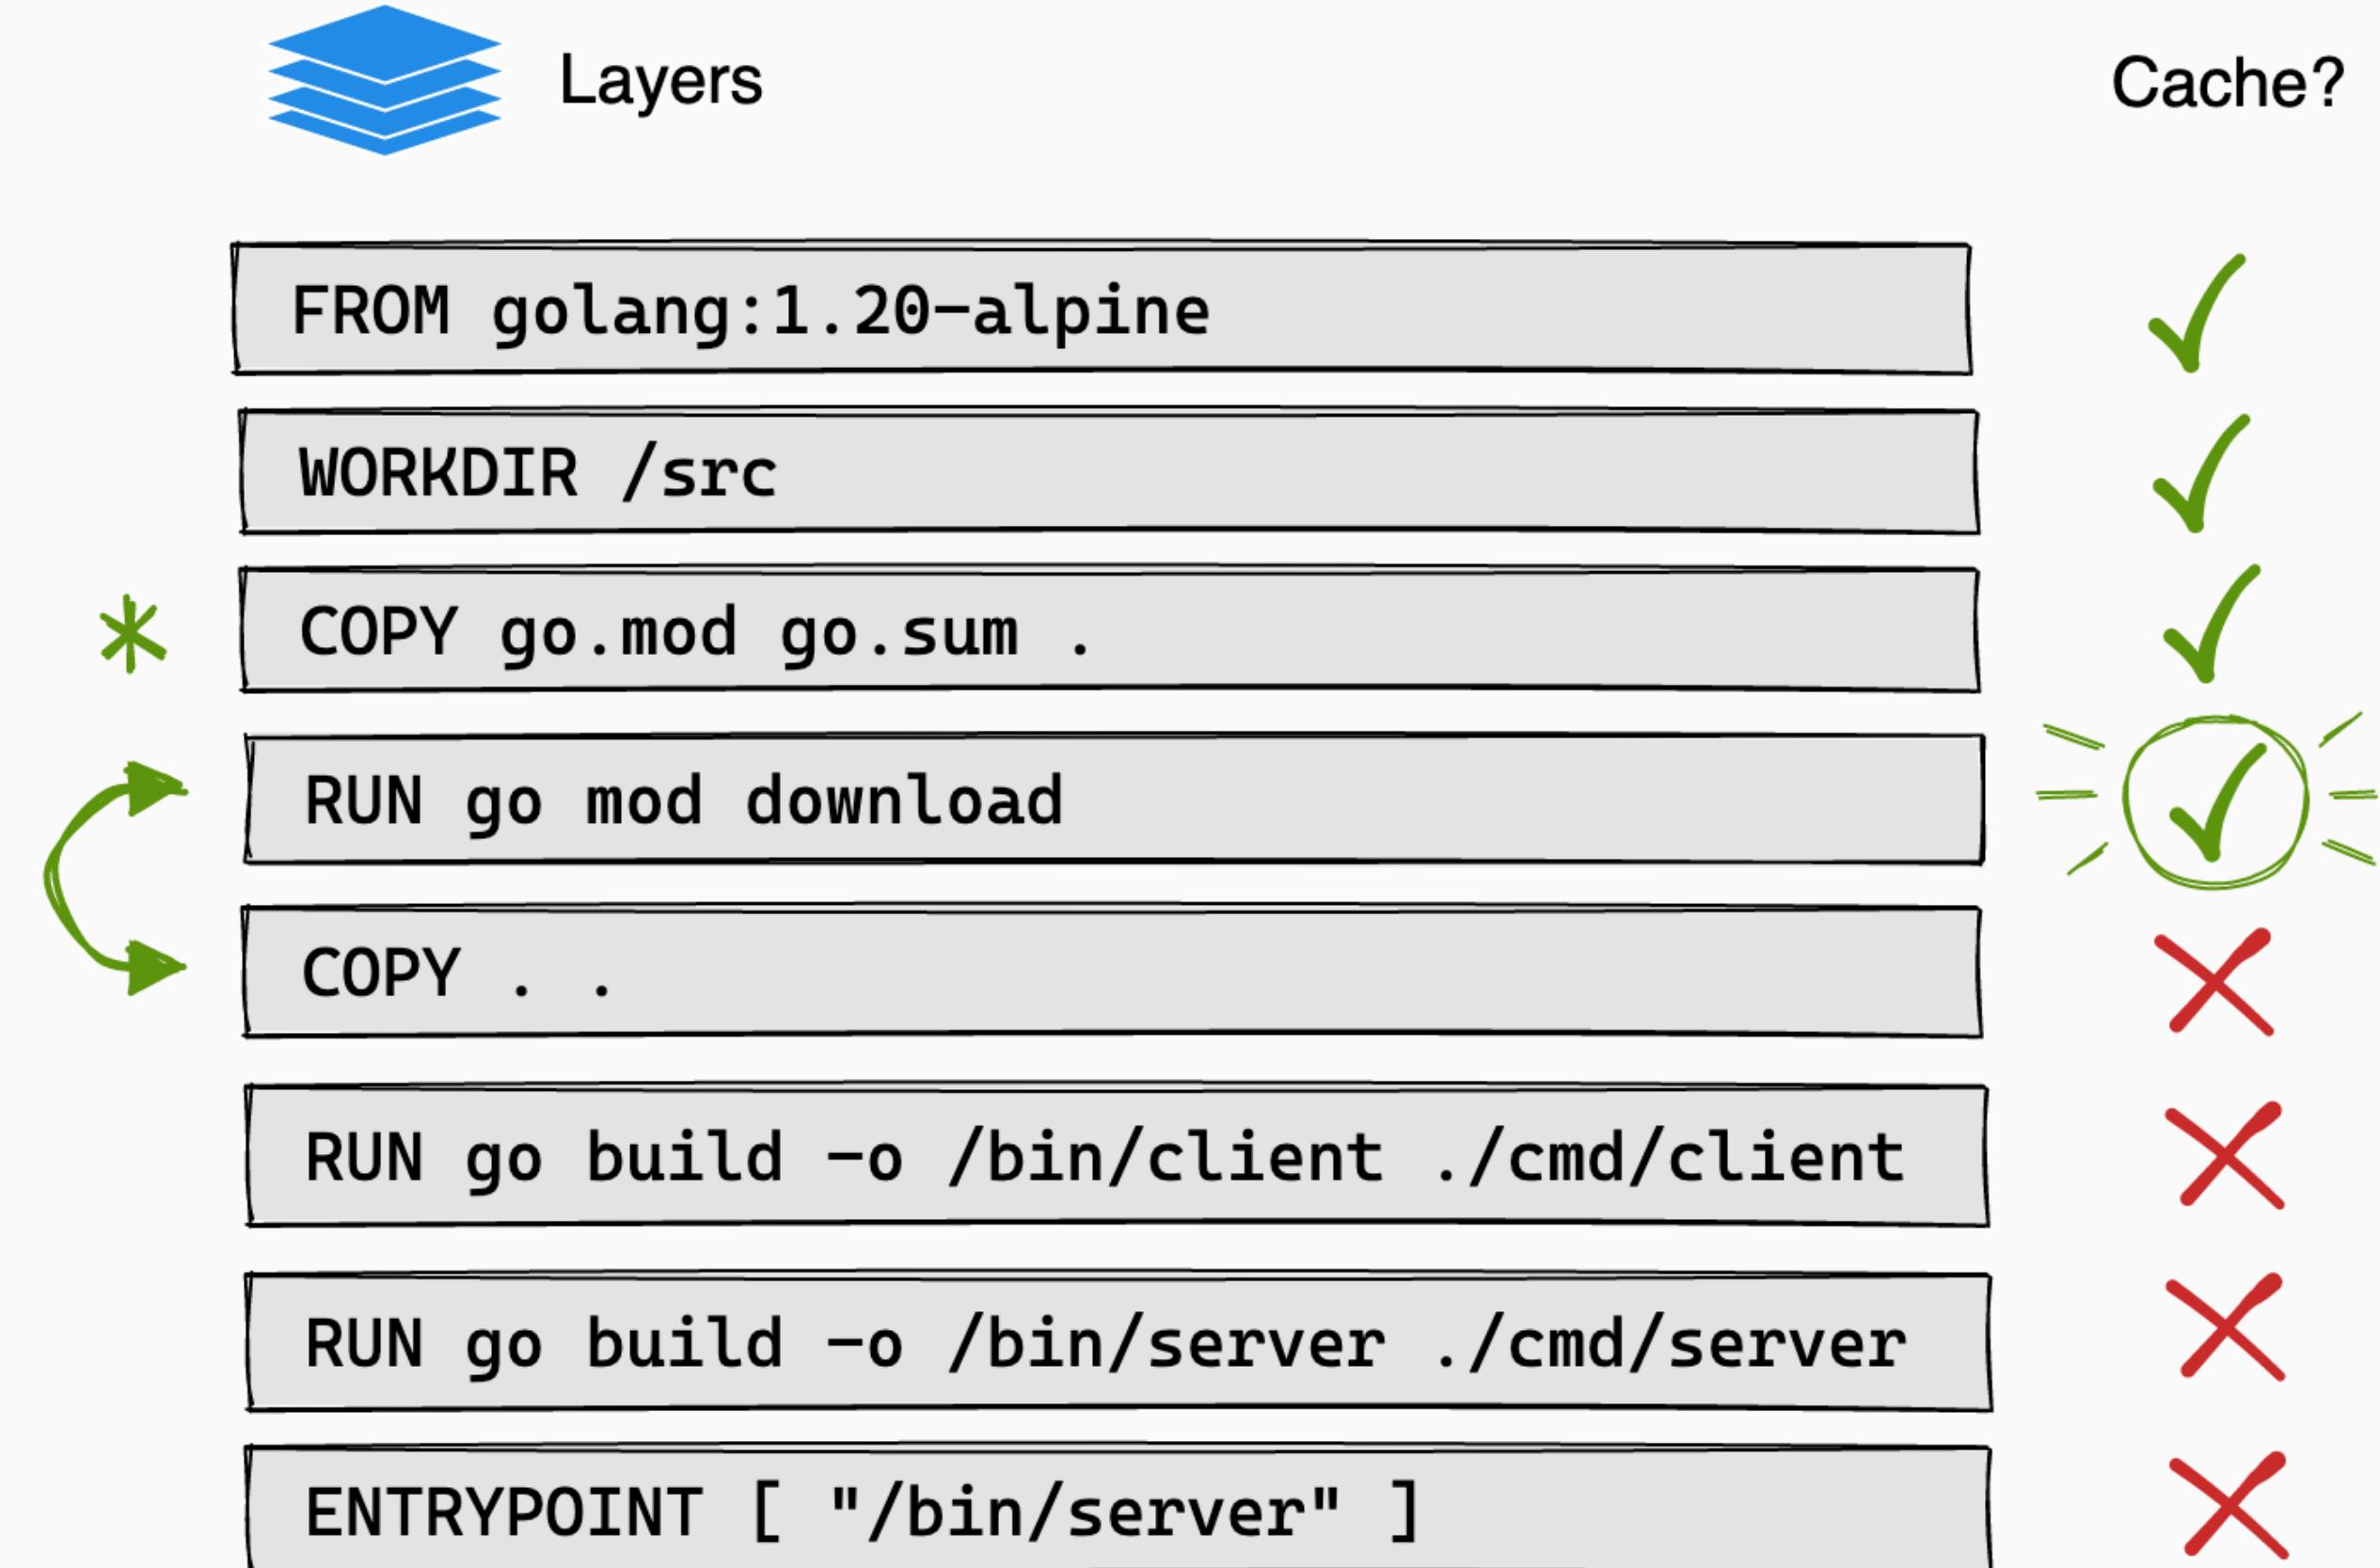
Caching, Efficient Storage, Incremental Changes

- Layers represent changes in a Docker image
 - Each command in a Dockerfile (like RUN, COPY, ADD) creates a new layer
- Layers are shared across images (for efficiency)
 - Layers are stored only once
 - Faster image builds
- Layer caching speeds up image builds
 - Unmodified layers are reused

Cached Layers

Layers	Cache?
FROM golang:1.20-alpine	✓
WORKDIR /src	✓
COPY . .	✗
RUN go mod download	✗
RUN go build -o /bin/client ./cmd/client	✗
RUN go build -o /bin/server ./cmd/server	✗
ENTRYPOINT ["/bin/server"]	✗

Instruction order matters



Joke time

Why do Docker containers never get lonely?

Because they always have a lot of
“friends” (processes) running inside.

Docker Networking

The Network Subcommand

```
docker network ls
```

```
docker network inspect bridge
```

```
docker network inspect host
```

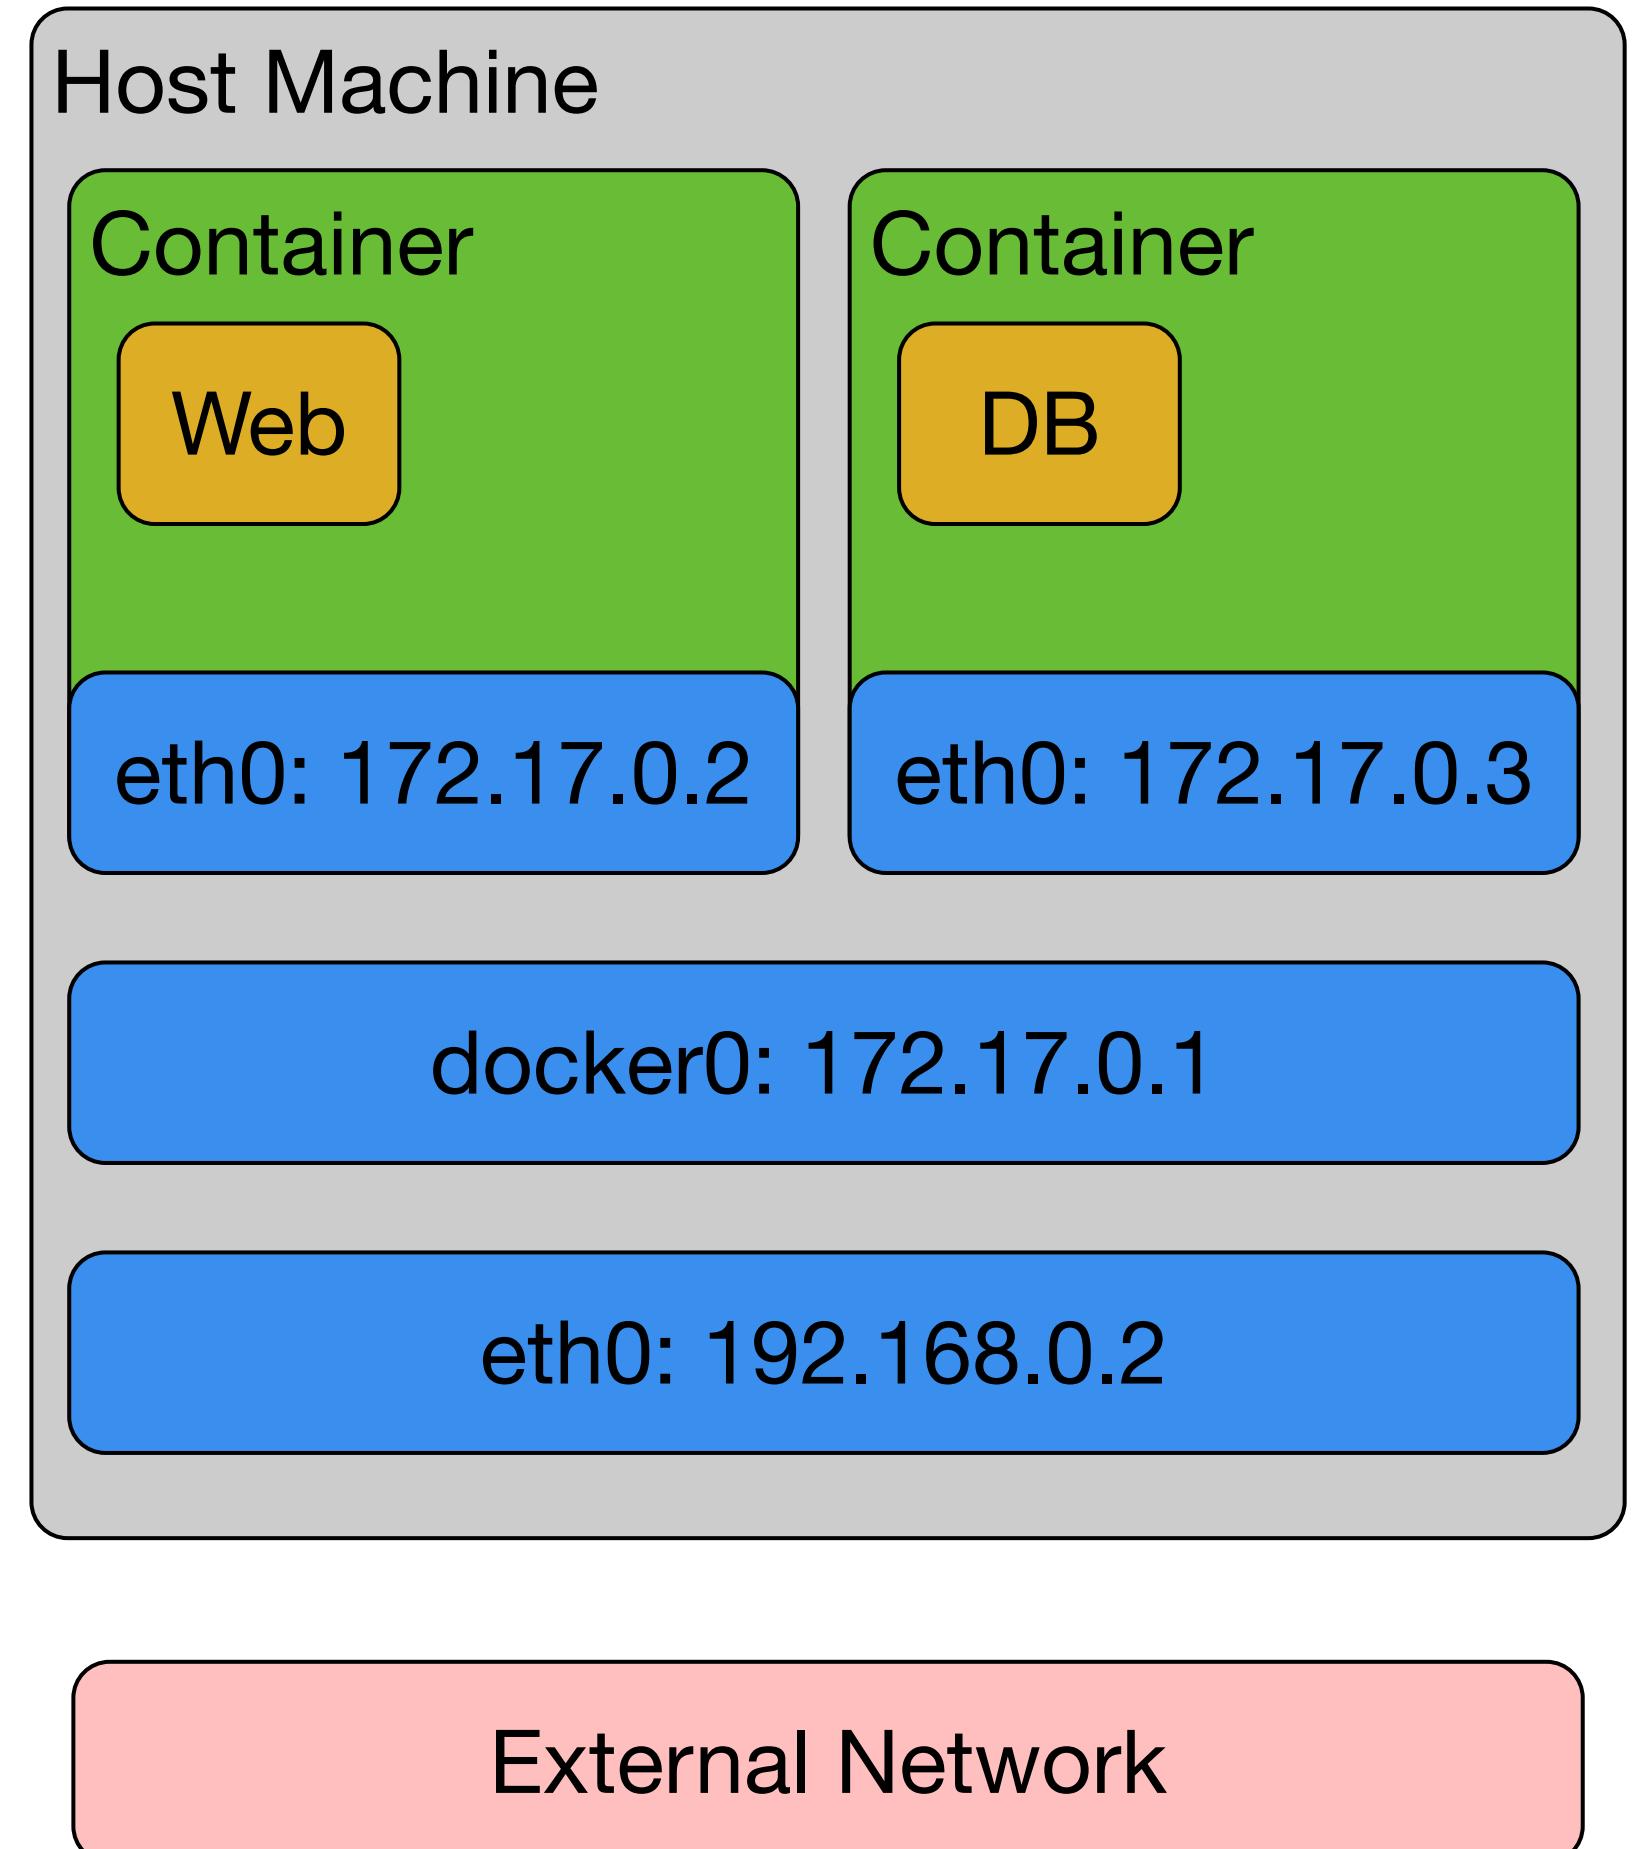
Default networks

```
docker network ls
```

NETWORK ID	NAME	DRIVER	SCOPE
9dcd3906d274	bridge	bridge	local
22096b6a3066	host	host	local
f10efee3f300	none	null	local

Bridge Networking (Default)

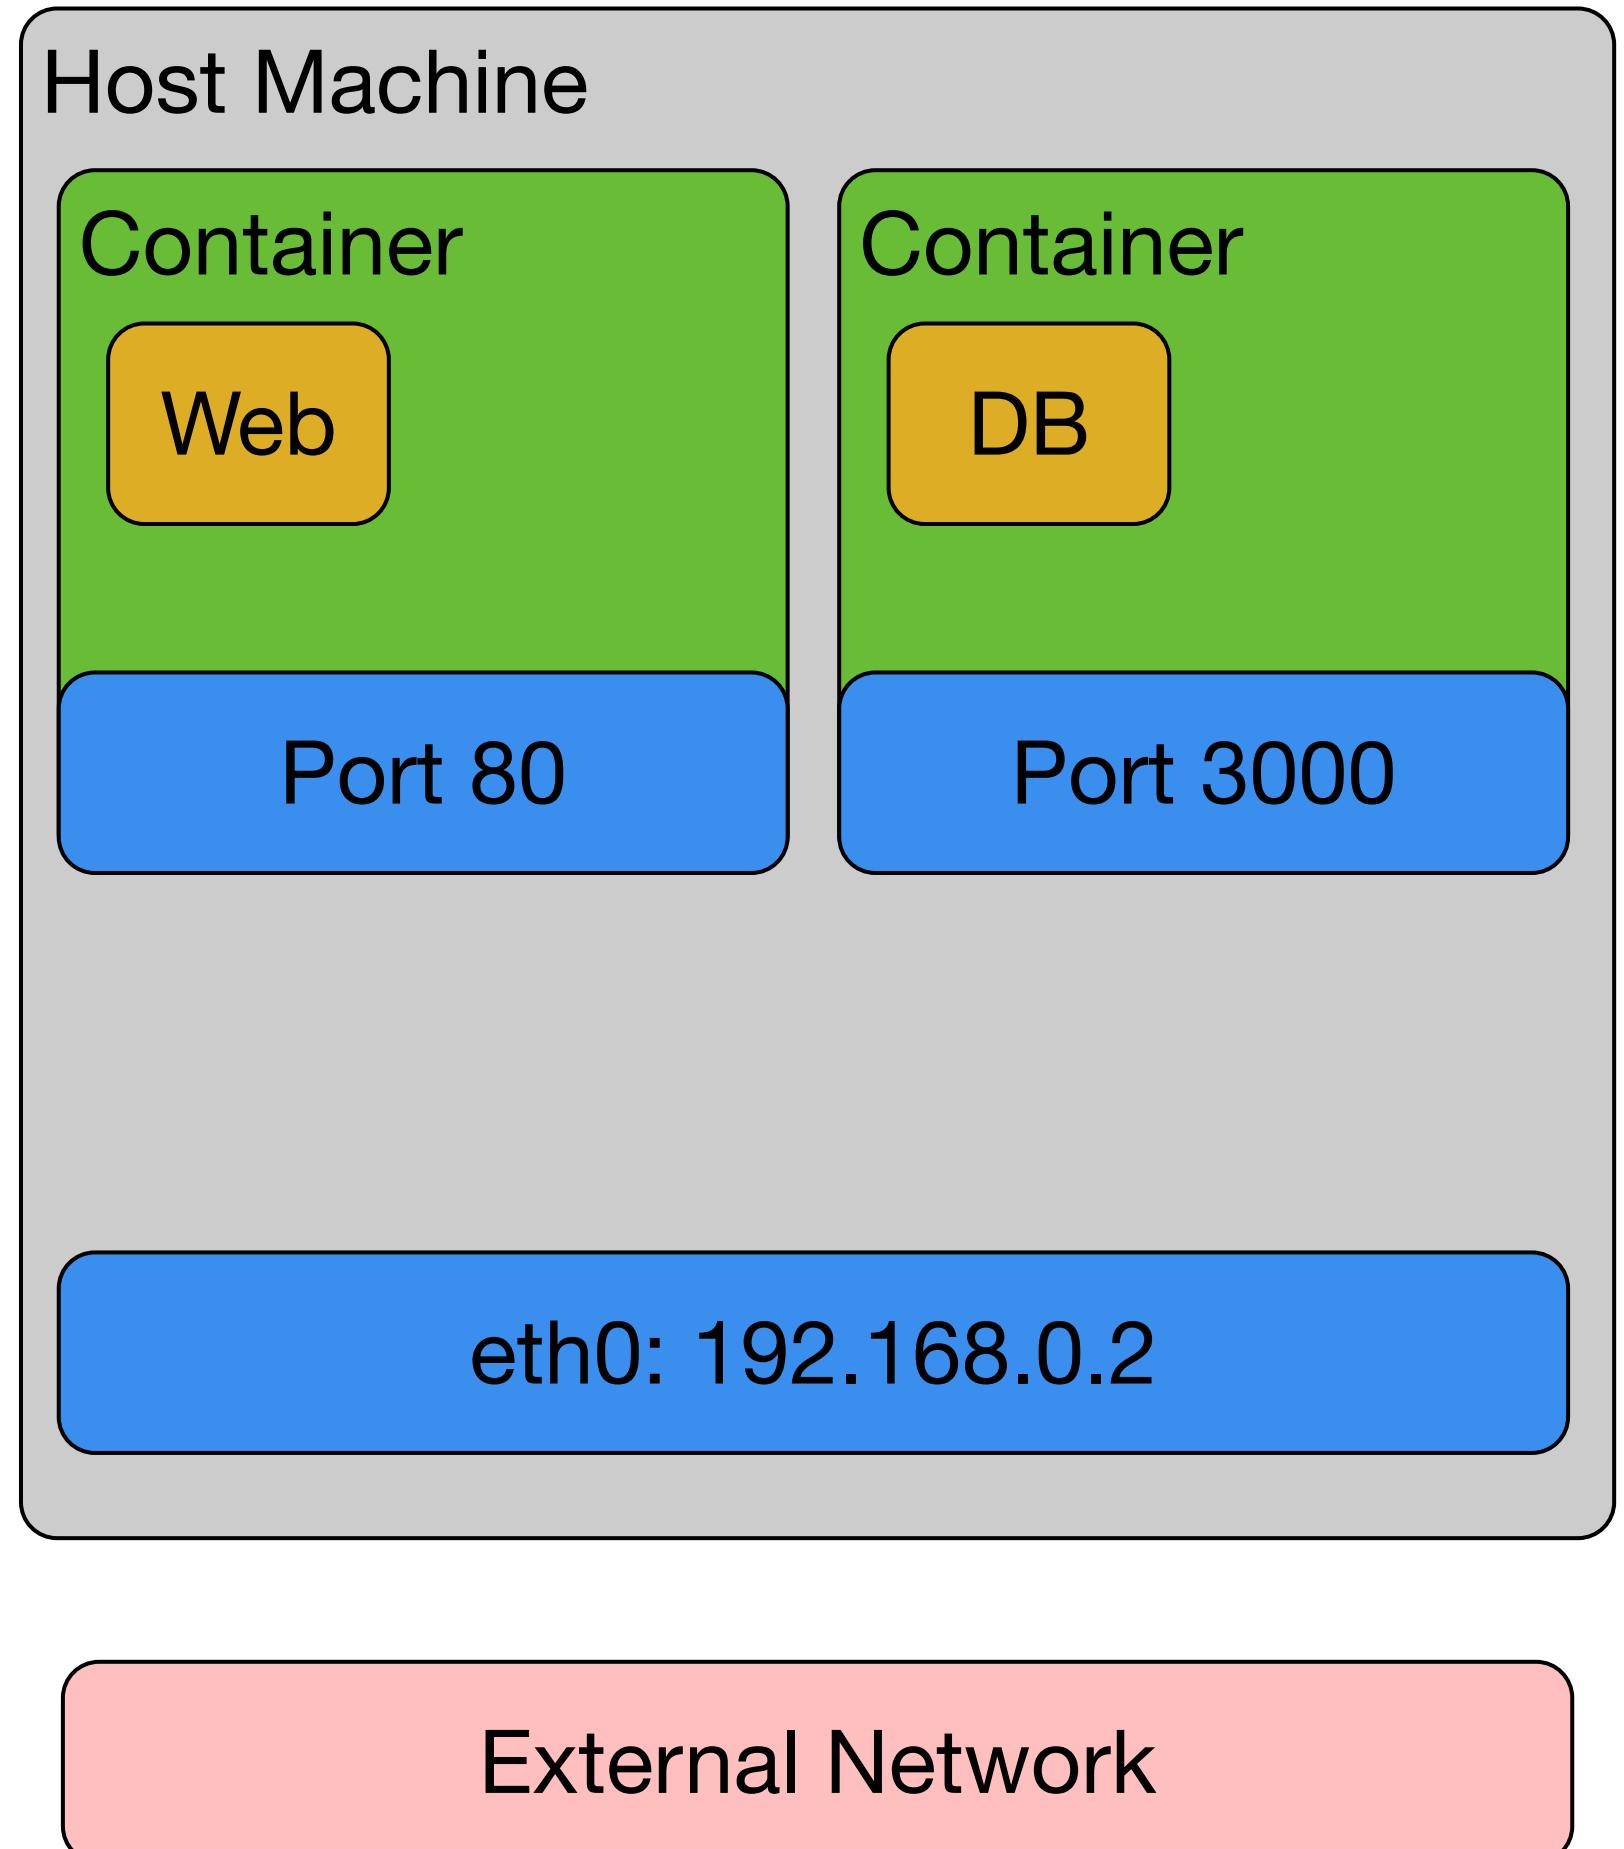
- Network for containers on a single host (few hosts)
- Provides **isolated network** among containers
 - Automatic IP addressing
- External access via **port mapping**
- Slight performance overhead compared to host networking



docker run ubuntu

Host Networking

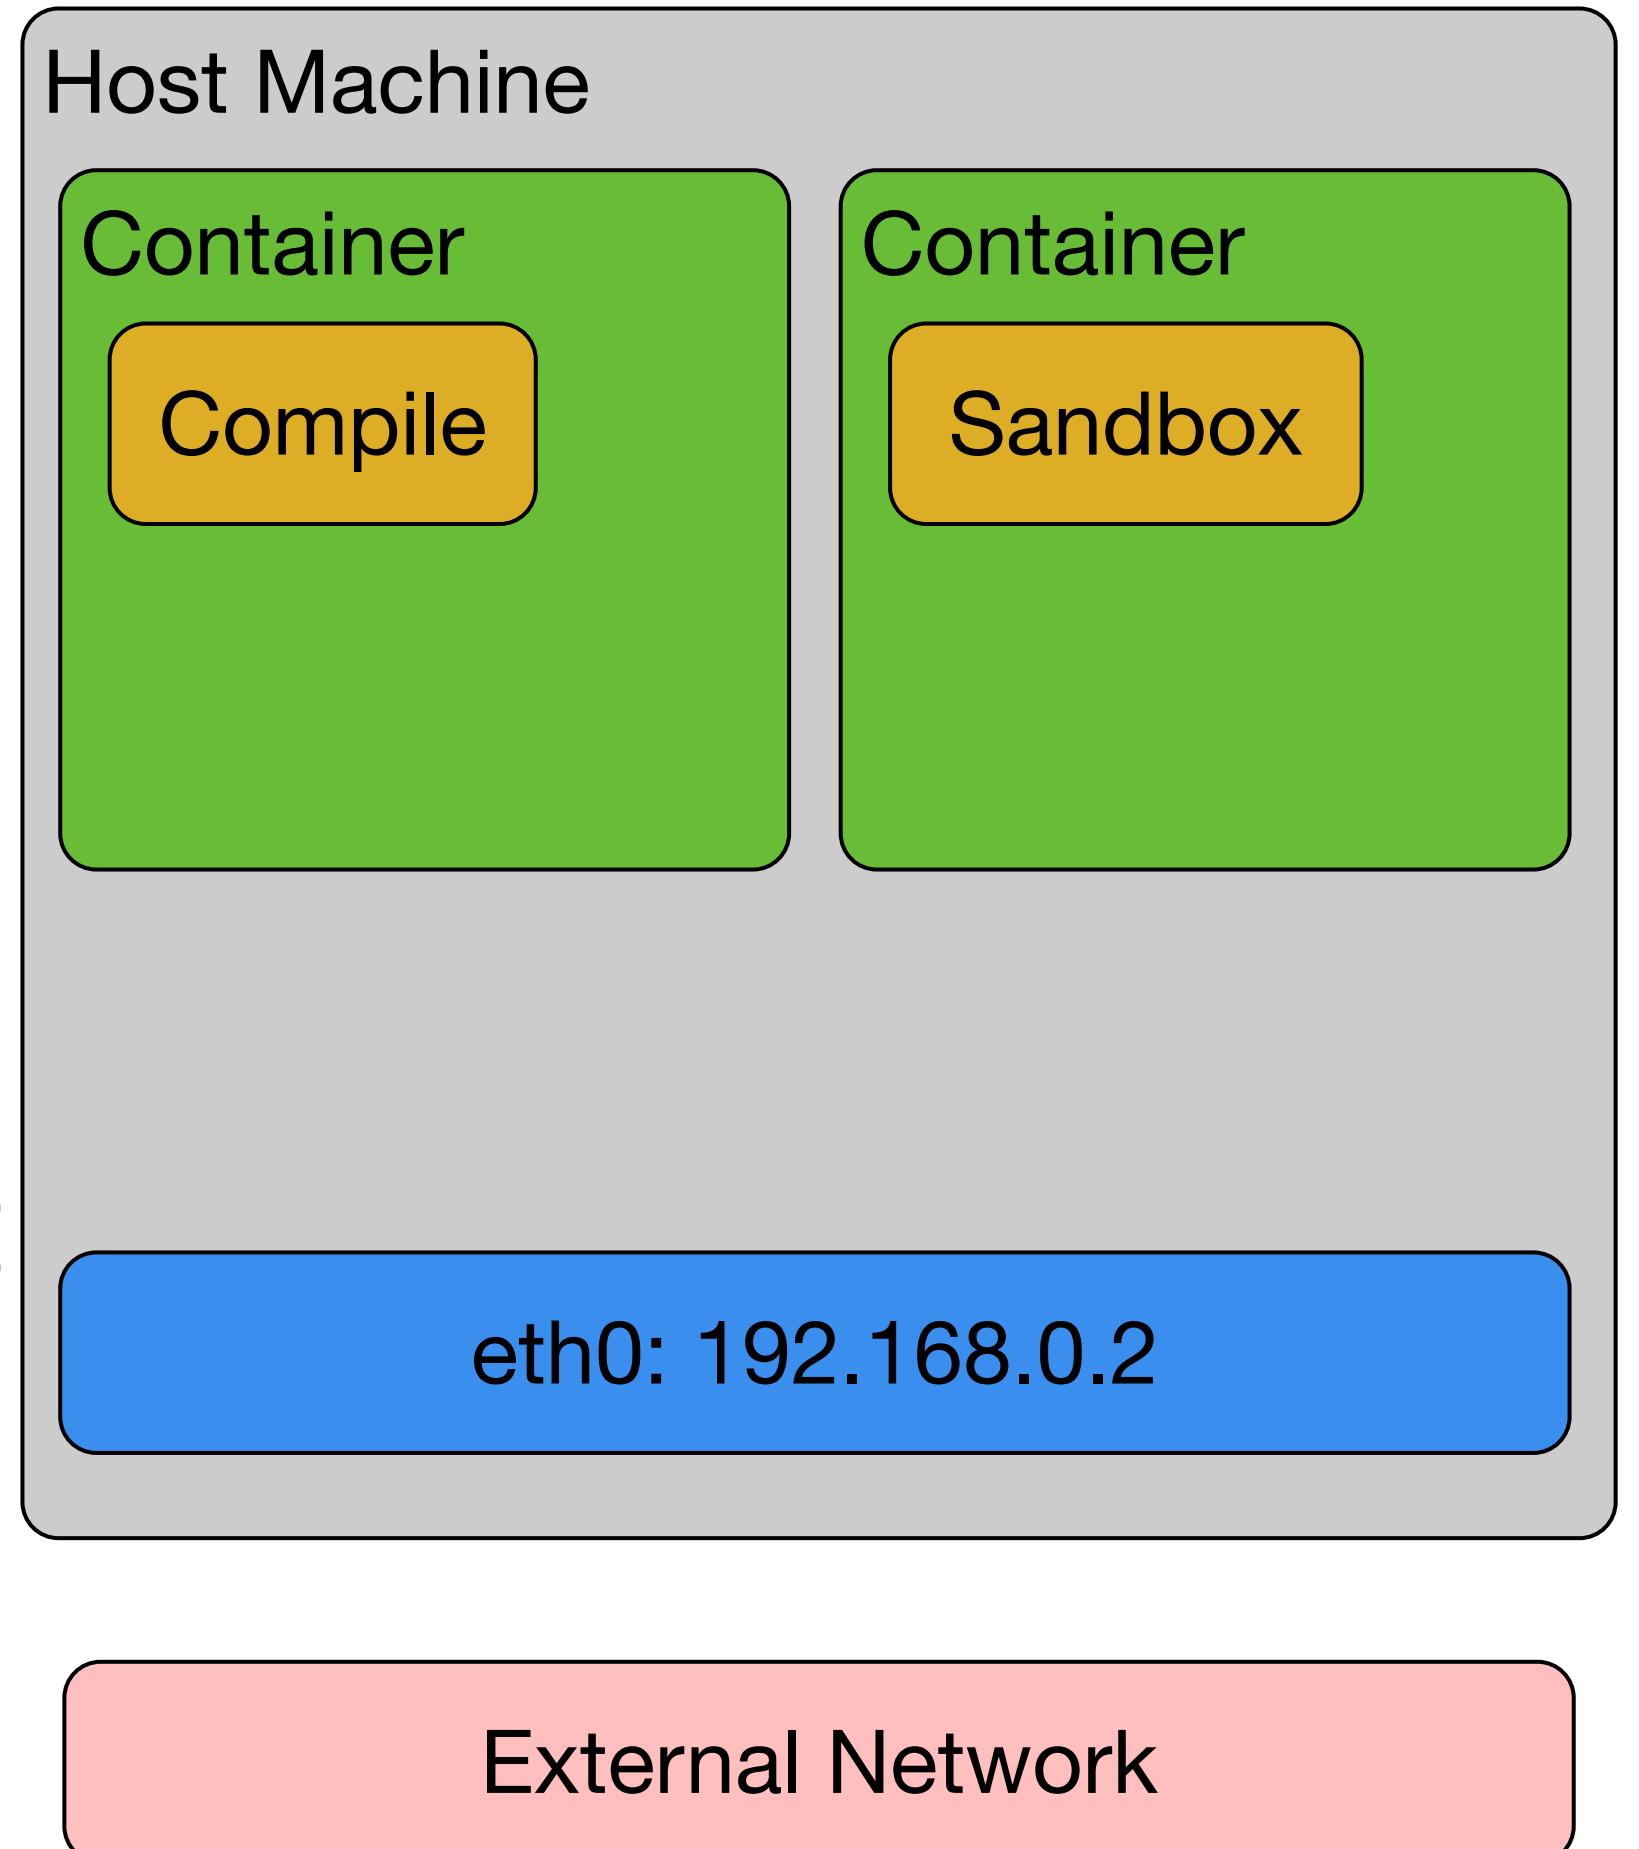
- Containers **share** the host's **network stack**
- Reduces network isolation
 - Share host's network namespace (IP, port)
 - Potentially **less secure**
- **Better performance** than bridge networking



```
docker run --network=host ubuntu
```

None Network

- Containers with **complete network isolation**
- Useful for tasks that don't require network access
 - Consistent and reproducible builds
 - Difficult to configure dependencies and environment
- **Output to file system** only

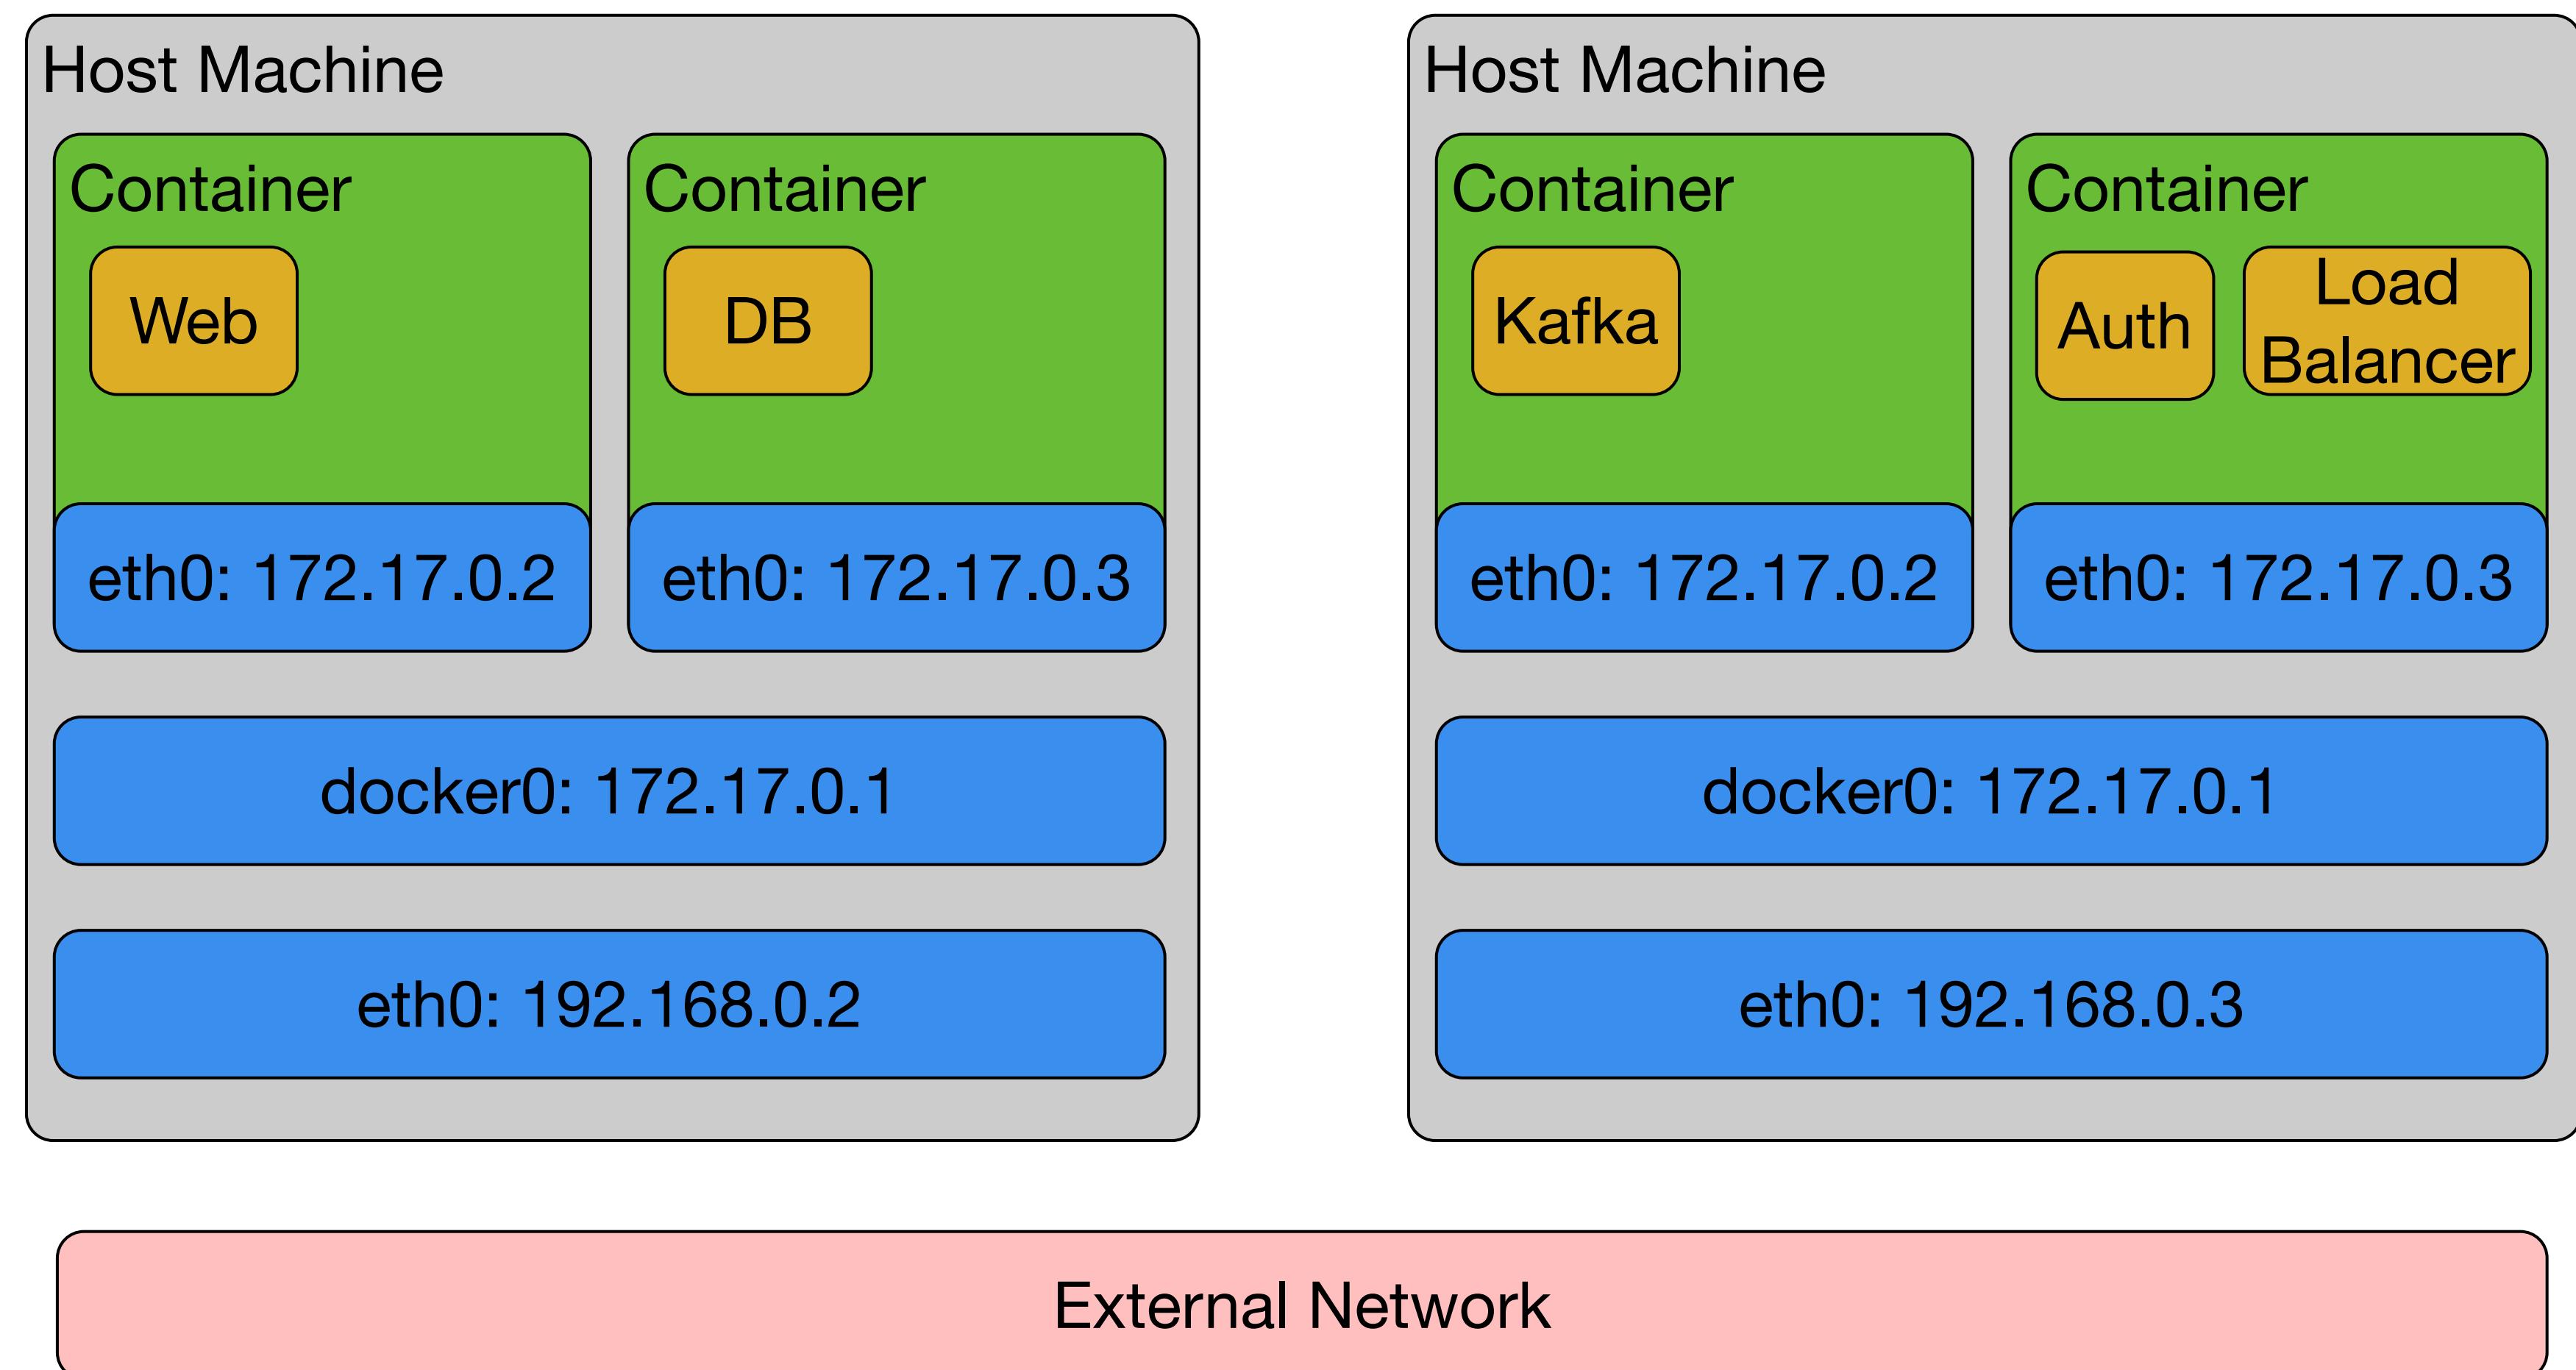


```
docker run --network=none ubuntu
```

Overlay Networking

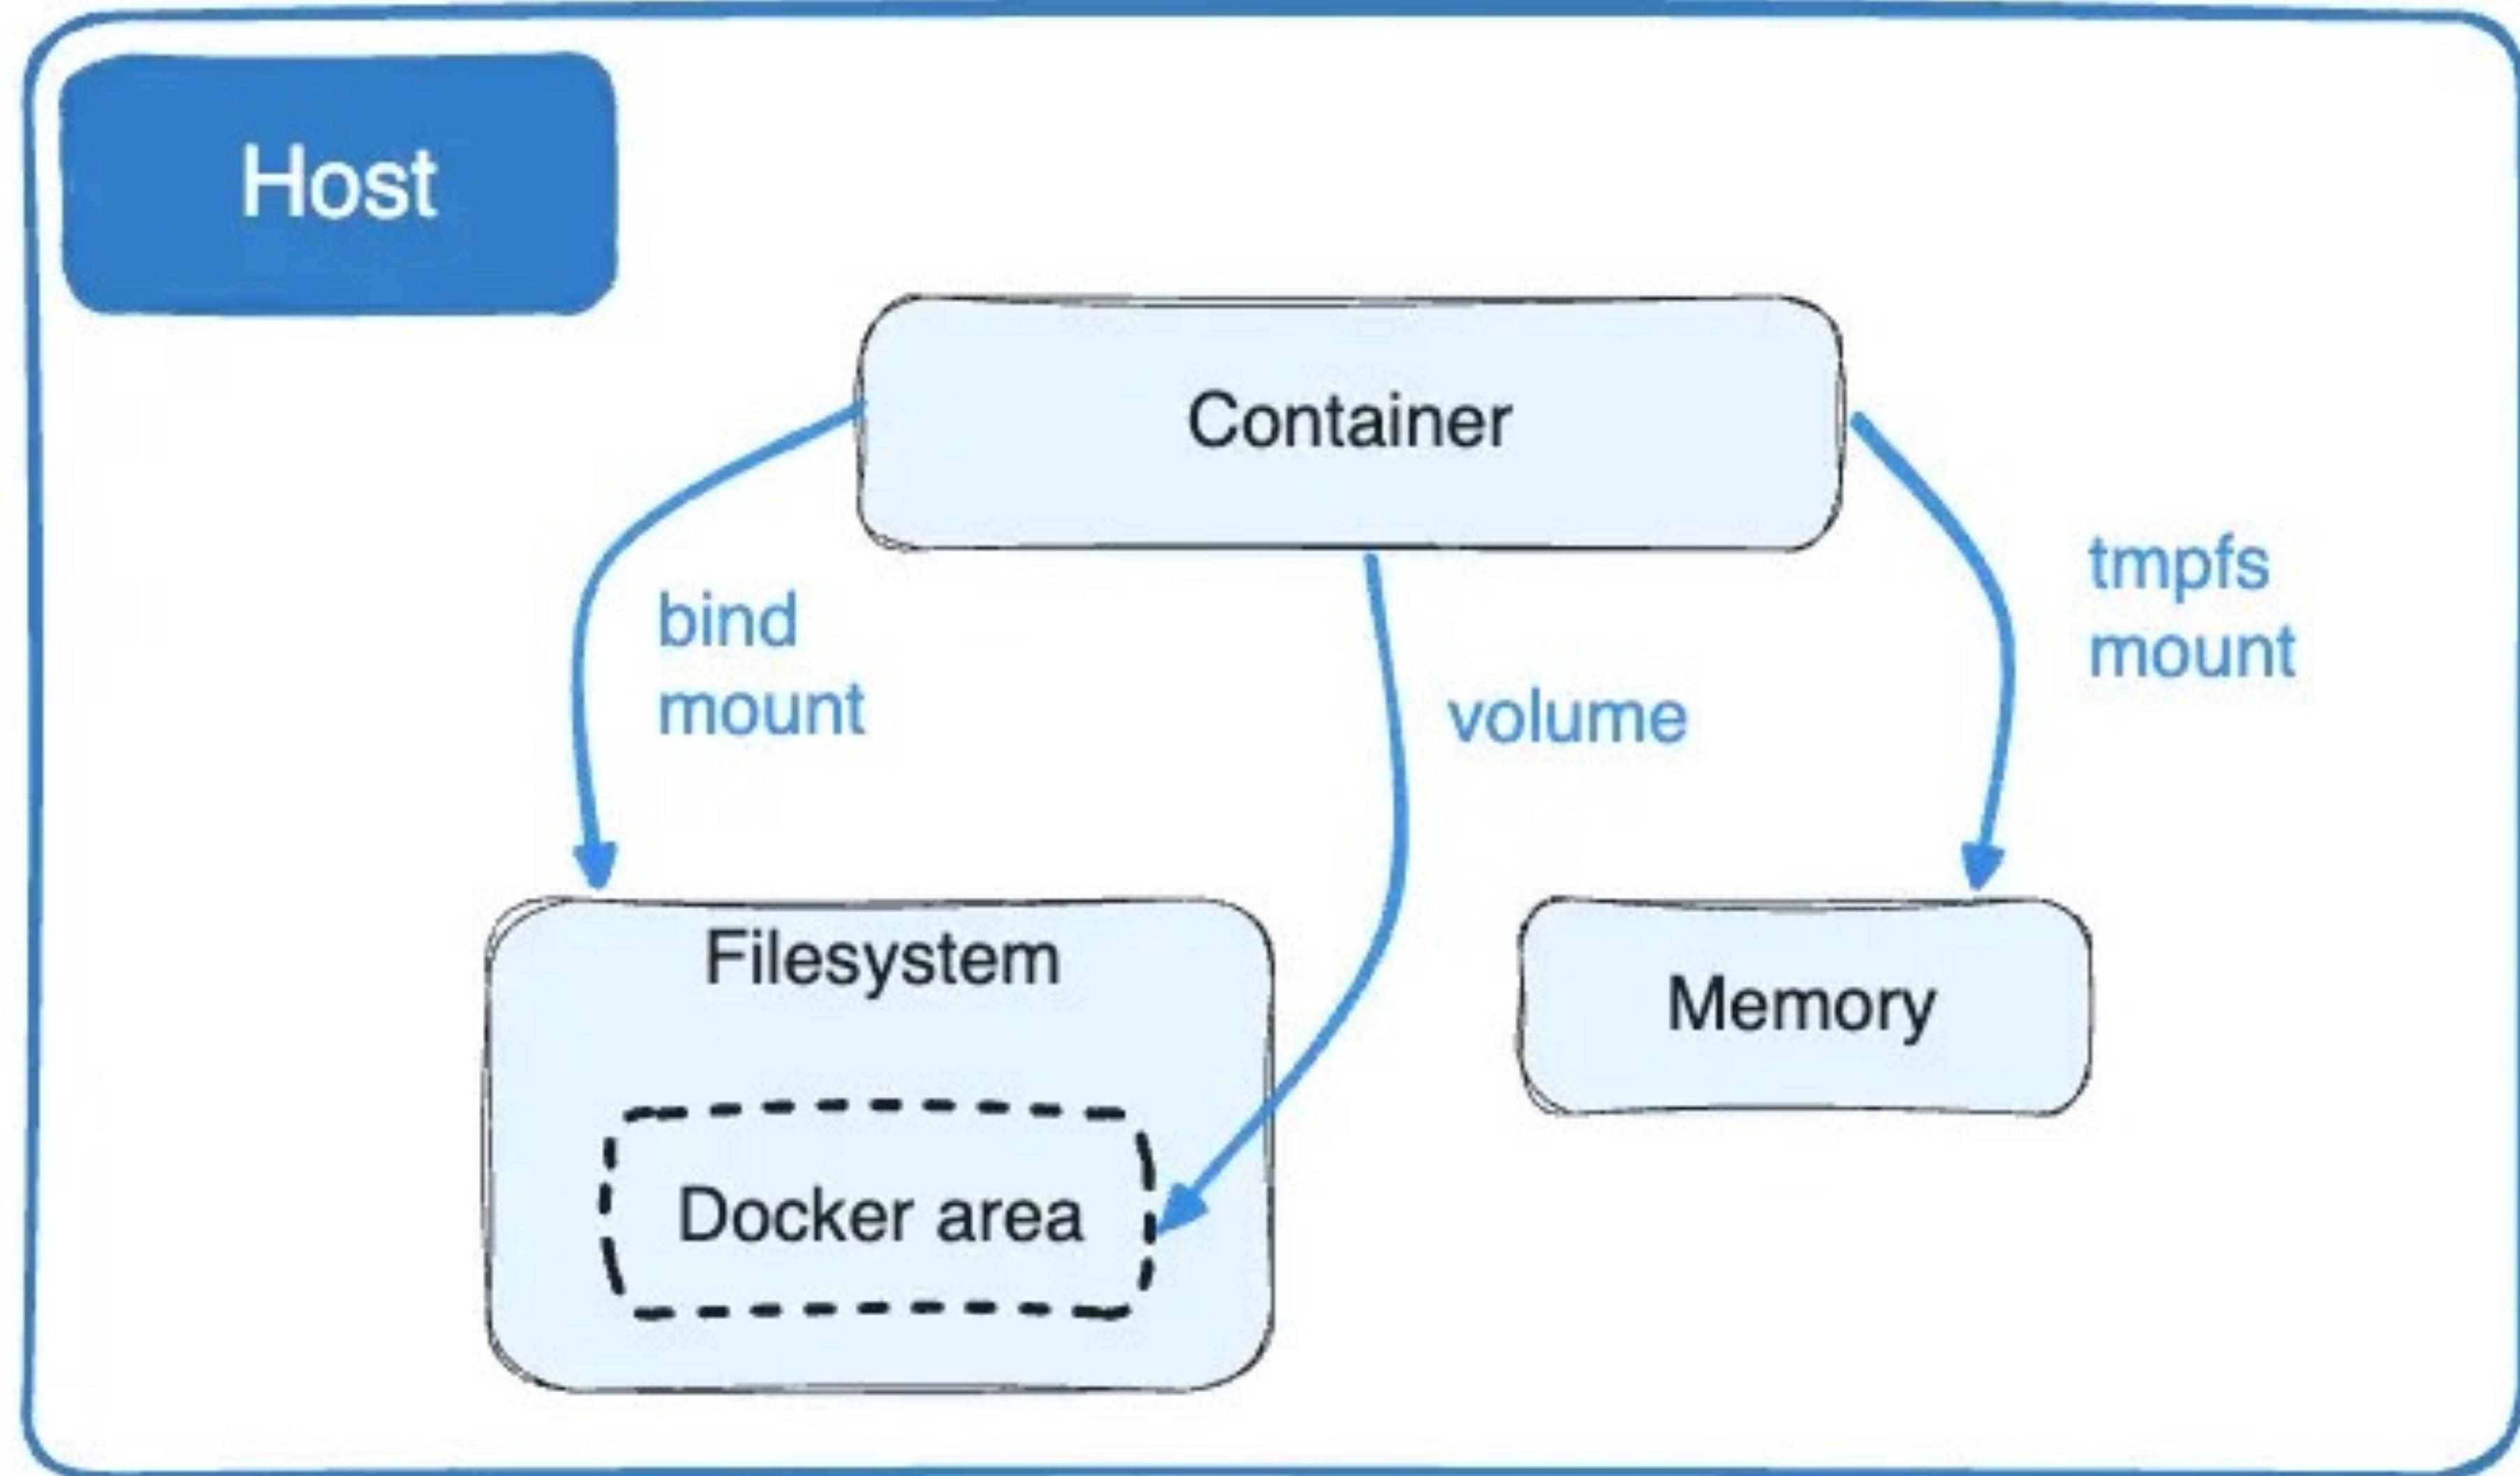
- Multi-host networking
 - Communication between **containers on different hosts**
- Automatic service discovery and configuration
- Hosts managed by Docker Swarm or Kubernetes clusters

Overlay Networking



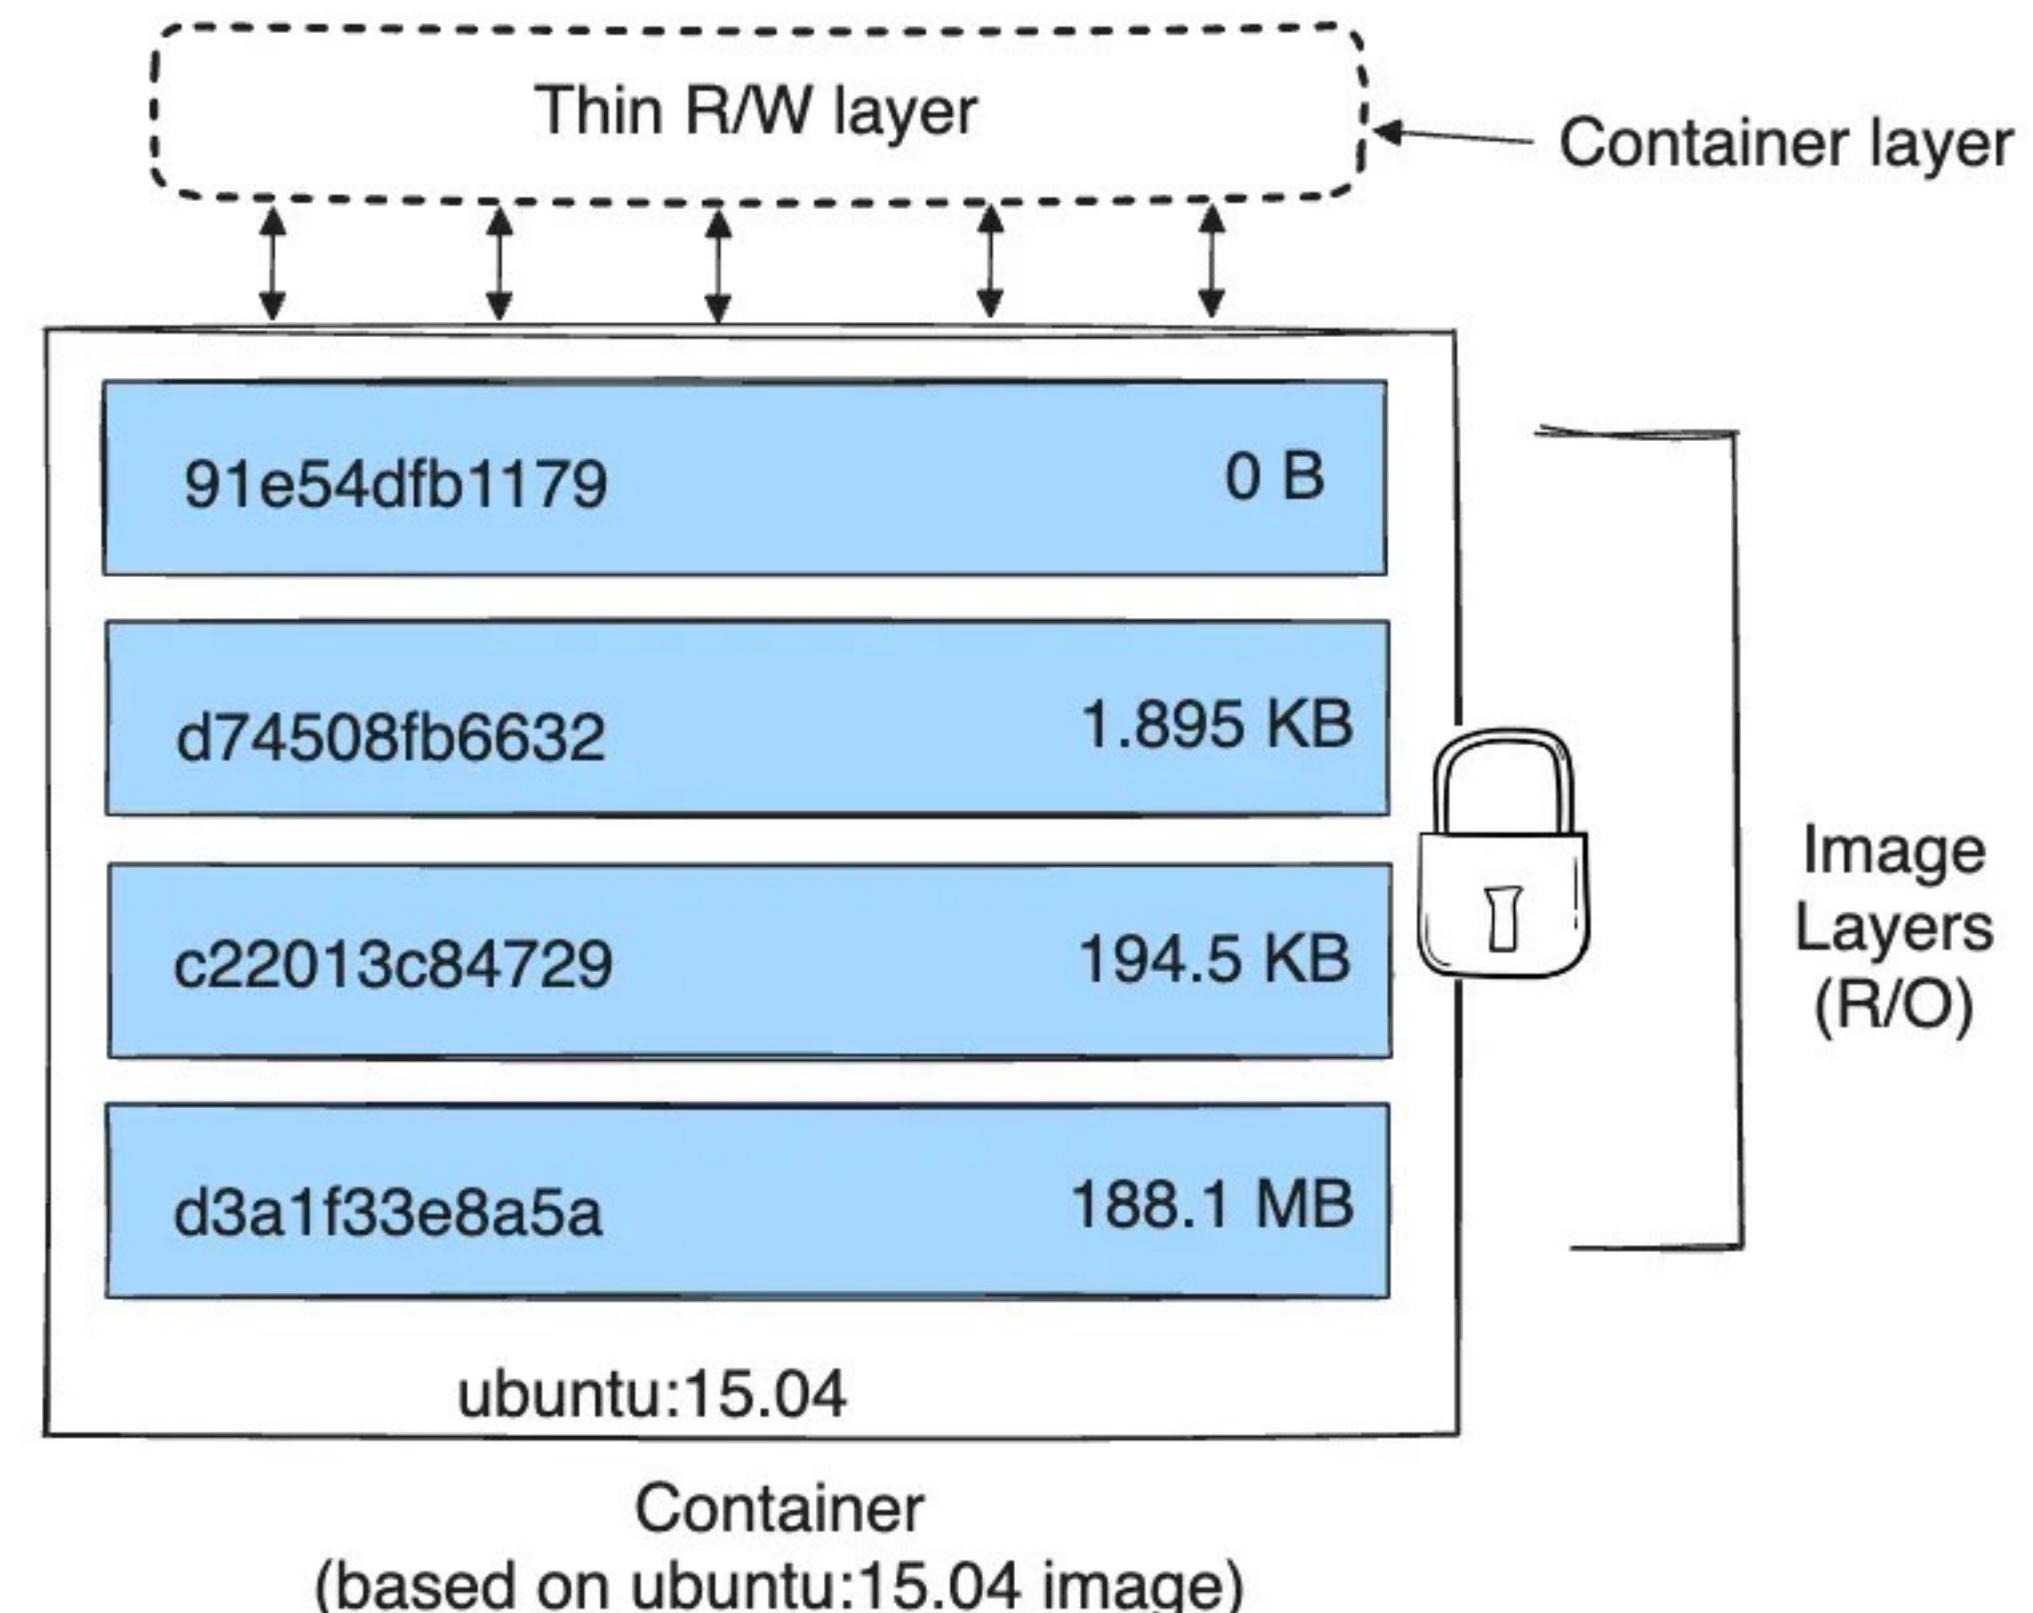
Docker Storage

Docker Storage Options



Layered File System

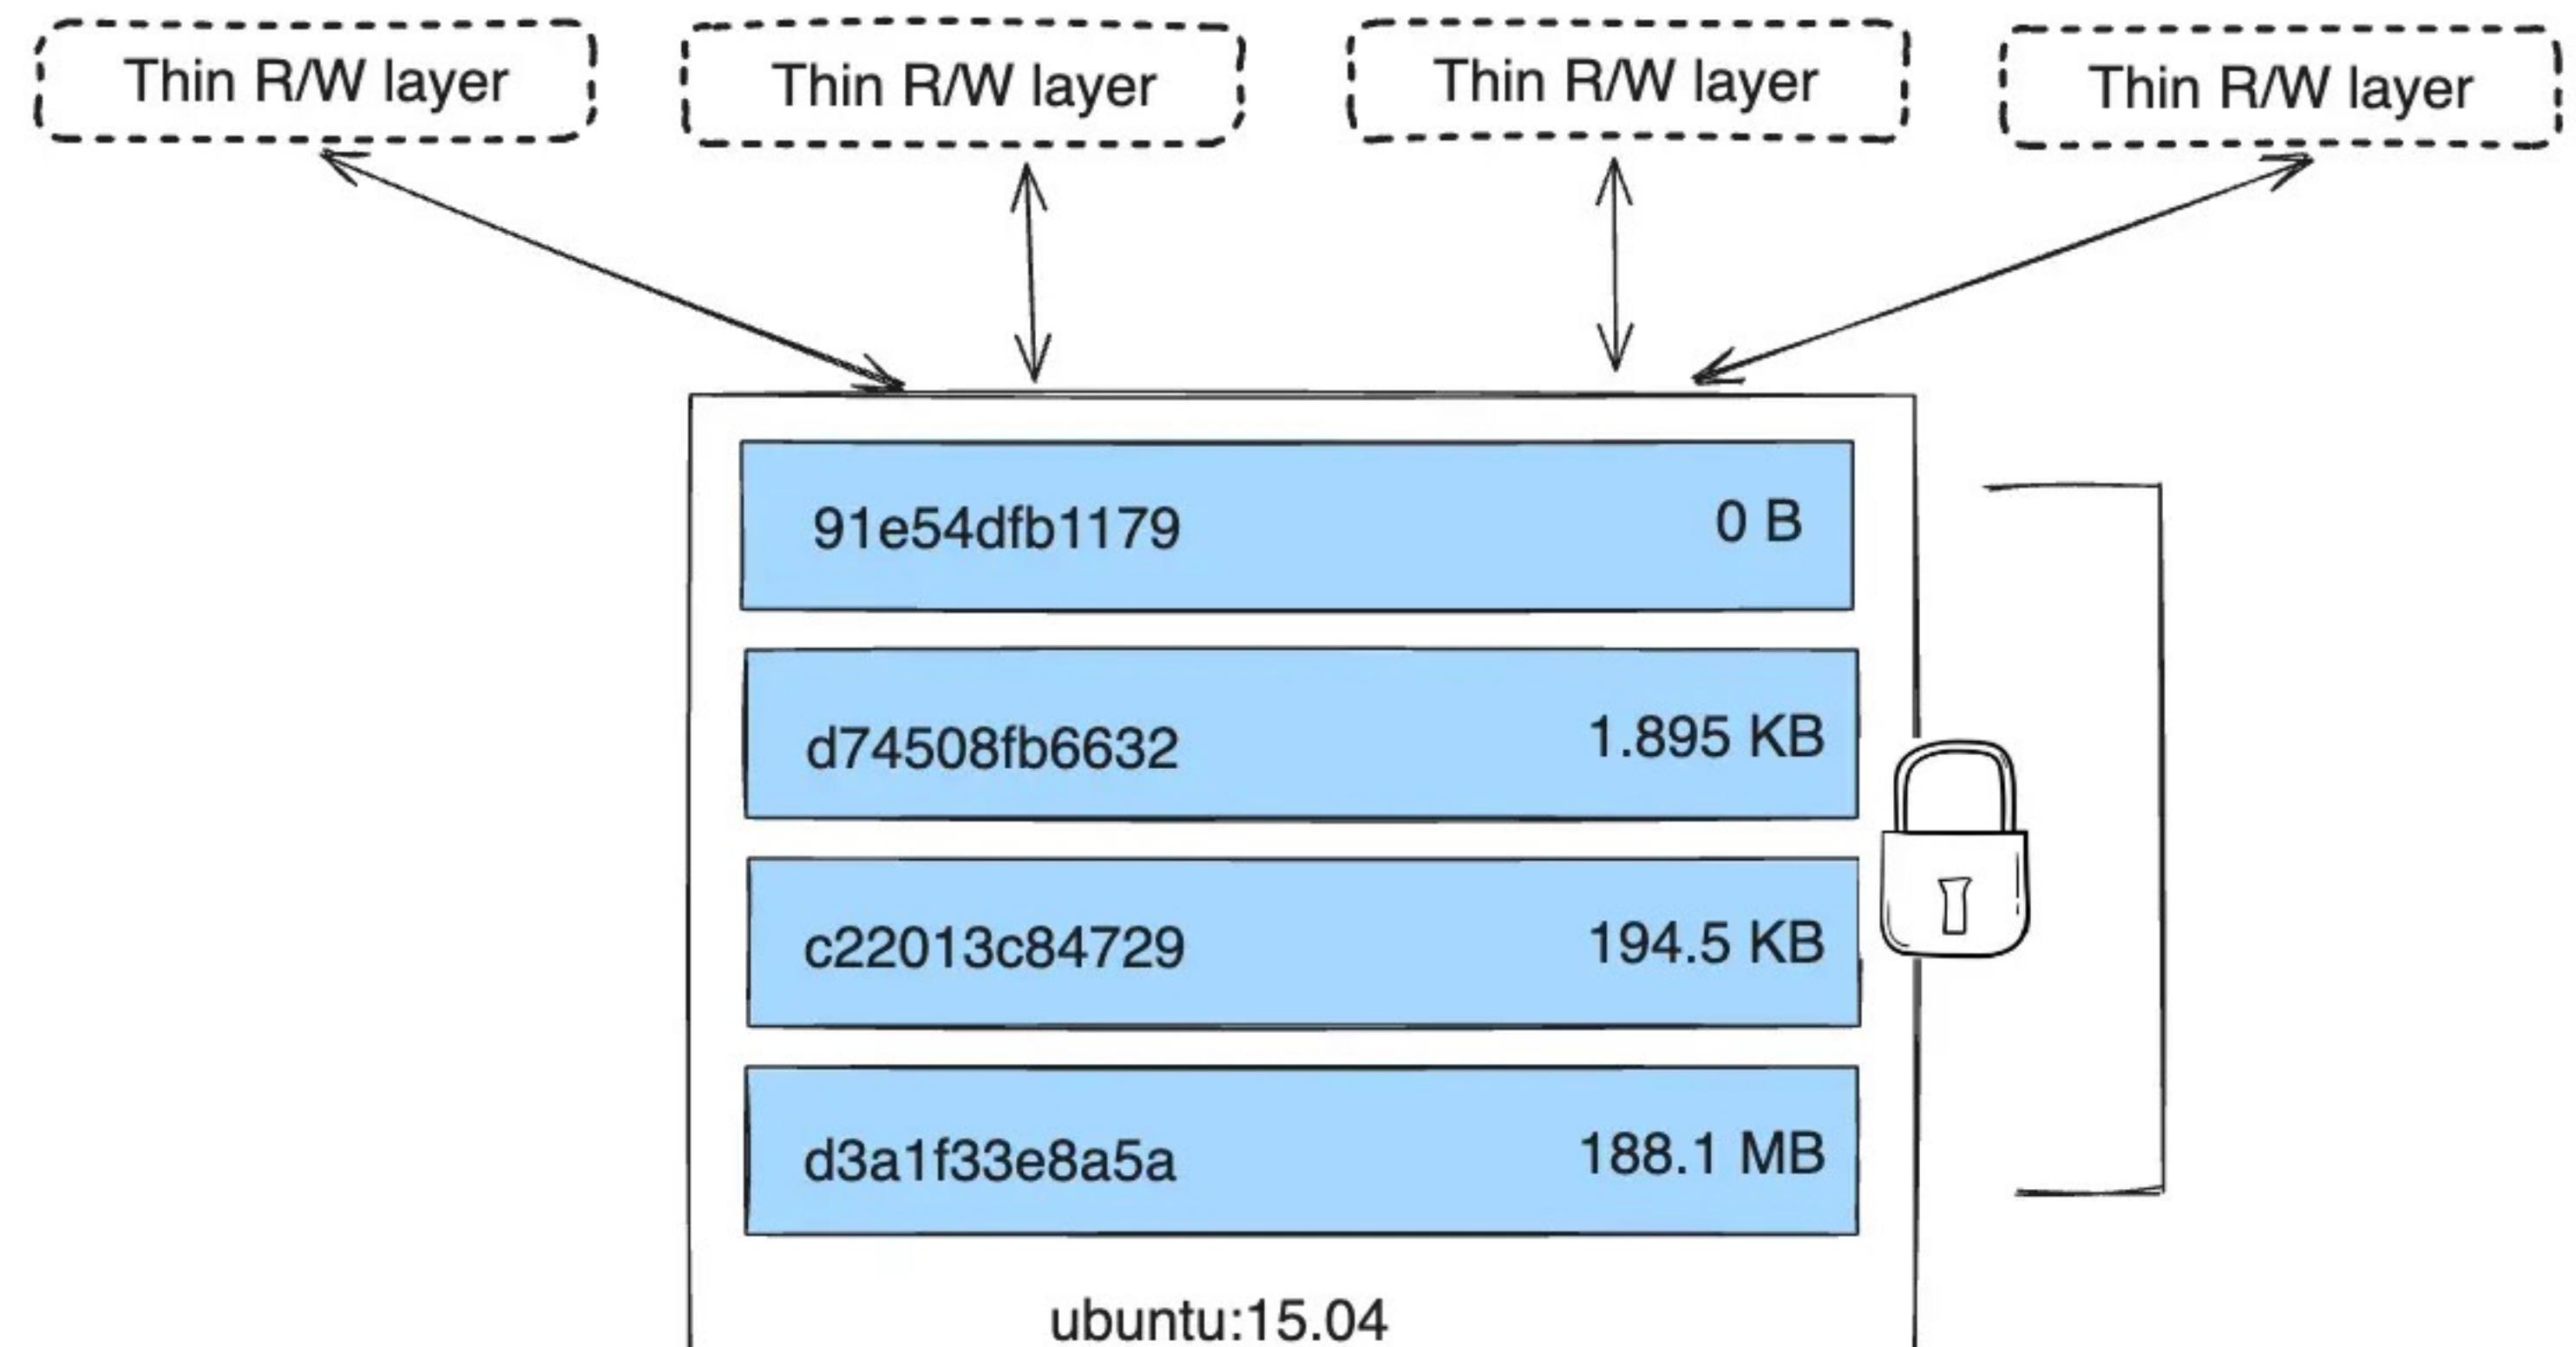
- Ephemeral storage
- **Data retained** across stops and **restarts** of the container
- **Data is lost** when the container is **deleted**
- Mainly useful for temporary data



Containers Reusing Lower Layers

Layered File System

- Different containers can use ubuntu:15.04
- Each container get their own writable layer



Data stored in the host's memory (RAM)

- Ephemeral storage
- **Data is lost** when the container is **stopped**
 - Useful for caches or session data

Persistent data storage outside container's filesystem

- **Persistent Storage**

- Ensure data persists across container restarts
- Managed by Docker /var/lib/docker/volumes

- **Shared Access**

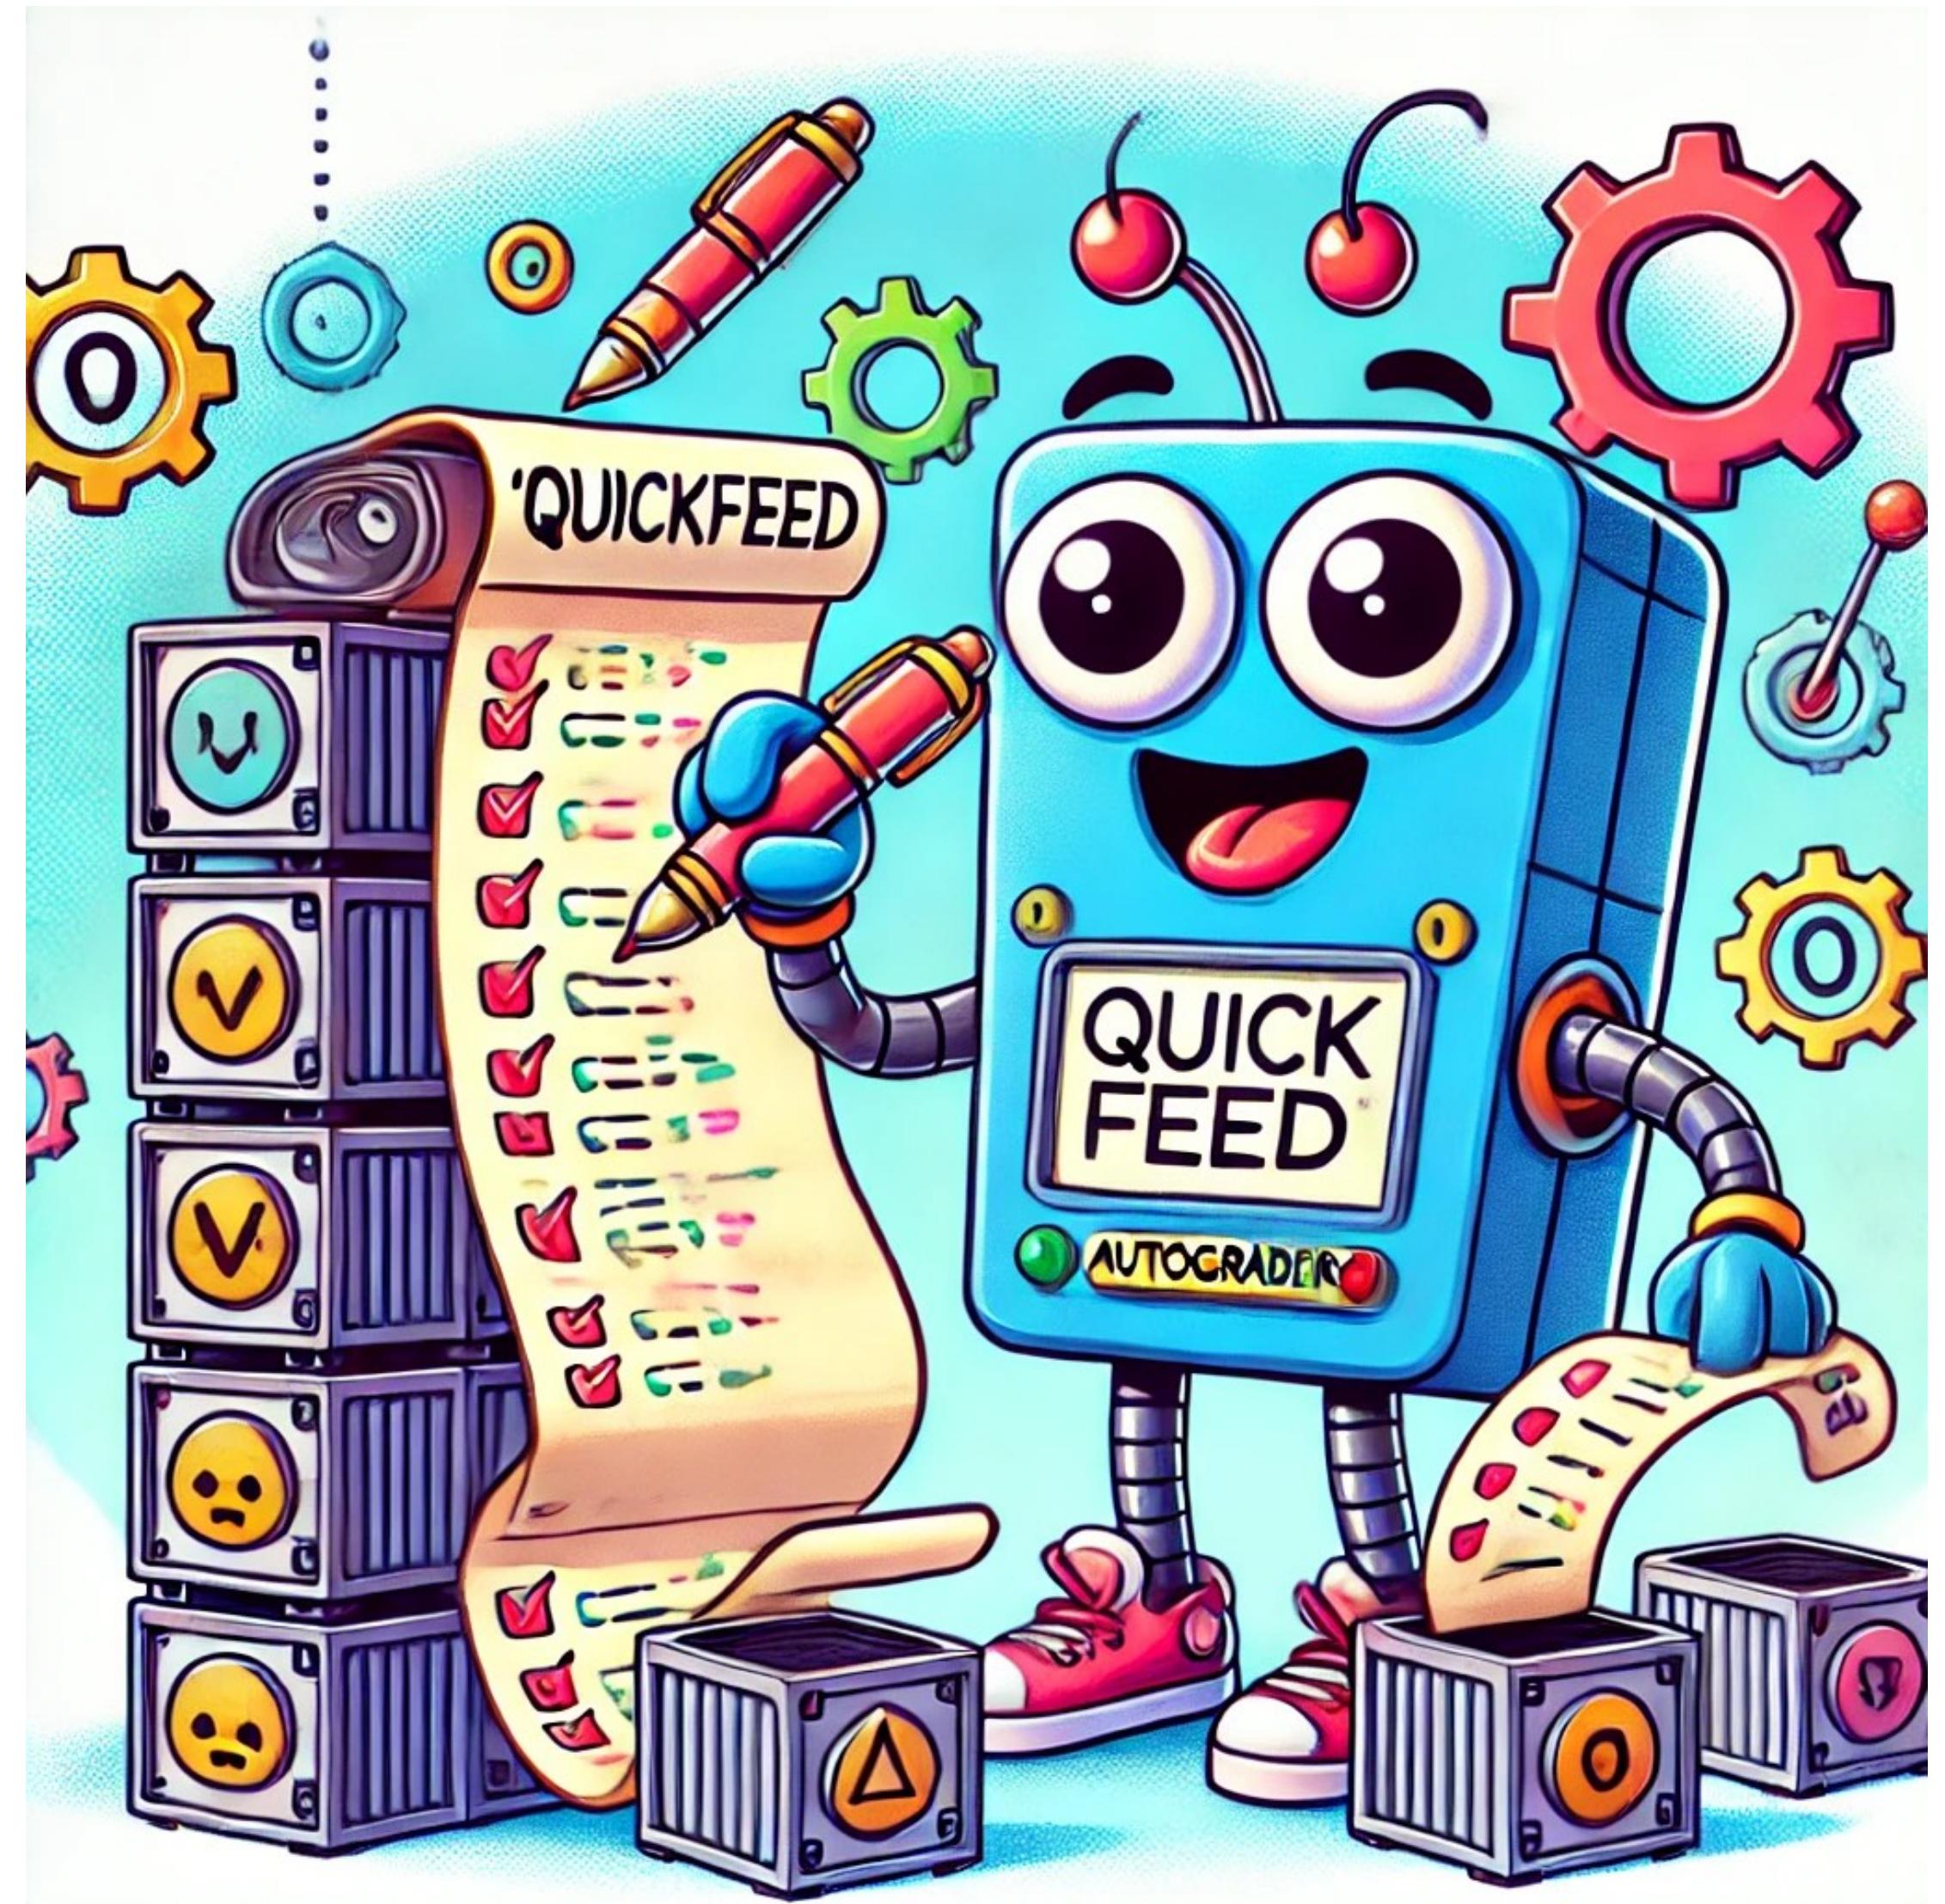
- Share data between containers by mounting the same volume
- Share data with host system
- Decouple data from container lifecycle, easy backups, and migration

- **Map specific file/directory** on **host's filesystem** to a path inside the container
- Use Cases: Access specific content on host file system
 - Source code and configuration files
 - (See QuickFeed example later)
- Tightly coupled to host's filesystem structure
- Less portable than Docker managed volumes

Summary of Storage Alternatives

Storage Option	Description	Usage	Persistence Across Start/Stop	Persistence After Container Removal
Volumes	Managed by Docker, stored outside the container's filesystem	Sharing data between containers, backups, migration, decoupling from lifecycle	Yes	Yes
Bind Mounts	Directly maps a host directory or file to a path inside a container	Accessing specific host files or directories, development (e.g., syncing code)	Yes	Yes
tmpfs Mounts	Stores data in the host's memory (RAM), not on disk	Temporary data that doesn't need persistence (e.g., caching, session data)	No	No
Layered Filesystem	Manages container's filesystem layers; data in the writable layer	Default container storage, ephemeral data storage within the container	Yes	No

QuickFeed Interlude



QuickFeed's Docker Bind Mount

```
tree -L1 courses
```

```
courses
├── dat240-2023
├── dat240-2024
├── dat310-2023
├── dat310-2024
├── dat320-2023
├── dat320-2024
├── dat515-2024
├── dat520-2023
├── dat520-2024
├── dat600-2024
└── qf-103
```

```
courses/dat515-2024/
```

```
├── assignments
│   ├── lab1
│   ├── lab2
│   ├── lab3
│   ├── lab4
│   ├── lab5
│   └── lab6
└── tests
    ├── lab1
    ├── lab2
    ├── lab3
    ├── lab4
    ├── lab5
    └── lab6
```

QuickFeed's Docker Bind Mount

```
tree -L1 courses
courses
├── dat240-2023
├── dat240-2024
├── dat310-2023
├── dat310-2024
├── dat320-2023
├── dat320-2024
├── dat515-2024
├── dat520-2023
├── dat520-2024
├── dat600-2024
└── qf-103
```

```
courses/dat515-2024/
├── assignments
│   ├── lab1
│   ├── lab2
│   ├── lab3
│   ├── lab4
│   ├── lab5
│   └── lab6
└── tests
    ├── lab1
    ├── lab2
    └── lab3

courses/dat320-2023/
├── assignments
│   ├── go.mod
│   ├── go.sum
│   └── lab1
│       ├── generate_token
│       ├── git_questions.md
│       ├── missing_semester_questions.md
│       ├── README.md
│       └── shell_questions.md
└── lab2
    ├── c_questions.md
    ├── hello
    └── README.md

└── lab3
    ├── cipher
    ├── cmd
    ├── errors
    ├── go_questions.md
    ├── multiwriter
    ├── README.md
    ├── sequence
    └── stringer

└── tests
    ├── internal
    │   ├── git
    │   ├── staff
    │   └── token
    ├── lab1
    │   ├── assignment.yml
    │   ├── git_questions_ag_test.go
    │   ├── main_ag_test.go
    │   ├── missing_semester_ag_test.go
    │   ├── shell_questions_ag_test.go
    │   ├── token_ag_test.go
    │   └── tokengen
    ├── lab2
    │   ├── assignment.yml
    │   ├── c_questions_ag_test.go
    │   ├── hello_world_ag_test.go
    │   ├── hello_world_init.go
    │   └── main_ag_test.go
    └── scripts
        ├── Dockerfile
        └── run.sh
```

The DAT320 Dockerfile

```
FROM golang:1.23-alpine

# Install bash and git (and build-base to get gcc)
# (this is required when building FROM: golang:alpine)
RUN apk update && apk add --no-cache git bash build-base docker openrc
RUN wget -O- -nv https://raw.githubusercontent.com/golangci/golangci-lint/master/install.sh | sh -s v1.41.1

WORKDIR /quickfeed
```

Script to Run and Score Student Code

The run.sh script

```
#image/dat320

cd "$ASSIGNMENTS/$CURRENT"
find . \( -name '*_test.go' -and -not -name '*_ag_test.go' \) -exec rm -rf {} \;
cp -r "$TESTS"/* "$ASSIGNMENTS"/

go get -t github.com/quickfeed/quickfeed/kit/score
go mod tidy
SCORE_INIT=1 go test -v ./... 2>&1 | grep TestName
```

Script to Run and Score Student Code

The run.sh script

```
#image/dat320

cd "$ASSIGNMENTS/$CURRENT"
find . \( -name '*_test.go' -and -not -name '*_ag_test.go' \) -exec rm -rf {} \;
cp -r "$TESTS"/* "$ASSIGNMENTS"/

go get -t github.com/quickfeed/quickfeed/kit/score
go mod tidy
SCORE_INIT=1 go test -v ./... 2>&1 | grep TestName

cd "$SUBMITTED/$CURRENT"
find . -name '*_test.go' -exec rm -rf {} \;
cp -r "$TESTS"/* "$SUBMITTED"/

go get -t github.com/quickfeed/quickfeed/kit/score
go mod tidy

go test -v -timeout 30s ./... 2>&1
```

Connect to Docker Daemon

```
// NewDockerCI returns a runner to run CI tests.
func NewDockerCI(logger *zap.SugaredLogger) (*Docker, error) {
    cli, err := client.NewClientWithOpts(
        client.FromEnv,
        client.WithAPIVersionNegotiation(),
    )
    if err != nil {
        return nil, err
    }
    return &Docker{
        client: cli,
        logger: logger,
    }, nil
}
```

```
// Run implements the CI interface. This method blocks until the job has been
// completed or an error occurs, e.g., the context times out.
func (d *Docker) Run(ctx context.Context, job *Job) (string, error) {
    if d.client == nil {
        return "", fmt.Errorf("cannot run job: %s; docker client not initialized", job.Name)
    }

    resp, err := d.createImage(ctx, job)
    if err != nil {
        return "", err
    }

    d.logger.Infof("Created container image '%s' for %s", job.Image, job.Name)
    if err = d.client.ContainerStart(ctx, resp.ID, container.StartOptions{}); err != nil {
        return "", err
    }
}
```

Wait for Container and Get Logs

```
d.logger.Infof("Waiting for container image '%s' for %s", job.Image, job.Name)
msg, err := d.waitForContainer(ctx, job, resp.ID)
if err != nil {
    return msg, err
}

d.logger.Infof("Done waiting for container image '%s' for %s", job.Image, job.Name)
// extract the logs before removing the container below
logReader, err := d.client.ContainerLogs(ctx, resp.ID, container.LogsOptions{
    ShowStdout: true,
})
if err != nil {
    return "", err
}
```

```
// Job describes how to execute a CI job.
type Job struct {
    // Name describes the running job; mainly used to name docker containers.
    Name string
    // Image names the image to use to run the job.
    Image string
    // Dockerfile contents.
    // If empty, the image is assumed to exist.
    // If non-empty, the image is built from this Dockerfile.
    Dockerfile string
    // BindDir is the directory to bind to the container's /quickfeed directory.
    BindDir string
    // Env is a list of environment variables to set for the job.
    Env []string
    // Commands is a list of shell commands to run as part of the job.
    Commands []string
}

// Runner contains methods for running user provided code in isolation.
type Runner interface {
    // Run should synchronously execute the described job and return the output.
    Run(context.Context, *Job) (string, error)
}
```

```
var hostConfig *container.HostConfig
if job.BindDir != "" {
    hostConfig = &container.HostConfig{
        Mounts: []mount.Mount{
            {
                Type:  mount.TypeBind,
                Source: job.BindDir,
                Target: QuickFeedPath,
            },
        },
    }
}

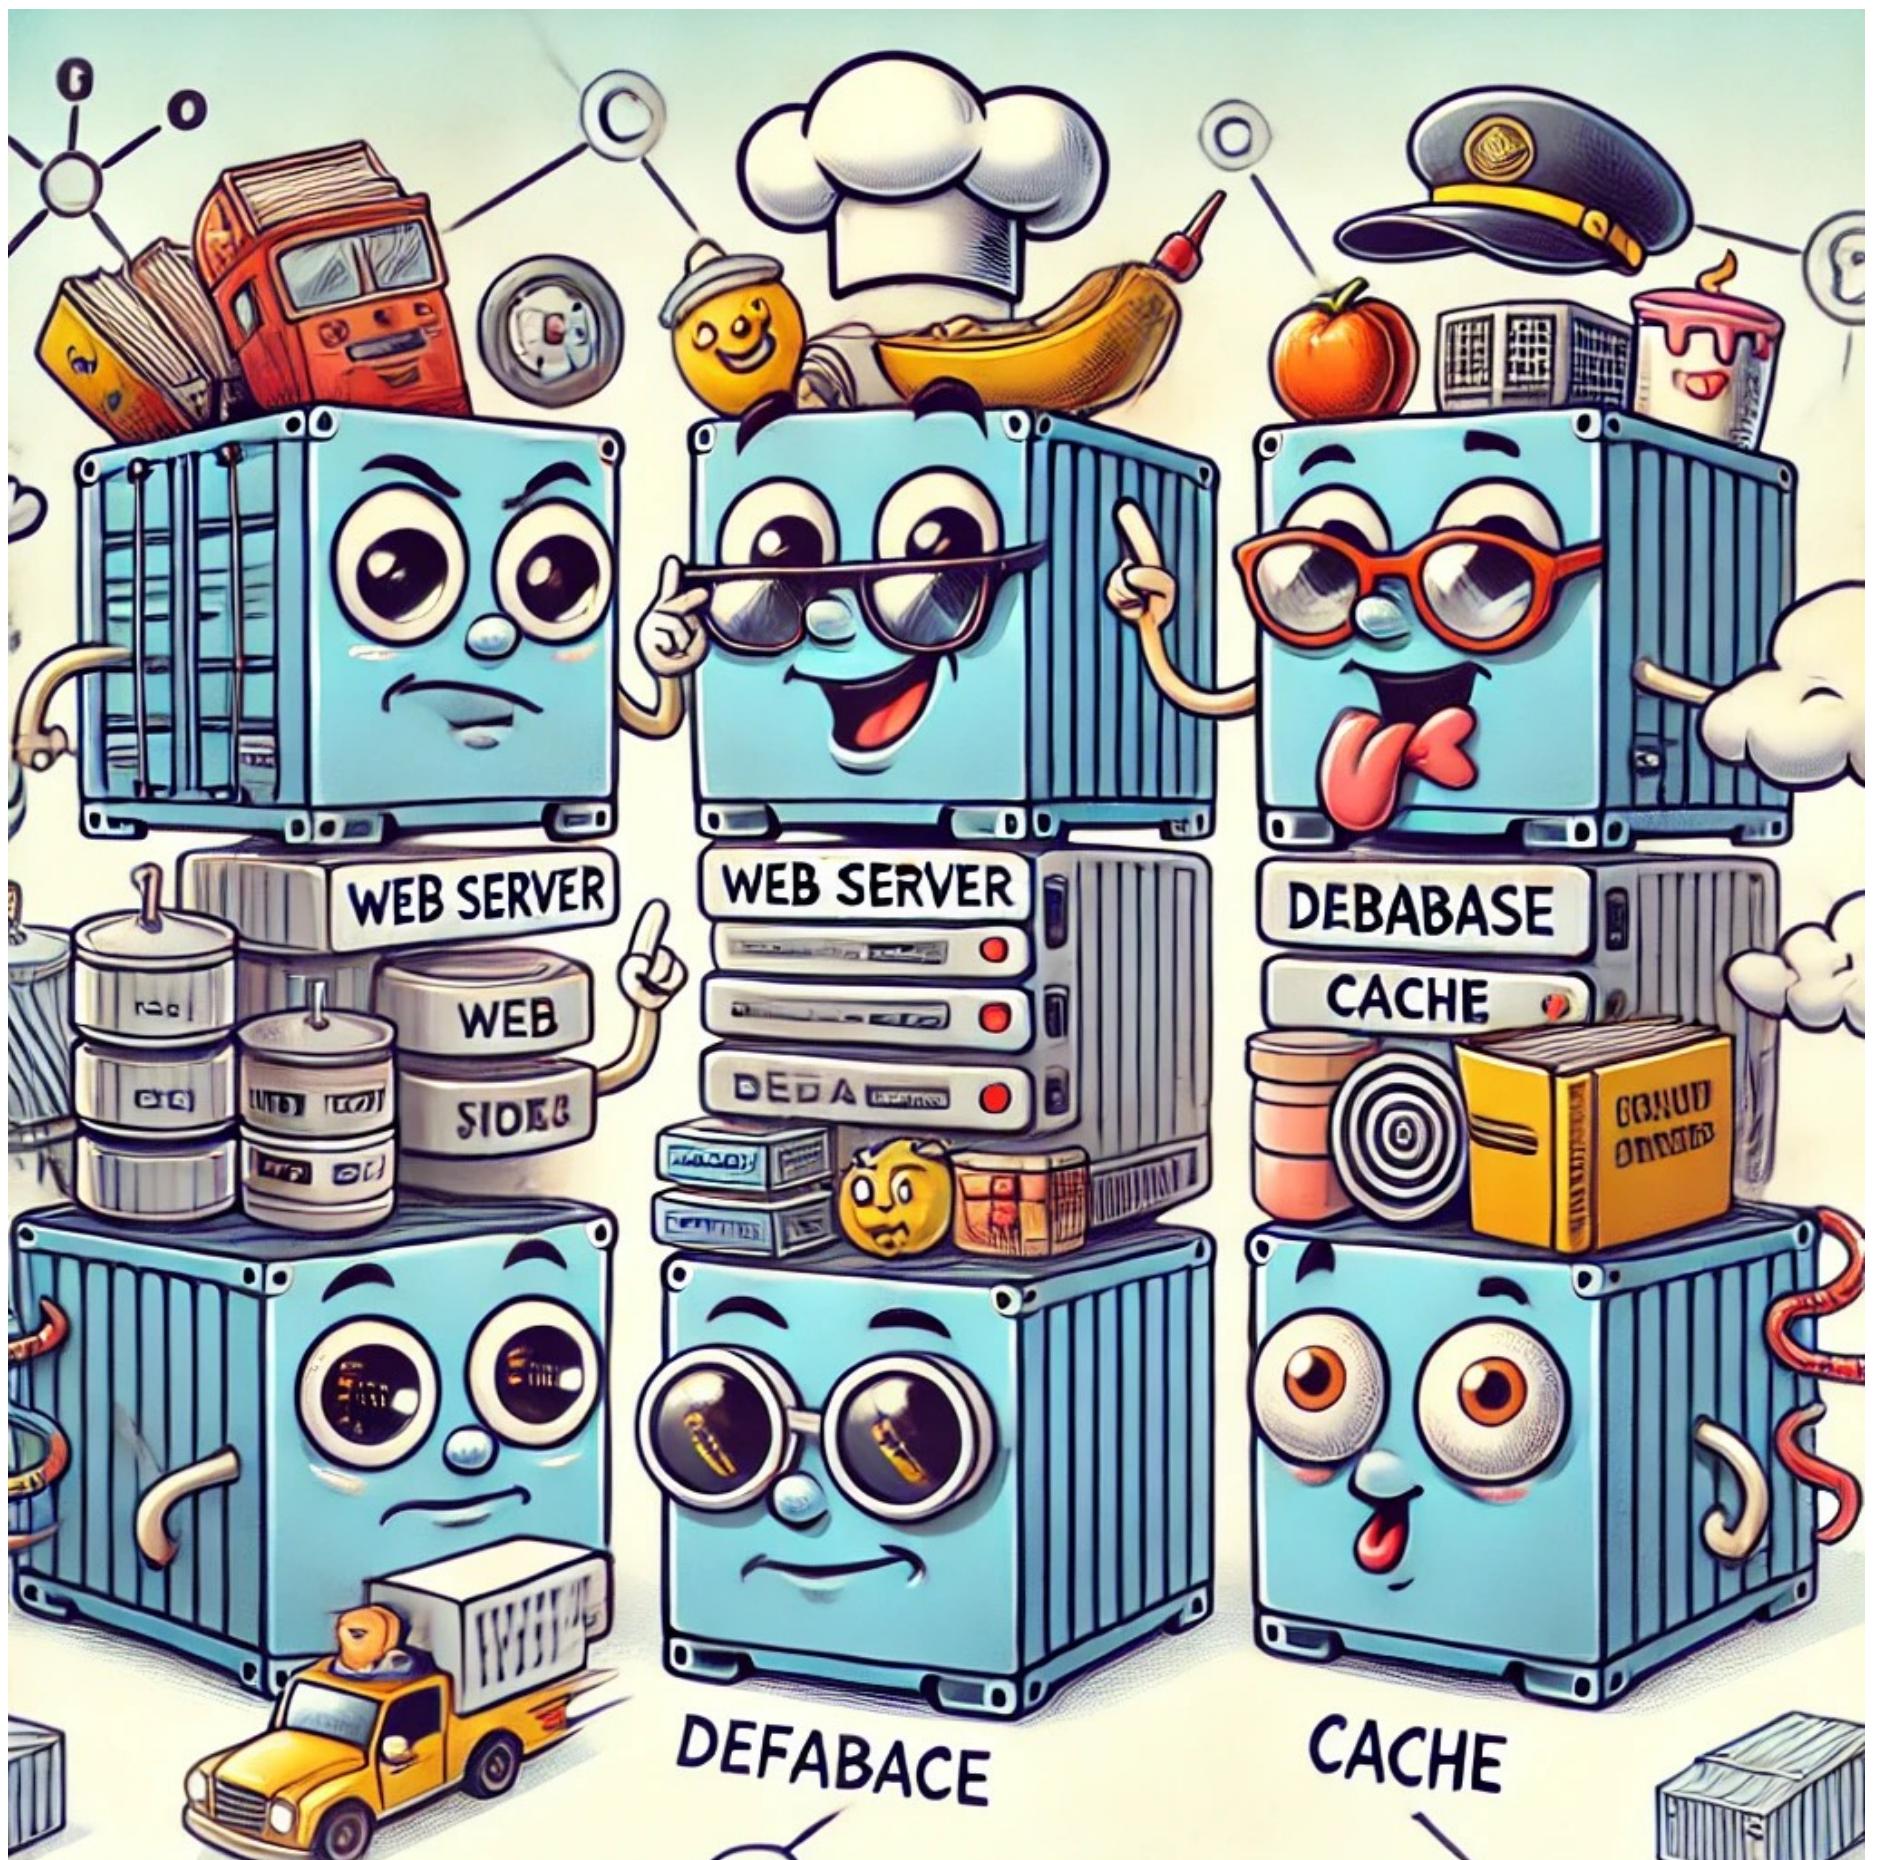
create := func() (container.CreateResponse, error) {
    return d.client.ContainerCreate(ctx, &container.Config{
        Image: job.Image,
        User:  fmt.Sprintf("%d:%d", os.Getuid(), os.Getgid()), // Run the image as the current user, e.g., quickfeed
        Env:   job.Env,                                         // Set default environment variables
        Cmd:   []string{"/bin/bash", "-c", strings.Join(job.Commands, "\n")},
    }, hostConfig, nil, nil, job.Name)
}

resp, err := create()
```

Multi-Container Applications

Microservice Architecture

- **Single Responsibility per Container**
 - Run only one service to ensure modularity and easier management
- **Inter-Container Communication**
 - Containers communicate with each other over a Docker network
 - Internal DNS to resolve service names to container IPs, making services discoverable

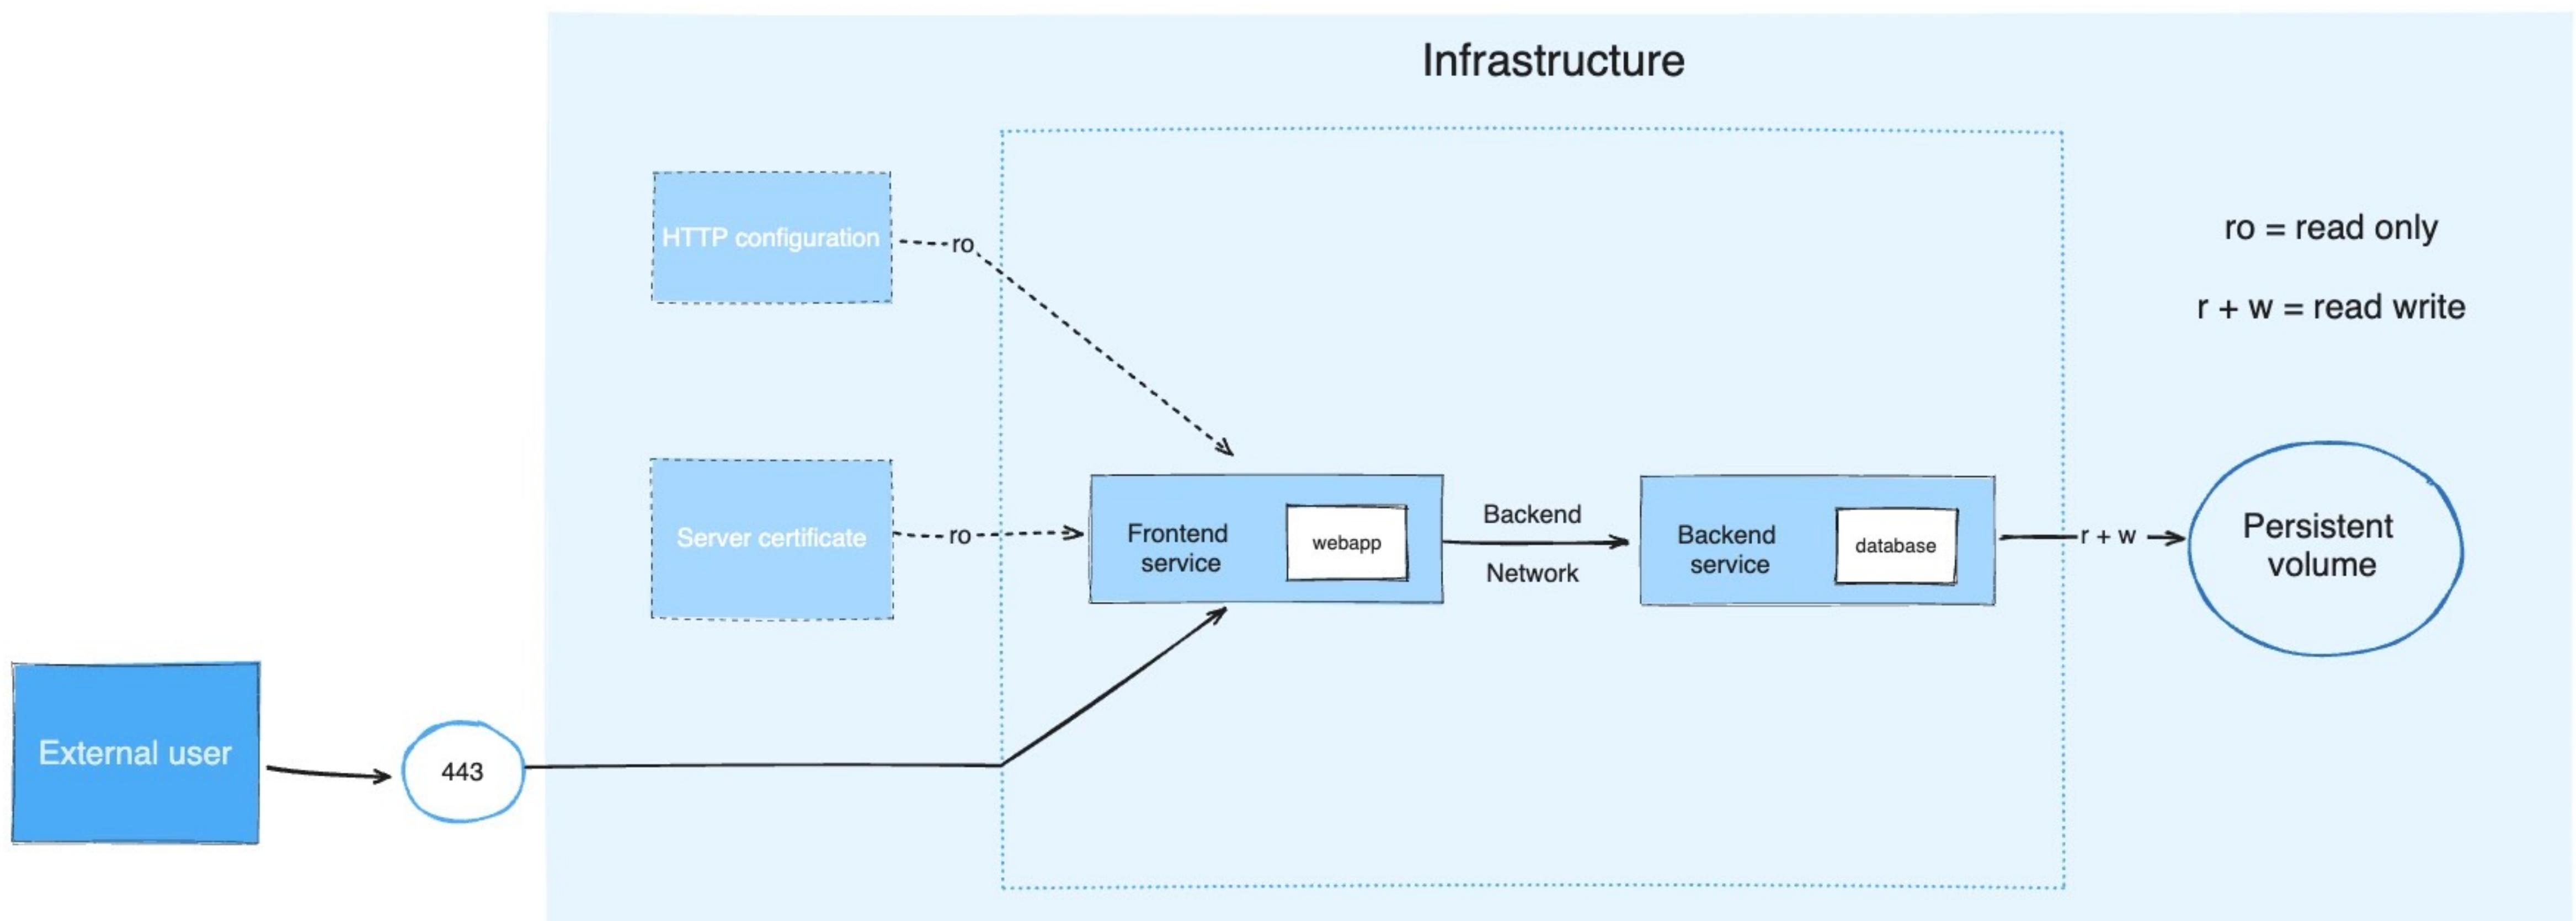


Microservice Architecture

- **Loose Coupling**
 - Isolating services into different containers
 - Updates can be made to one service without affecting other containers
 - More maintainable and resilient applications
- **Scaling and Resilience**
 - Individual services can be scaled independently based on load
 - Load balancer, multiple web servers, a single database

- Docker Compose is a tool that simplifies managing multi-container applications
- A `compose.yml` file to define and run multiple containers
 - As a single service stack,
 - handling container creation,
 - networking, and
 - volume management

Docker Compose Example



Docker Compose Example

```
services:  
  frontend:  
    image: example/webapp  
    ports:  
      - "443:8043"  
    networks:  
      - front-tier  
      - back-tier  
    configs:  
      - httpd-config  
    secrets:  
      - server-certificate  
  
  backend:  
    image: example/database  
    volumes:  
      - db-data:/etc/data  
    networks:  
      - back-tier  
  
volumes:  
  db-data: {}  
  
configs:  
  httpd-config:  
    external: true  
  
secrets:  
  server-certificate:  
    external: true  
  
networks:  
  front-tier: {}  
  back-tier: {}
```

Separation of Secret Management from Container Spec

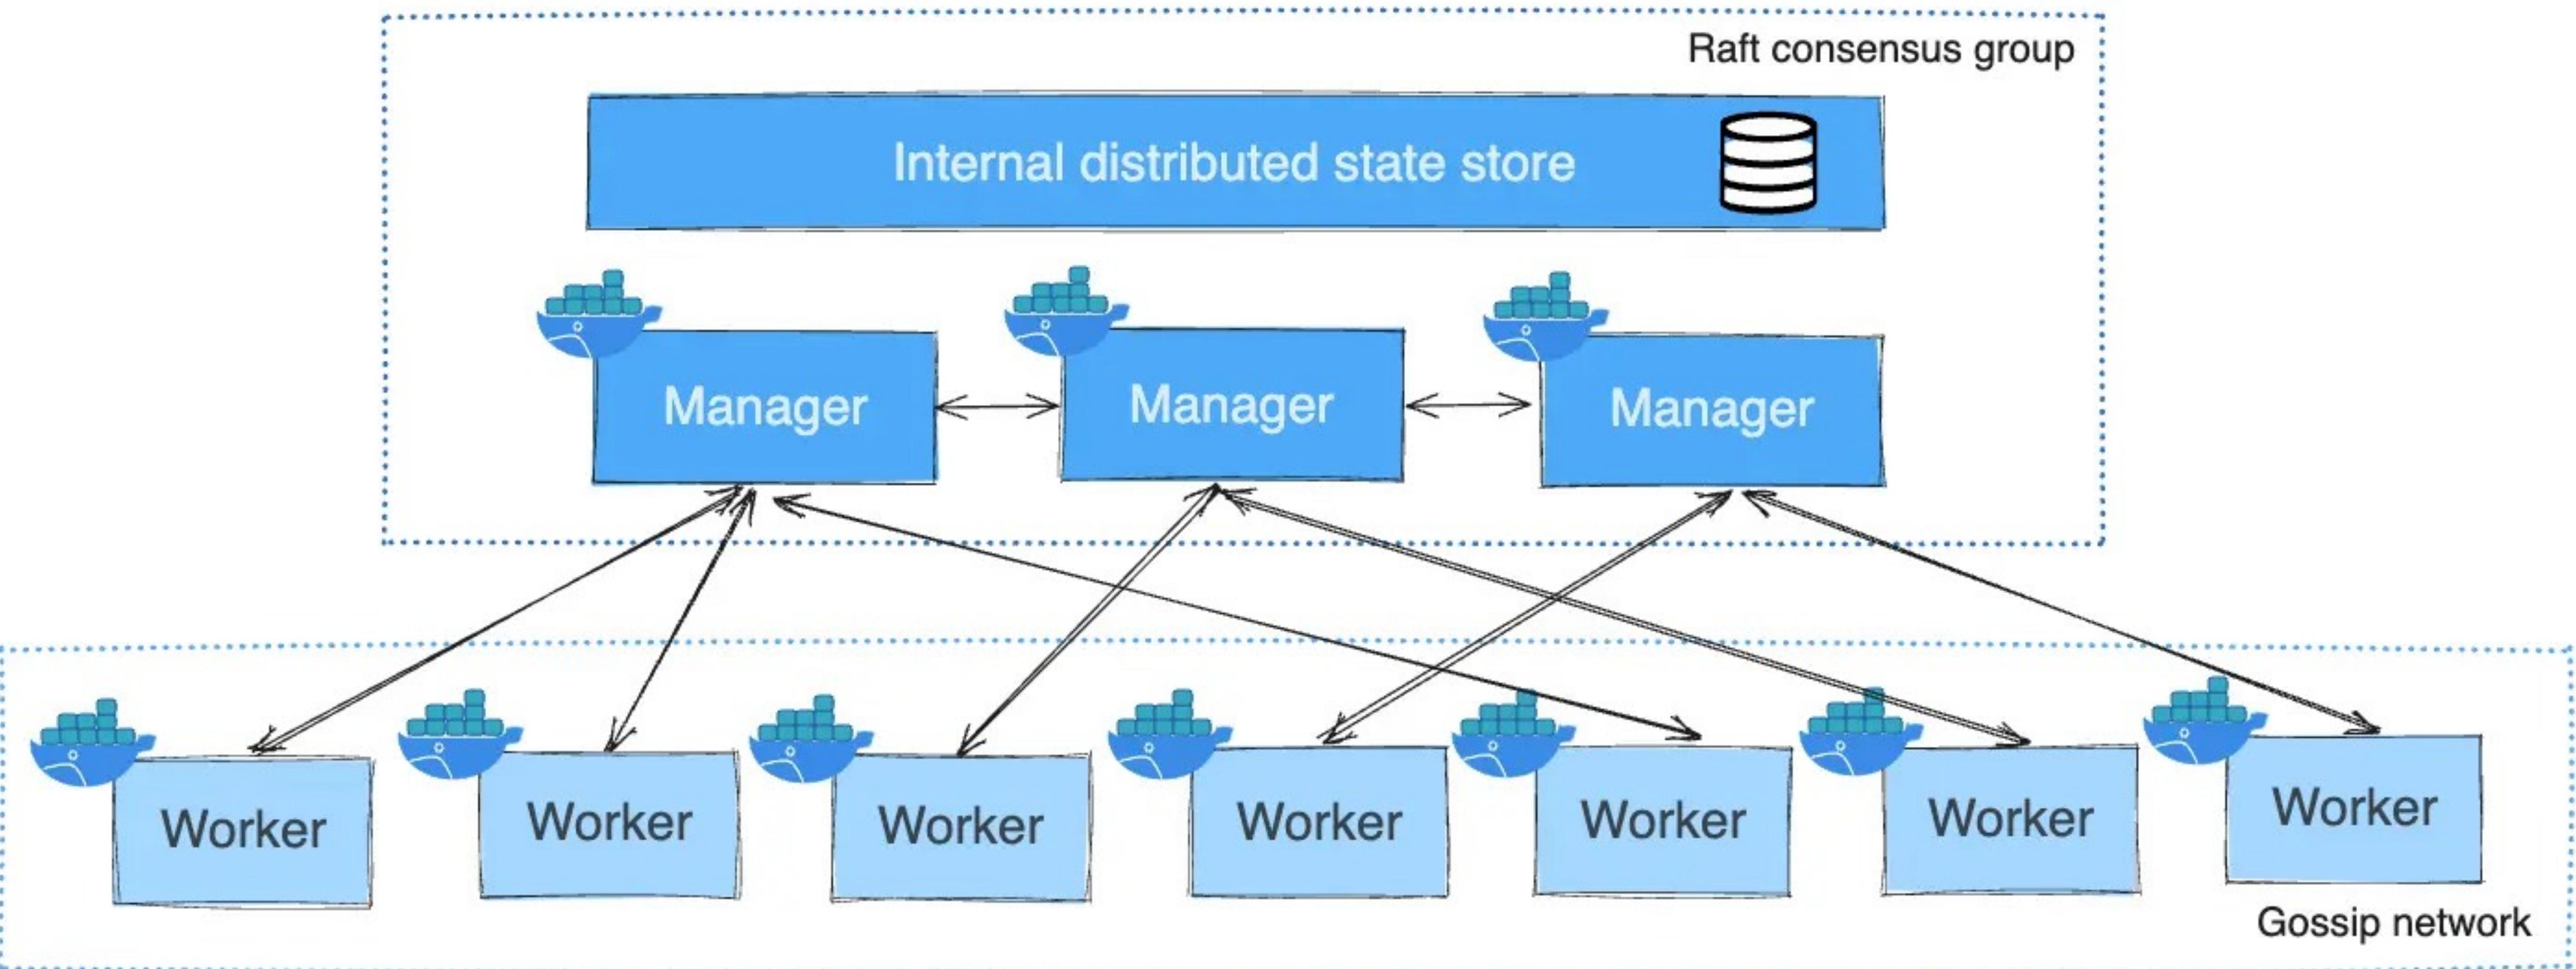
- Avoid hardcoding into the compose file
- Make the application more secure and flexible
- Secret resources managed externally via Docker Swarm

```
docker swarm init
docker secret create server-certificate /path/to/server-certificate.pem
docker secret ls
docker-compose up -d
```

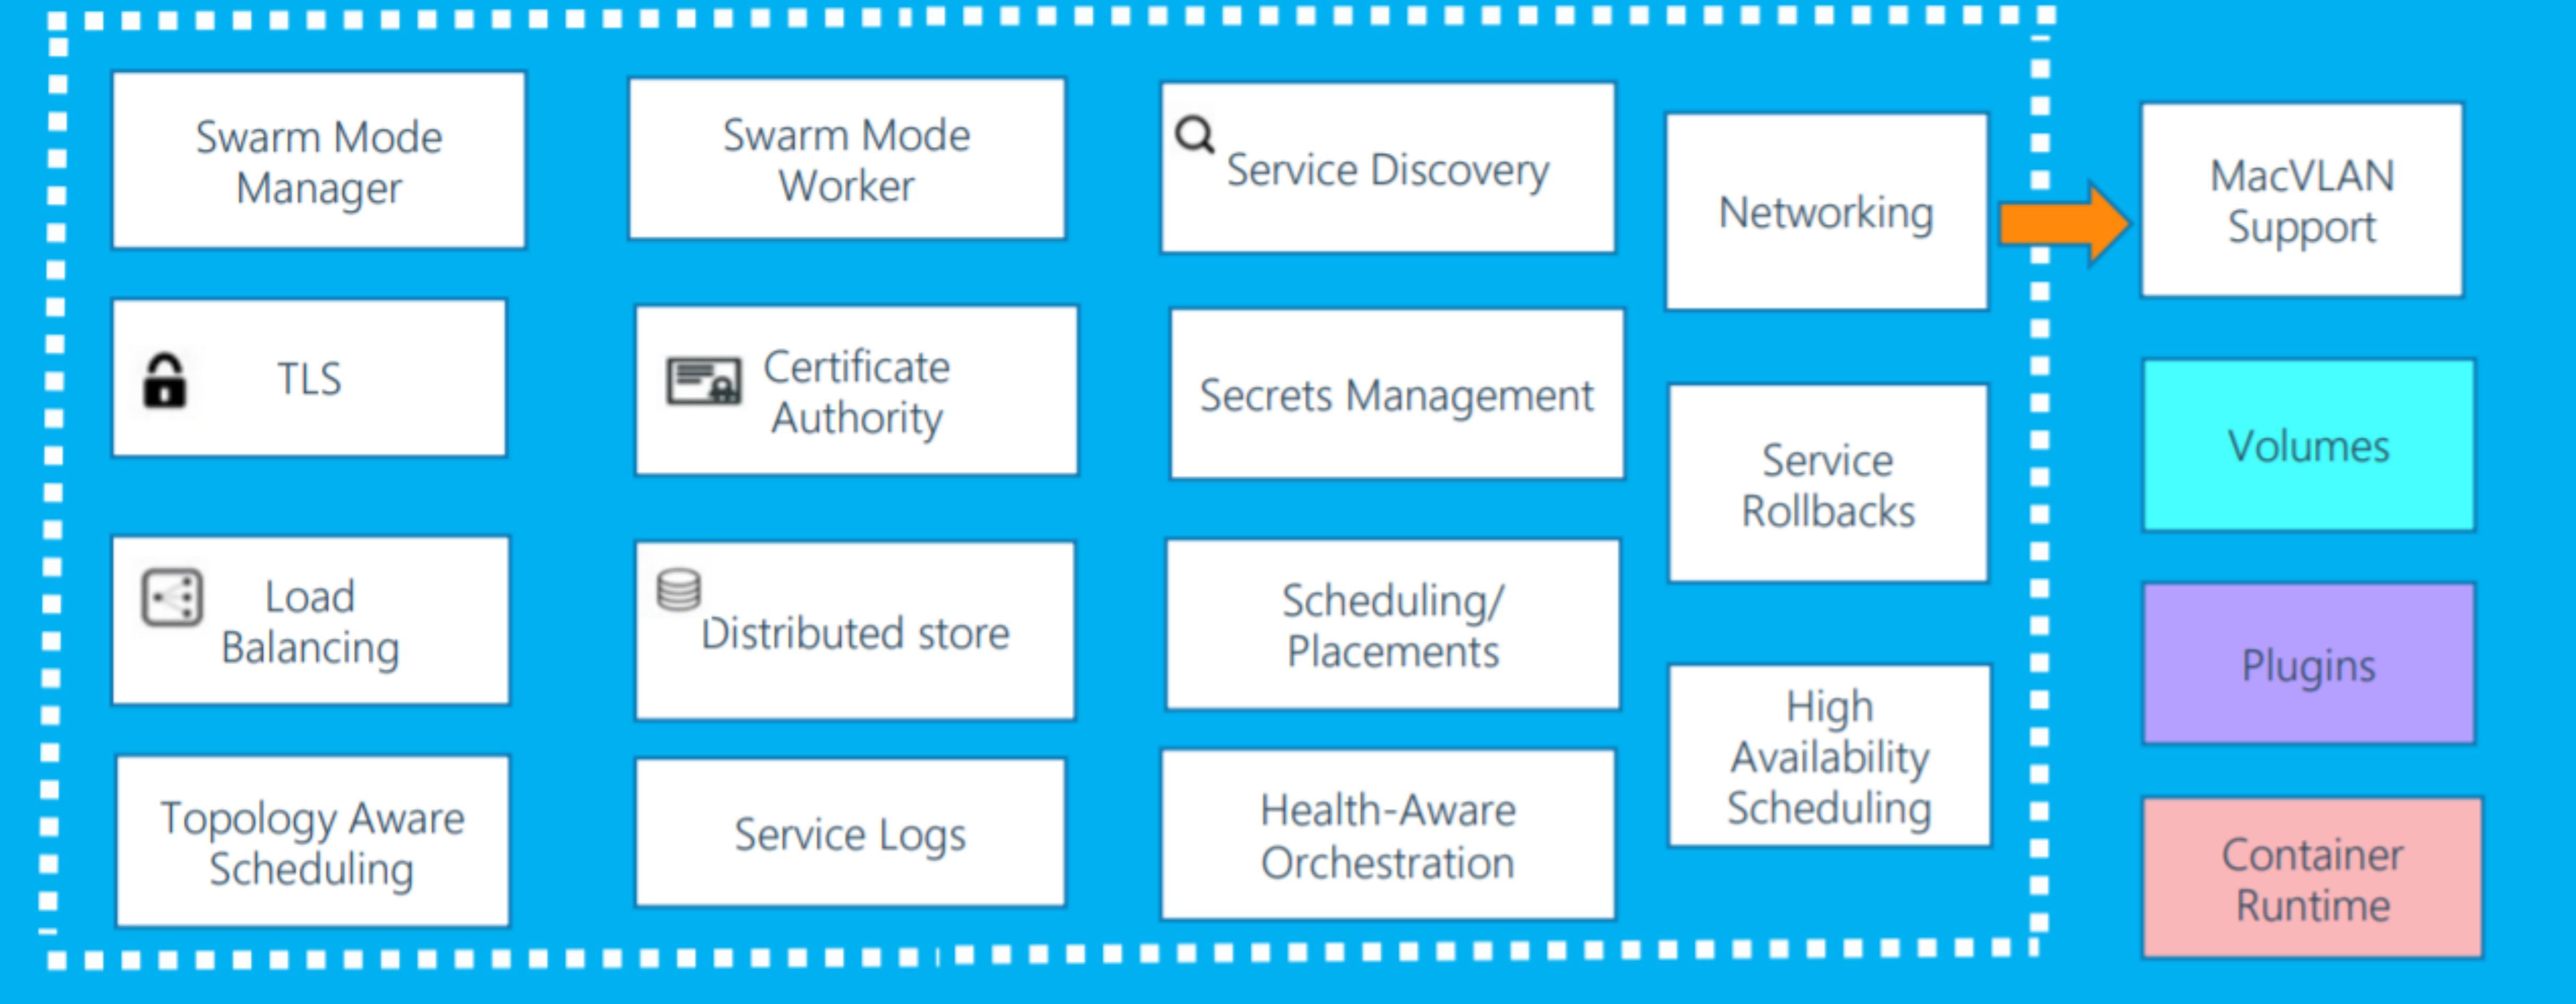
Docker Container Orchestration

Container Orchestration Services

- **Scaling**
 - Automatically scale containers across multiple nodes in a Docker Swarm
- **Load Balancing**
 - Distribute incoming requests evenly across containers within the service
- **Declarative Management**
 - Define desired state of services
 - Docker ensures the correct number of containers are running



Orchestration Components



Questions?