

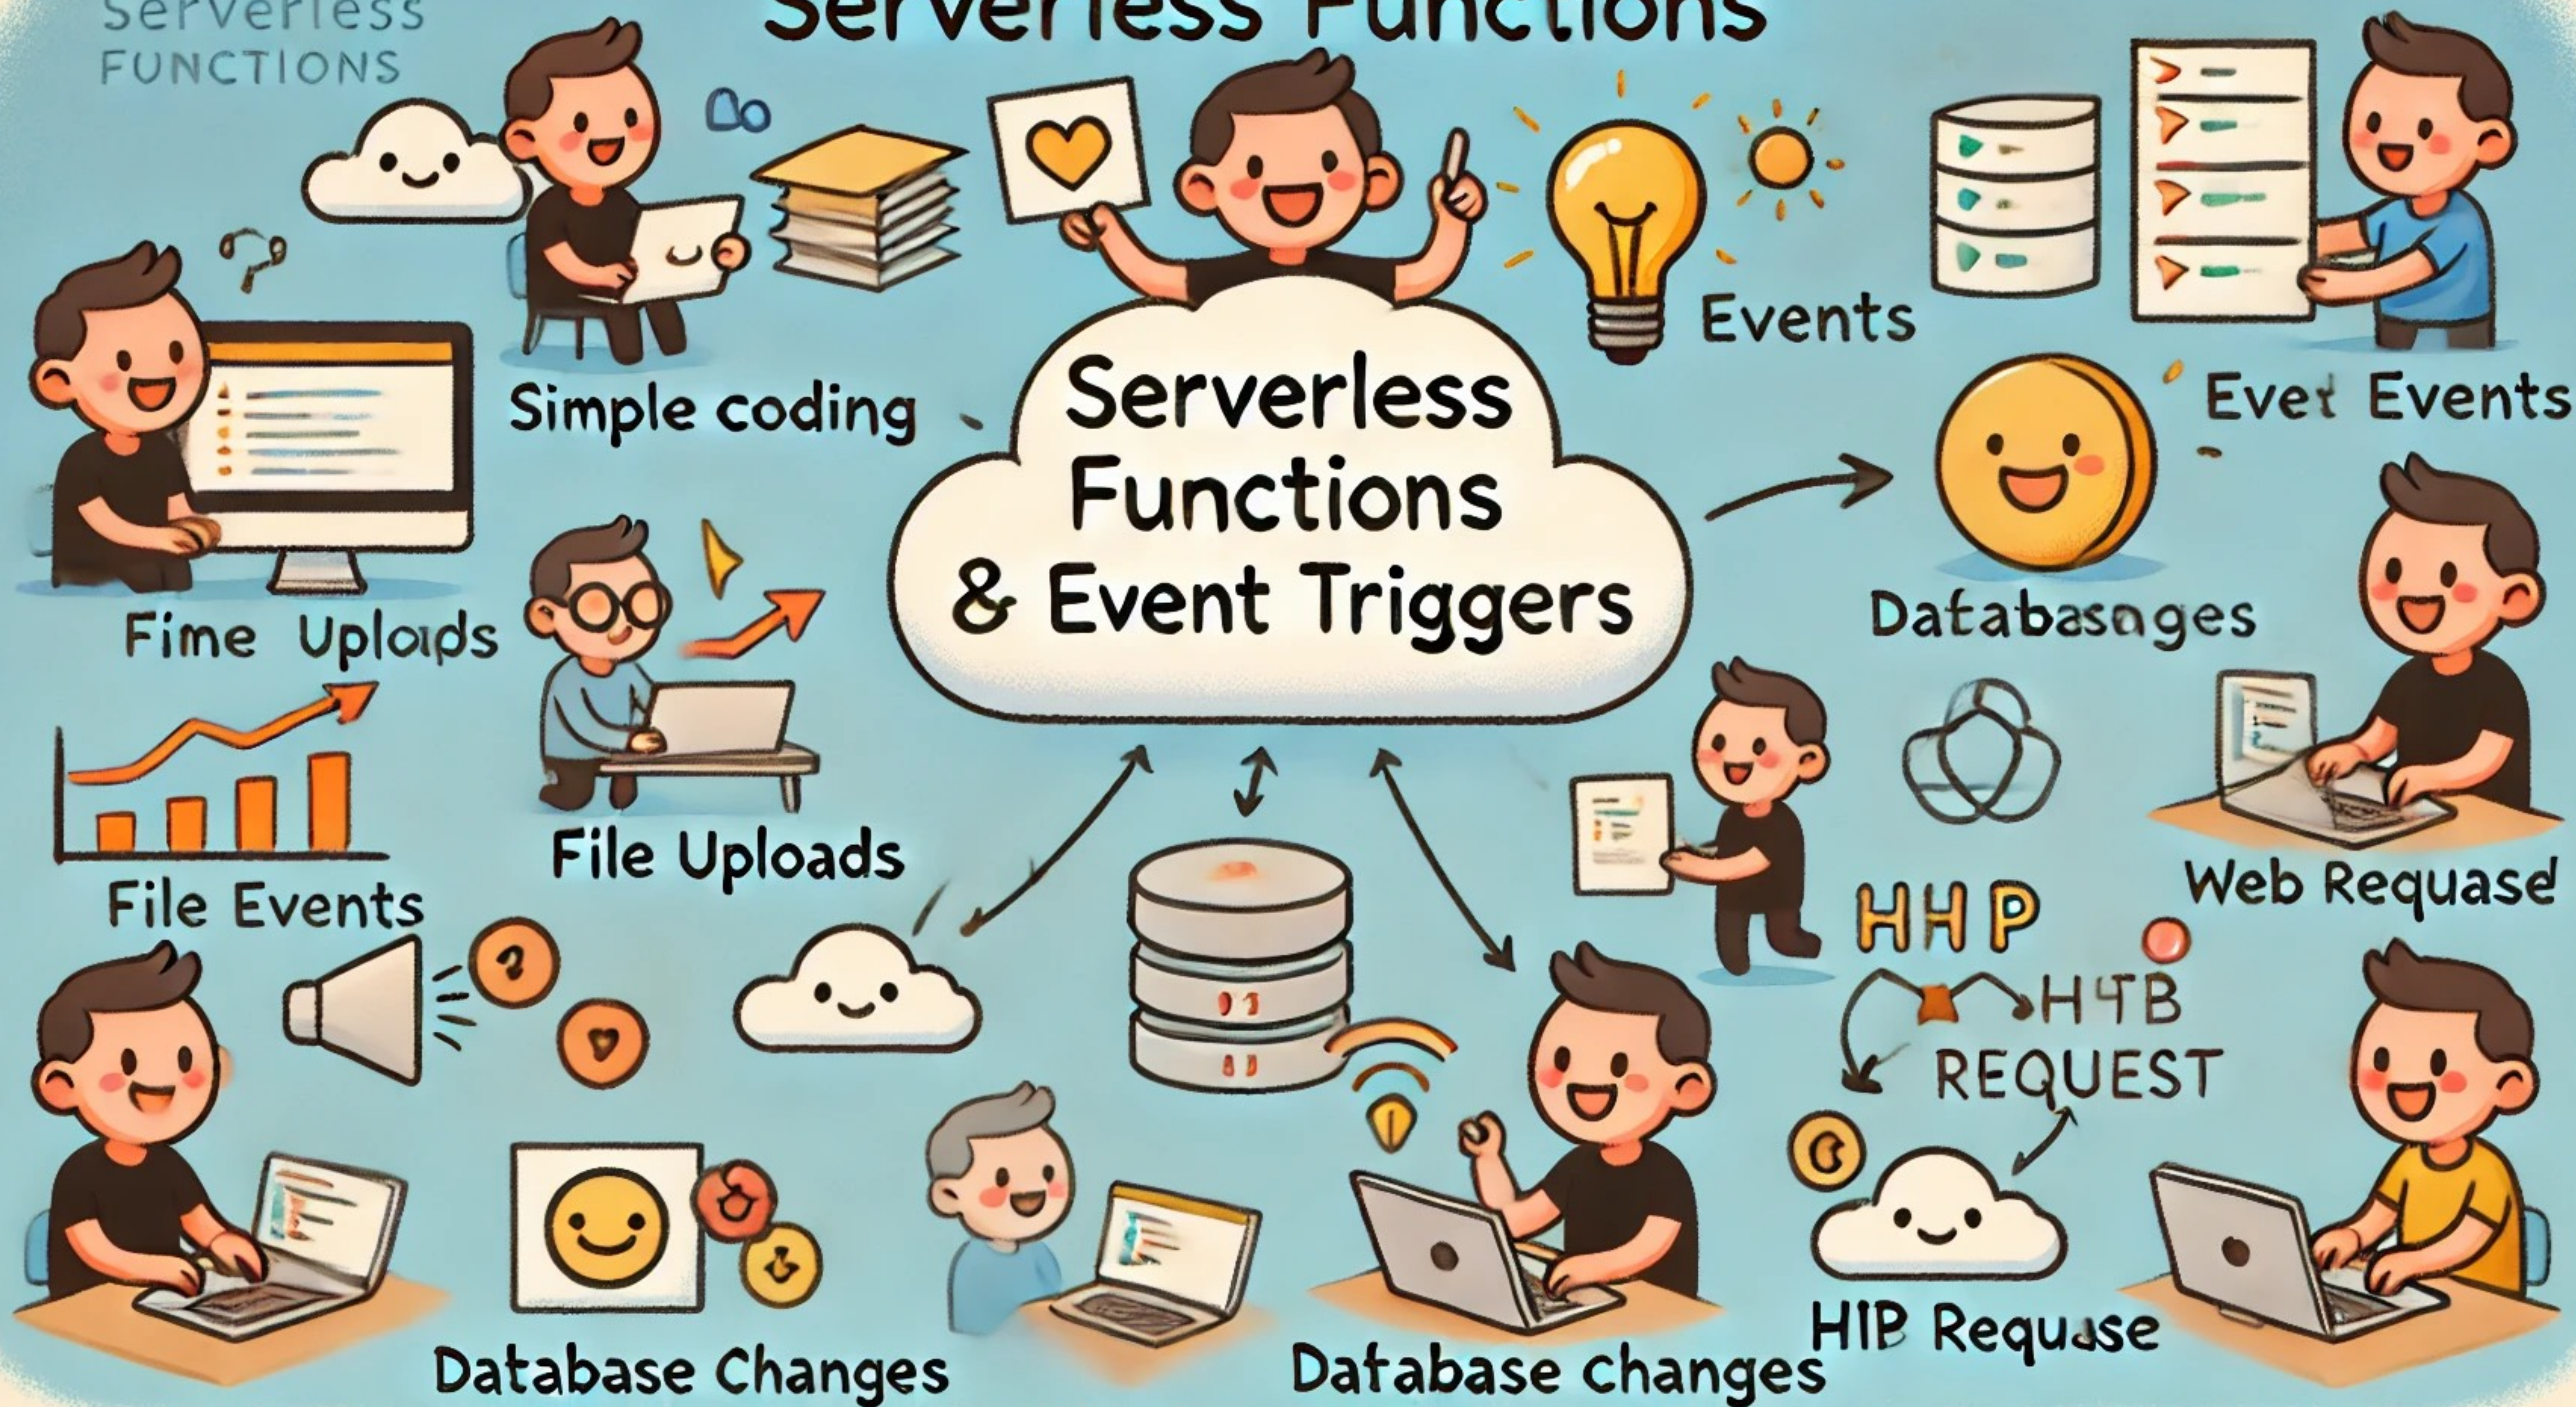
Cloud Computing Technologies

DAT515 - Fall 2024

Serverless Functions

Prof. Hein Meling

Serverless FUNCTIONS




What are Serverless Functions?

Definition of the Serverless Model

- Functions-as-a-Service (FaaS)
 - Specify **function** definition and **events** that triggers it
- Automate
 - **Provisioning** and **scaling** of machines
 - **Distribution** of functions over machines

Benefits of the Serverless Model

- Reduces development effort
- Provides elastic scaling
- Cost-effective



| On-site | IaaS | PaaS | SaaS |
|------------------|------------------|------------------|------------------|
| Applications | Applications | Applications | Applications |
| Data | Data | Data | Data |
| Runtime | Runtime | Runtime | Runtime |
| Middleware | Middleware | Middleware | Middleware |
| Operating System | Operating System | Operating System | Operating System |
| Virtualization | Virtualization | Virtualization | Virtualization |
| Servers | Servers | Servers | Servers |
| Storage | Storage | Storage | Storage |
| Networking | Networking | Networking | Networking |

Service Model Comparison

| On-site | IaaS | PaaS | FaaS | SaaS |
|------------------|------------------|------------------|------------------|------------------|
| Functions | Functions | Functions | Functions | Functions |
| Applications | Applications | Applications | Applications | Applications |
| Data | Data | Data | Data | Data |
| Runtime | Runtime | Runtime | Runtime | Runtime |
| Middleware | Middleware | Middleware | Middleware | Middleware |
| Operating System | Operating System | Operating System | Operating System | Operating System |
| Virtualization | Virtualization | Virtualization | Virtualization | Virtualization |
| Servers | Servers | Servers | Servers | Servers |
| Storage | Storage | Storage | Storage | Storage |
| Networking | Networking | Networking | Networking | Networking |

Example Serverless Function and Trigger

```
import azure.functions as func
import urllib.request

def main(req: func.HttpRequest):
    url = req.get_body().decode("utf-8")
    fid = urllib.request.urlopen(url)
    webpage = fid.read().decode('utf-8')
    found_free = webpage.find("free") >= 0
    return str(found_free)
```

```
"bindings": [ { "authLevel": "anonymous",
                  "name": "req",
                  "type": "httpTrigger",
                  "direction": "in",
                  "route": "checkWebPage",
                  "methods": [ "get" ] },
                { "name": "$return",
                  "type": "http",
                  "direction": "out" } ]
```

- Can scale out automatically to handle 100s of kops per second
- Operate cheaply under low load — due to load-based billing

Ideal Use Cases

- Compute-intensive highly parallelizable workloads
- Workflow processing
- Event-driven applications
- Microservices
- Real-time data processing

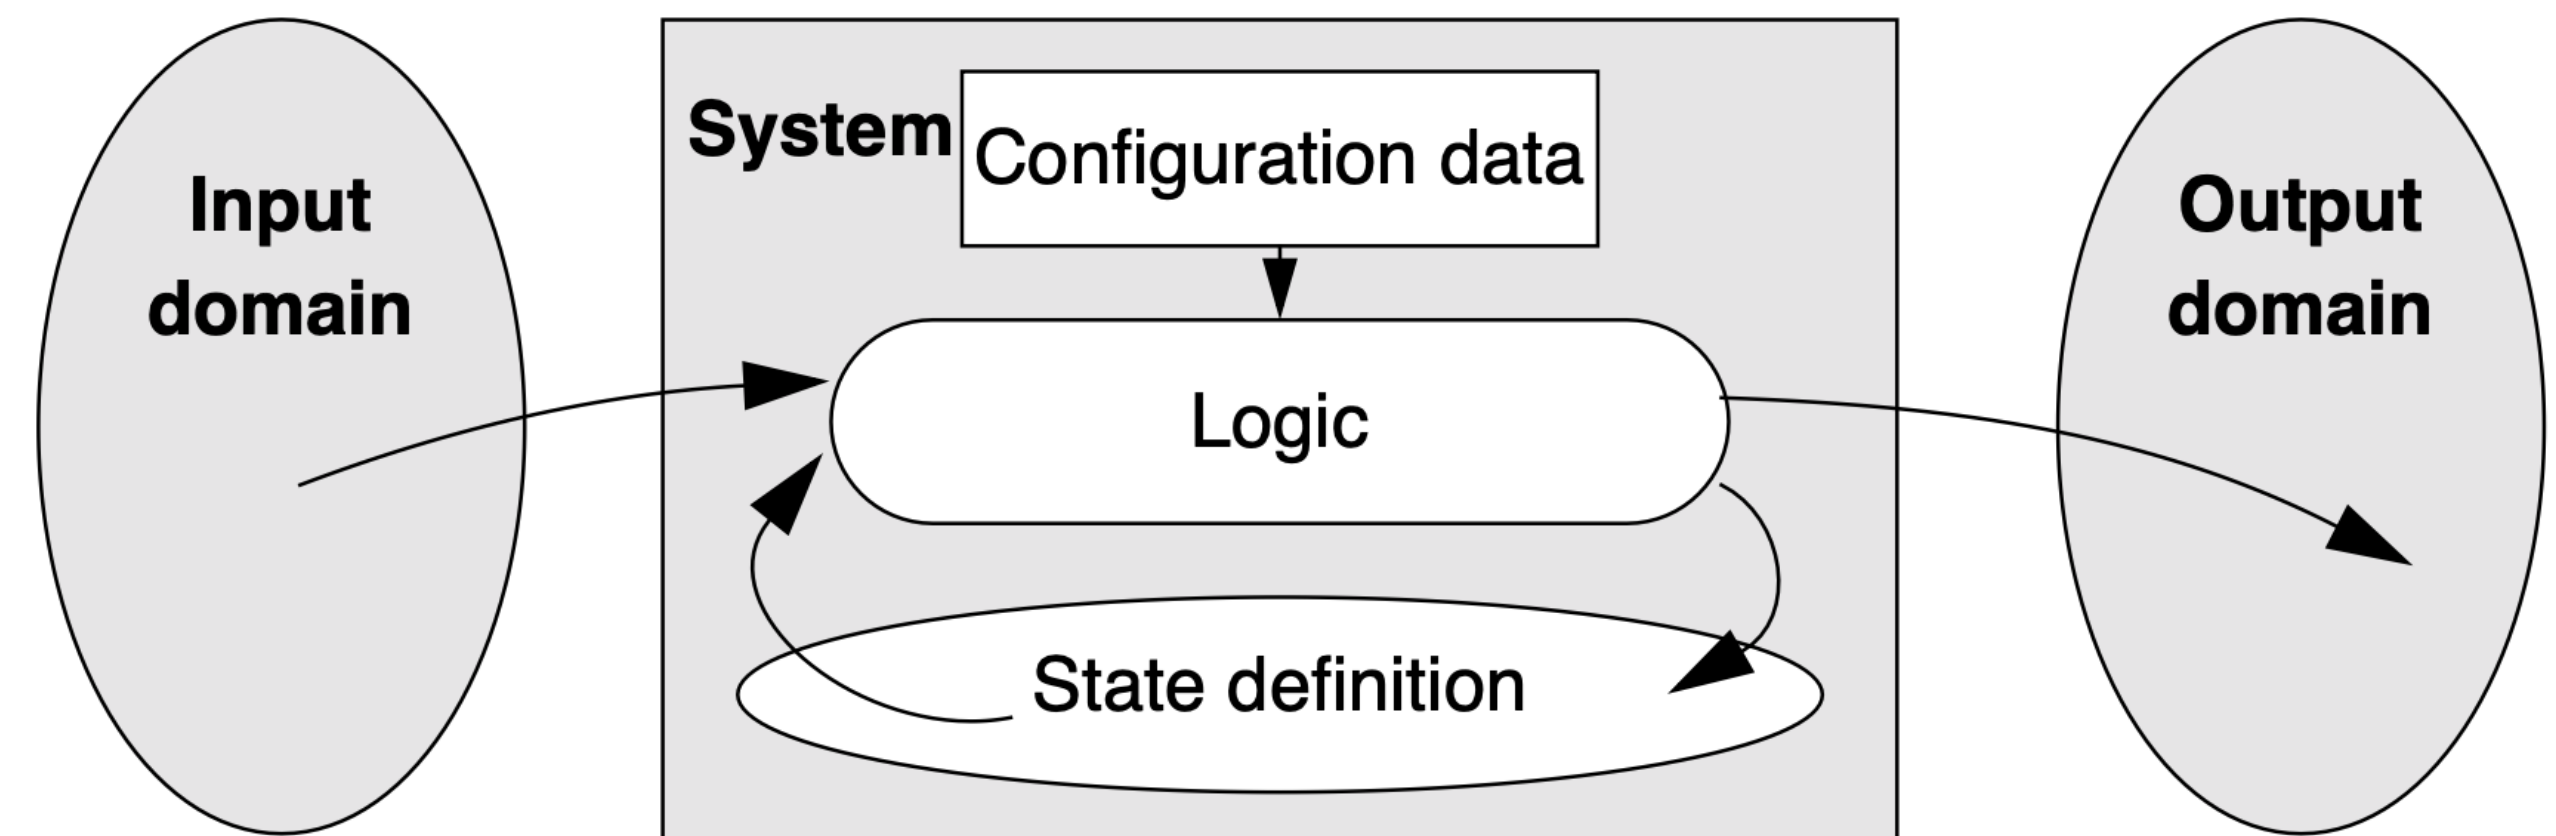

Not Ideal For

- Long-running processes
- High performance (over time)

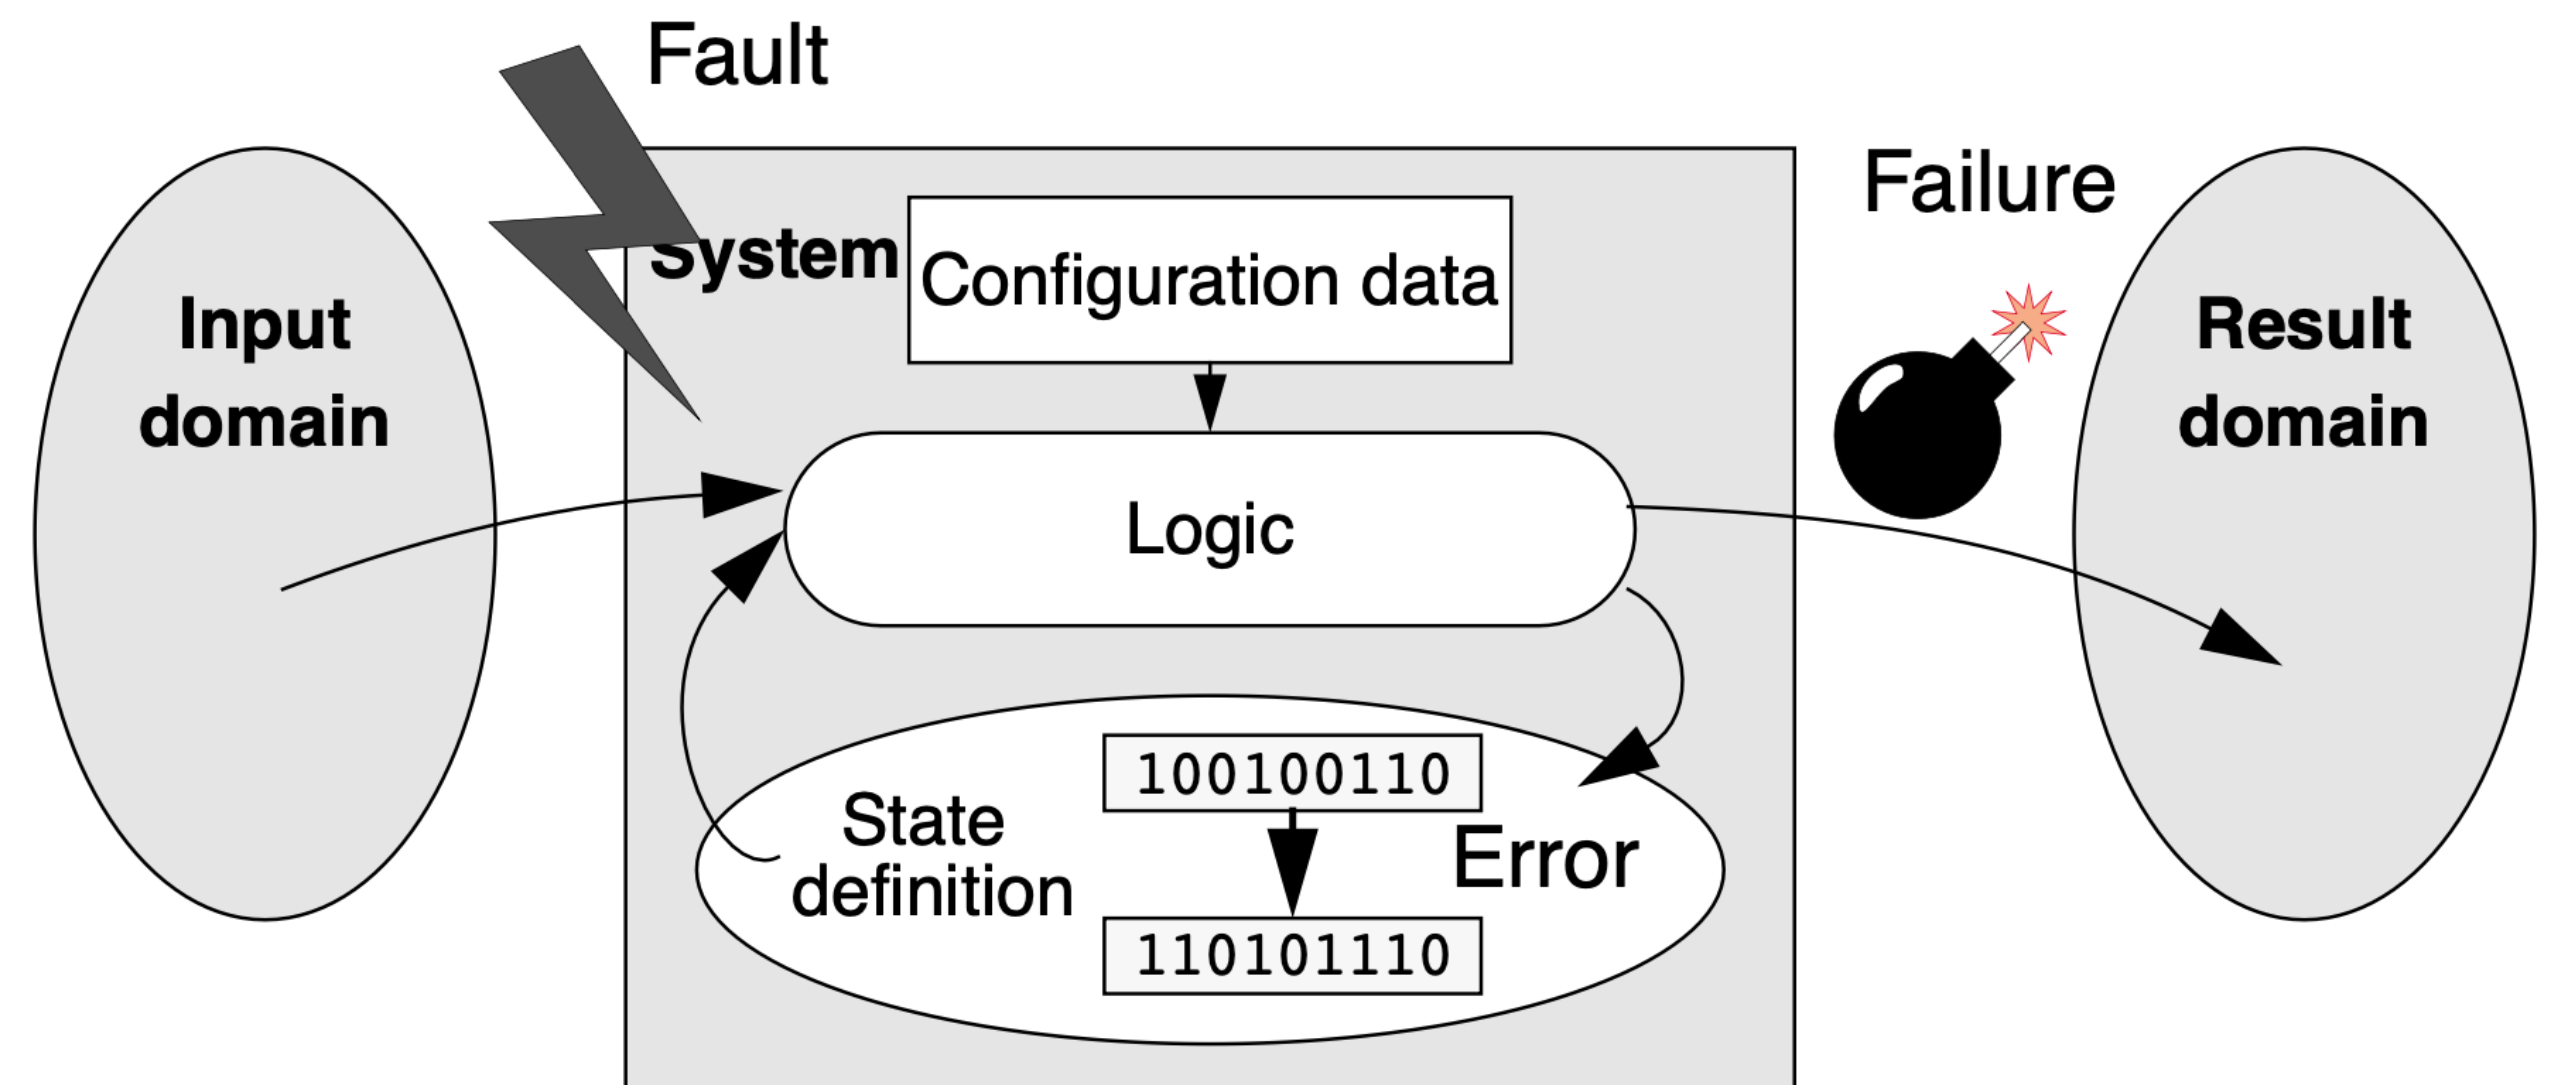

Decision Factors

- Scalability needs
- Cost considerations
- Development speed and flexibility

Function Machine vs Moore-Mealy Machine

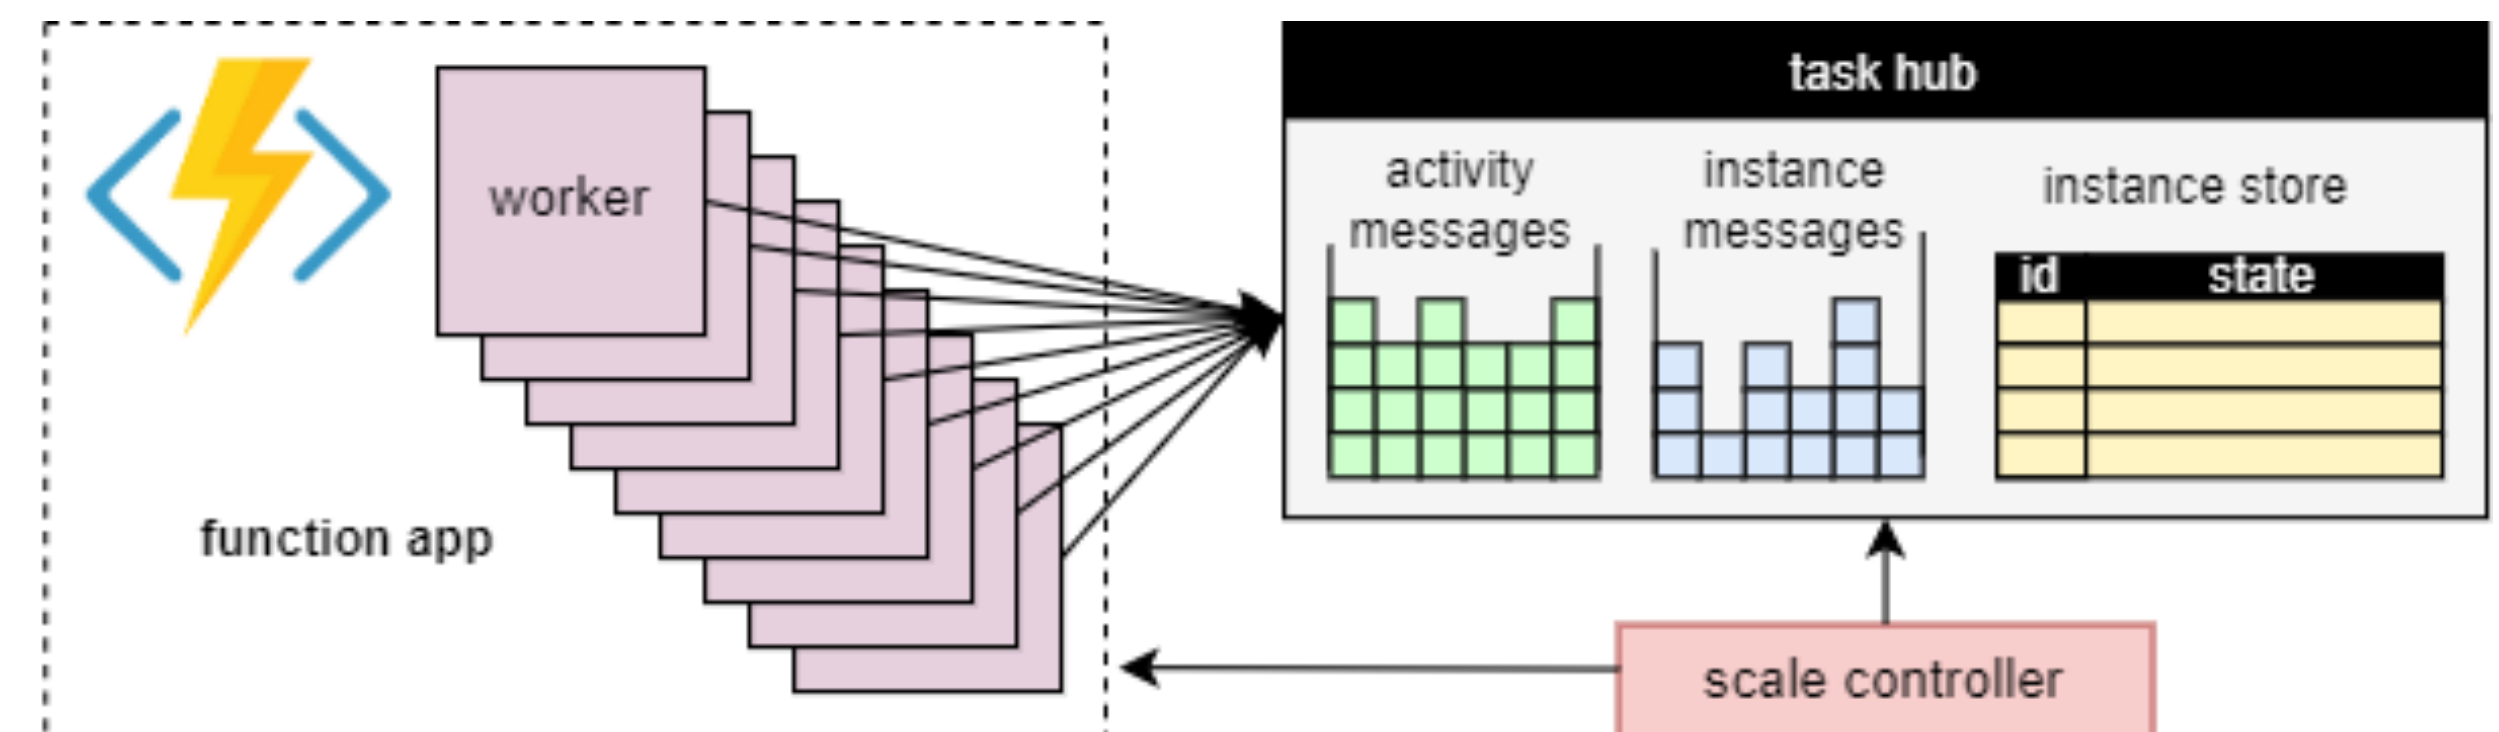


Function Machine vs Moore-Mealy Machine



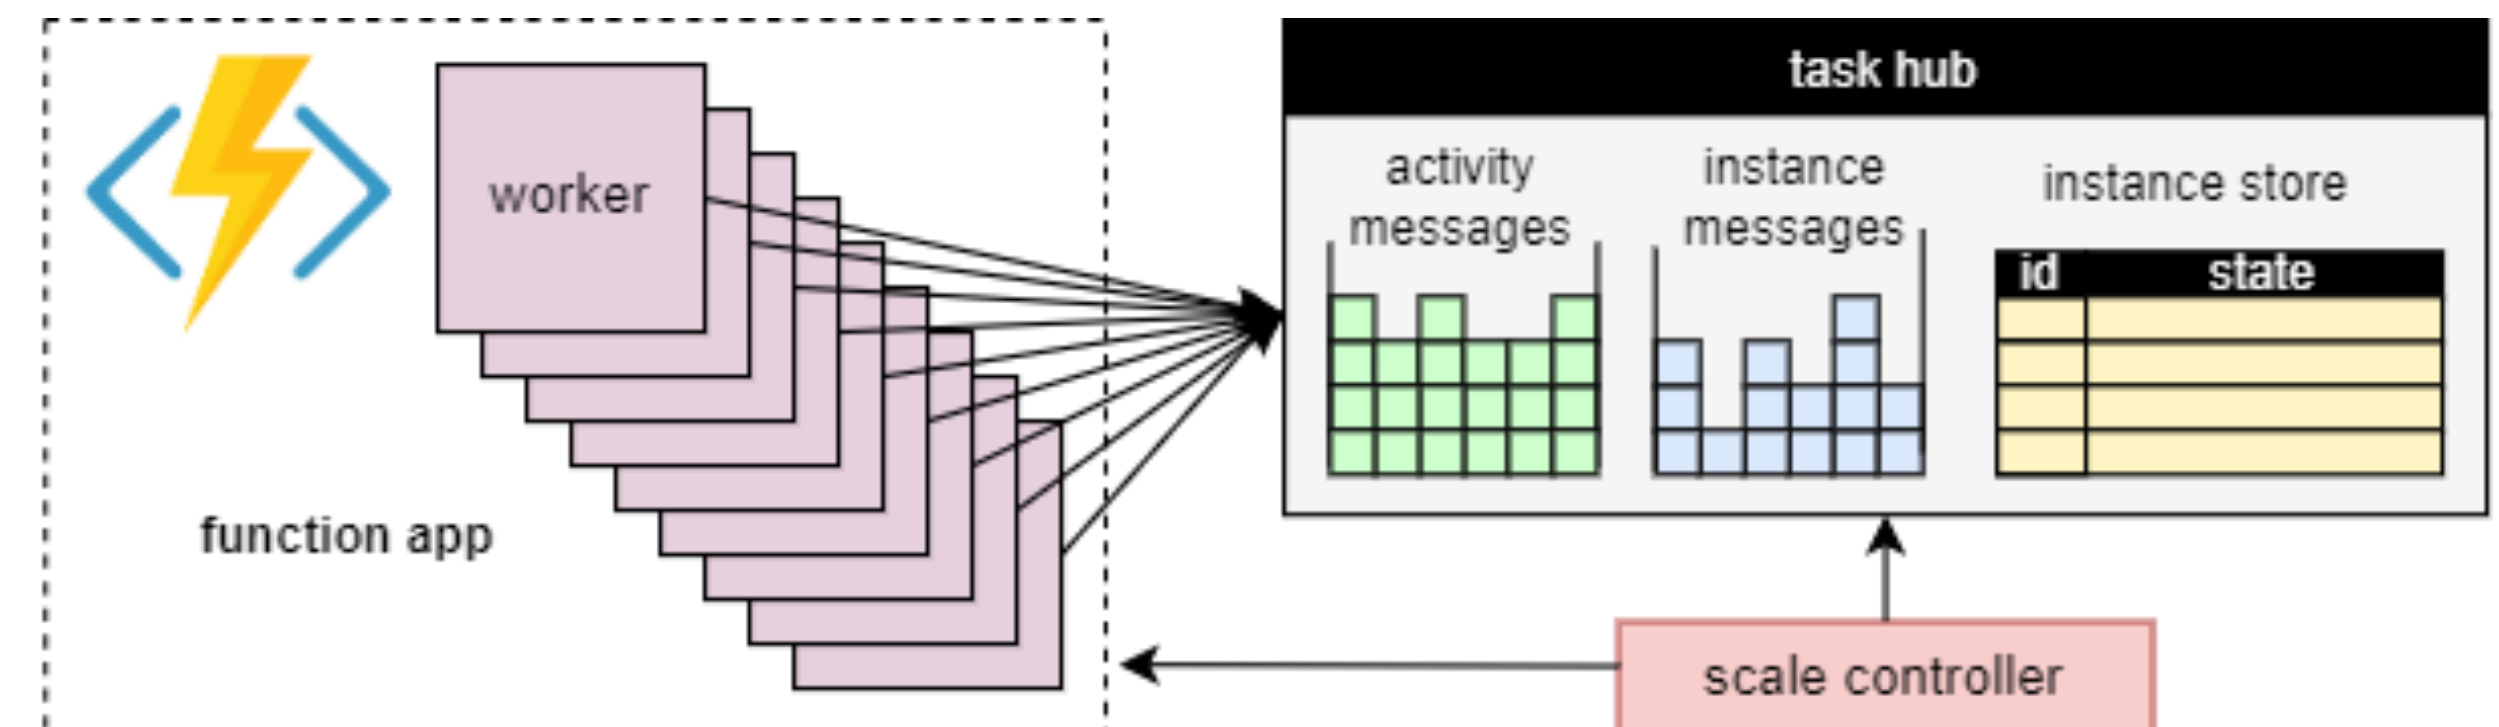

The Compute-Storage Separation

- Is *becoming* a dominant **architectural principle** for cloud services
- Separating the application into **ephemeral workers** and a **storage** service
 - State of the application decoupled from compute
 - Making the state available when a worker crashes (or is shut down due to low load)
- Scale workers independently of storage



The Compute-Storage Separation

- Not limited to stateless input/output computation
- Functions can call **external services**
 - Key-value stores, queues, or databases
- Compute-storage separation can be **cumbersome** for developers



Joke time

Why did the developer go
serverless?

Because she didn't want to
waste cycles on managing
servers!

Serverless Developer Challenges

Cloud Environment is Inherently and Pervasively

- Concurrent
- Parallel
- Distributed
- Failure-prone

Weak Execution Guarantees for Serverless Functions

- **Execution Time Limit**, e.g., 5 minutes
 - Problem for long-running computations that don't finish
- **Partial Execution**
 - Functions are subject to **failures** (OOM error, VM shut down, HW)
- **At-Least-Once Triggers**
 - Triggers designed to **retry** if uncertain if Function completed
 - Single event can result in **multiple Function executions**

Threats to Consistency: Non-Atomic Updates

- Example
 - Function must update two account balances (in multiple steps)
 - Partial execution may update only one account

Threats to Consistency: Concurrency Control

- Functions are inherently parallel, which can cause **races**

- Example


```
def main(message: QueueTrigger) -> bool:
    current = storage.read("messagecount");
    current = current + 1;
    storage.write("messagecount", current);
```

- Two concurrent invocations may interleave so that
- Counter incremented only once
- Need **external synchronization** mechanism

Threats to Consistency: Effect Duplication

- At-least-once triggers can lead to effect duplication

- Example


```
def main(message: QueueTrigger) -> bool:
    current = storage.read("messagecount");
    current = current + 1;
    storage.write("messagecount", current);
```

- Over-count number of messages if
 - Invoked multiple times for same message
- Standard advice: Make function **idempotent**
 - Not always easy!


Stateful Abstractions

Cloud Environment is Inherently and Pervasively

- Concurrent
- Parallel
- Distributed
- Failure-prone

Challenging to Achieve

- Clean separation of
 - Application logic from
 - Concurrency, parallelism, distribution, and failures



Motivates Augmented Programming Model

- **Durable Functions** need
- Abstractions for **state** and **synchronization**
 - Hide challenges in the runtime
 - Simplify application development

Joke time

What did the event say to
the serverless function?

“I’ve triggered something
special in you!”

Durable Functions

Microsoft's FaaS Platform

- Durable Functions is a component of Azure Functions
- Supports: Python, C#, JavaScript, PowerShell...

Programming Model

- **Activities** — durable equivalent of stateless FaaS functions
- **Entities** — actors that
 - Encapsulate application state
 - Process operations one at a time
- **Orchestrations** — task-parallel async-await style code to
 - Coordinate **activities** and **entities**

Programming Model

- Durability is implicit with Durable Functions
 - **State** of entities **automatically** saved to storage
 - **Transparently** restored after faults

Activities

- Automatically retried — under partial execution
- Unless time limit exceeded

```
import azure.functions as func
import urllib.request

def main(req: func.HttpRequest):
    url = req.get_body().decode("utf-8")
    fid = urllib.request.urlopen(url)
    webpage = fid.read().decode('utf-8')
    found_free = webpage.find("free") >= 0
    return str(found_free)
```

```
"bindings": [ { "authLevel": "anonymous",
                 "name": "req",
                 "type": "httpTrigger",
                 "direction": "in",
                 "route": "checkWebPage",
                 "methods": [ "get" ] },
                { "name": "$return",
                 "type": "http",
                 "direction": "out" } ]
```


Orchestrations

- Decompose computation into tasks: *activities, entity operations, timers, sub-orchestrations*
- Example: Sequencing three activities DownloadData, Process, and Summarize

```
def orchestrate_pipeline(context: df.DurableOrchestrationContext):
    try:
        dataset = yield context.call_activity('DownloadData')
        outputs = yield context.call_activity('Process', dataset)
        summary = yield context.call_activity('Summarize', outputs)
        return summary
    except Exception as exc:
        yield context.call_activity('CleanUp')
        return f"Something went wrong" {exc}"
```

Entities

- Allow applications to **encapsulate** durable state
- Define **operations** on them
 - Operation requests stored in a queue
 - Execute them one at time

```
class Account:
    # Initialize the entity state, which defaults to 0 balance
    def __init__(self):
        self.balance = 0

    # Return the balance
    def get(self):
        return self.balance

    # Deposit an amount, increasing the balance
    def deposit(self, amount):
        self.balance += amount
        return self.balance

    # Withdraw an amount, reducing the balance
    def withdraw(self, amount):
        self.balance -= amount
        return self.balance
```

Critical Sections

Concurrent Invocations Require Atomicity

- Invocations across **entities** must happen at once, atomically to
 - Prevent violating consistency
- Critical sections — regions of code where
 - Only **one orchestration** can call specific entities

Concurrent Invocations Require Atomicity

```
def transfer_safe_orchestration(context: df.DurableOrchestrationContext):
    # From input, get entity ids for source and destination account, and the amount
    [source, dest, amount] = context.get_input()
    # Critical Section: Acquire a lock for each entity
    with (yield context.lock([source, dest])):
        # Make sure that the source account has adequate balance to avoid overdraft
        source_balance = yield context.call_entity(source, "get")
        if source_balance < amount:
            return False
        else:
            # Wait for transfer to complete
            yield context.task_all([context.call_entity(source, "withdraw", amount),
                                   context.call_entity(dest, "deposit", amount)]);
            return True
```

Serverless Technologies

Proprietary

- AWS Lambda
- Azure Functions
- Google Cloud Run

Edge Computing

- Cloudflare Workers
- Fastly Compute@Edge
- AWS Lambda@Edge

Open-Source

- OpenFaaS
- Apache OpenWhisk
- Kubeless
- Knative
- Fission
- Serverless Framework
(Framework Agnostic)

Questions?

- Burckhardt et al, Durable Functions: Semantics for Stateful Serverless, OOPSLA 2021